

LIVER CIRRHOSIS STAGE PREDICTION with 95% Accuracy

Abstract

Liver cirrhosis is a progressive and irreversible condition that results from chronic liver diseases, leading to fibrosis, functional decline, and eventual organ failure (Schuppan & Afdhal, 2008). The accurate staging of cirrhosis is critical for clinical decision-making, guiding therapeutic interventions, and improving patient prognosis (Tsochatzis et al., 2014). Traditional diagnostic methods such as biopsy, imaging, and laboratory tests, while effective, present limitations in accuracy, invasiveness, and inter-observer variability (Marcellin & Kutala, 2018). Machine learning (ML) has emerged as a powerful tool in medical diagnostics, offering automated, accurate, and scalable solutions for disease classification (Esteva et al., 2019). This study explores the potential of ML models in classifying the histologic stages of liver cirrhosis using a dataset from a Mayo Clinic study on primary biliary cirrhosis (PBC), collected between 1974 and 1984, and publicly available on Kaggle (Dickson et al., 1989). The dataset contains clinical and laboratory parameters, including demographic details (age, sex), biochemical markers (bilirubin, albumin, cholesterol, SGOT, platelets), and disease-related attributes (ascites, edema, hepatomegaly, and spider angiomas). Before model training, data preprocessing was performed, including handling missing values, encoding categorical variables, and normalizing numerical attributes. Additionally, synthetic data augmentation techniques were applied to address class imbalances and improve model robustness (He et al., 2009). Logistic Regression yielded an accuracy of 60%, indicating limited efficacy in this context. The Decision Tree Classifier improved performance with an 88% accuracy but posed a risk of overfitting. The Random Forest Classifier achieved the highest accuracy at 95%, demonstrating superior precision, recall, and F1-scores across all classes, and

effectively mitigating overfitting. The K-Nearest Neighbors (KNN) algorithm also performed well with an 89% accuracy, though it is computationally intensive for larger datasets. Gradient Boosting Classifier attained an 83% accuracy, outperforming Logistic Regression but underperforming compared to other models. These findings underscore the efficacy of ensemble tree-based methods, particularly Random Forest, in accurately staging liver cirrhosis. The effectiveness of ML-based approaches in staging liver cirrhosis, particularly ensemble tree-based methods such as Random Forest, which demonstrated superior predictive performance. The study underscores the potential of ML in augmenting traditional diagnostic frameworks, offering a non-invasive, data-driven alternative for early cirrhosis detection and staging. Future research should explore deep learning architectures and external validation using independent datasets to enhance the robustness and clinical applicability of ML models in hepatology (Liu et al., 2022).

Introduction

Liver cirrhosis is a progressive and irreversible condition characterized by extensive fibrosis and the formation of regenerative nodules, which ultimately impair liver function. It represents the final stage of chronic liver diseases (CLDs) and is a leading cause of morbidity and mortality worldwide (Schuppan & Afdhal, 2008). The disease is a consequence of sustained liver injury due to various etiologies, including chronic hepatitis B and C infections, excessive alcohol consumption, and metabolic dysfunction-associated steatotic liver disease (MASLD) (Tsochatzis et al., 2014). Liver cirrhosis is a major global health concern, with approximately 2 million deaths annually attributed to cirrhosis-related complications such as liver failure, portal hypertension, and hepatocellular carcinoma (Asrani et al., 2019).

The pathophysiology of cirrhosis is driven by a persistent cycle of liver injury, inflammation, and wound-healing responses, leading to excessive deposition of extracellular matrix (ECM) proteins, primarily collagen (Bataller & Brenner, 2005). Hepatic stellate cells (HSCs) play a central role in fibrosis progression, as they transform into myofibroblast-like cells that secrete ECM components in response to chronic injury (Hernandez-Gea & Friedman, 2011). Over time, fibrotic tissue disrupts the normal hepatic architecture, leading to increased intrahepatic resistance and portal hypertension, a hallmark of cirrhosis.

Clinically, cirrhosis progresses through compensated and decompensated stages. The compensated stage is often asymptomatic, with normal liver function despite fibrotic damage. As the disease advances, patients develop complications such as ascites, hepatic encephalopathy, and variceal bleeding, indicating decompensation and an increased risk of liver-related mortality (D'Amico et al., 2018). Diagnosis of cirrhosis typically involves a combination of clinical assessment, biochemical markers, imaging techniques (e.g., elastography), and liver biopsy when necessary (European Association for the Study of the Liver [EASL], 2015).

Problem Statement

Liver cirrhosis is a chronic and progressive liver disease that leads to liver failure if not detected and managed in its early stages. Traditional diagnostic approaches, such as biopsies and blood tests, often have limitations in accuracy, invasiveness, and cost-effectiveness. With the advancement of machine learning (ML), automated models can assist in classifying liver cirrhosis stages more accurately and efficiently. However, selecting the best ML model remains a challenge due to differences in model performance, feature importance, and computational efficiency. This study evaluates multiple ML models to determine the most effective approach for Liver Cirrhosis Stage Prediction using a dataset containing key clinical features such as Bilirubin, Albumin, Platelets, and Prothrombin levels, as well as patient demographics and medical conditions. The dataset used in this study contains various medical and demographic attributes that serve as potential predictors for liver cirrhosis stages. These attributes include patient demographics (age, sex), clinical symptoms (ascites, hepatomegaly, edema, and spider angiomas), liver function biomarkers (bilirubin, albumin, SGOT, and prothrombin), metabolic markers (cholesterol, triglycerides, alkaline phosphatase, and copper), and hematological factors (platelet count) (Pimpin et al., 2018).

N_Days – Represents the number of days the patient has been under observation in the study. This can be useful for understanding disease progression over time.

Status – Indicates the patient's current condition, which can be:

- *Alive (patient is still living)
- *Dead (patient has passed away)
- *Transplanted (patient has undergone a liver transplant)

Drug – Specifies the type of treatment the patient has received. This can help analyze whether certain treatments affect disease progression.

Age – The patient's age at the time of observation, which may be a risk factor for disease severity.

Sex – The biological sex of the patient (Male/Female), which can influence disease presentation and progression.

Ascites – A binary indicator (Yes/No) of fluid accumulation in the abdominal cavity, a common symptom of liver cirrhosis.

Hepatomegaly – A binary indicator (Yes/No) of an enlarged liver, often associated with liver disease.

Spiders – A binary indicator (Yes/No) for the presence of spider angiomas, which are small, dilated blood vessels seen in liver disease patients.

Edema – A categorical variable indicating the presence of swelling due to fluid retention. Possible values:

- * No edema
- * Edema not responding to diuretics
- * Edema responding to diuretics

Bilirubin – A liver function biomarker indicating the level of bilirubin in the blood. High bilirubin levels suggest impaired liver function.

Cholesterol – Measures the cholesterol level in the blood. Liver dysfunction can cause abnormal cholesterol metabolism.

Albumin – A protein produced by the liver, essential for maintaining blood volume and transporting substances. Low levels indicate poor liver function.

Copper – Measures copper levels in the blood. The liver plays a key role in copper metabolism, and abnormal levels may indicate liver disease.

Alk_Phos (Alkaline Phosphatase) – An enzyme that indicates liver function. Elevated levels may signal bile duct obstruction or liver damage.

SGOT (Serum Glutamic Oxaloacetic Transaminase) – Also known as **AST (Aspartate Aminotransferase)**, this enzyme is released when liver cells are damaged.

Tryglicerides – The level of triglycerides (a type of fat) in the blood, which may be affected by liver disease.

Platelets – The number of platelets in the blood, important for clotting. Low platelet count is common in liver cirrhosis.

Prothrombin – The time taken for blood to clot. A prolonged prothrombin time indicates impaired liver function and a higher risk of bleeding.

Stage – The target variable representing different stages of liver cirrhosis. The number of stages may vary depending on the classification system used (e.g., mild, moderate, severe).

This dataset includes a mix of categorical and numerical variables, all of which are crucial for predicting liver cirrhosis stages using machine learning models

Methodology

The following Python libraries were utilized for various stages of data processing, analysis, and model building

Libraries Used

Pandas: A powerful Python library used for data manipulation and analysis, particularly for handling tabular data. It was used for loading, cleaning, and transforming the dataset.
NumPy: A library for numerical operations and handling arrays. It was used for mathematical functions and numerical computations on data.
Matplotlib: A plotting library used for creating static, interactive, and animated visualizations, including histograms, scatter plots, and box plots, to explore and visualize data.
Seaborn: A statistical data

visualization library built on top of Matplotlib, used for creating advanced visualizations such as heatmaps and pair plots to analyze relationships in data. Scikit-learn: A machine learning library used for model building and evaluation. It provides tools for training and evaluating classifiers, including Decision Trees, Logistic Regression, SVM, Random Forest, and K-Nearest Neighbors (KNN). Plotly: A graphing library used for creating interactive visualizations. It was used for generating subplots to represent complex data in an interactive format. itertools: A Python module used for creating iterators that efficiently loop through data. It was utilized to generate combinations and permutations during the analysis. Warnings: A module used to filter out unnecessary warnings during the execution of the code, ensuring cleaner output and focus during the analysis. Scikit-learn: A comprehensive machine learning library that provides tools for data preprocessing, model building, and evaluation. It was used for training various classifiers such as Decision Tree, Logistic Regression, Support Vector Machine (SVM), Random Forest, and K-Nearest Neighbors (KNN). It also provided functions for model evaluation through metrics like confusion matrix, ROC curve, precision-recall curve, and AUC.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import itertools
import warnings
warnings.filterwarnings("ignore")
```

Data Collection

The dataset for this project was collected from two reliable sources:

1. **Cirrhosis Patient Survival Prediction Dataset** from the UCI Machine Learning Repository, which provides data on patients diagnosed with cirrhosis. This dataset contains information related to the survival prediction of these patients based on various medical features. The dataset can be accessed [here](#).
2. **Liver Cirrhosis Stage Classification Dataset** from Kaggle, which offers data about liver cirrhosis stages, containing records related to different stages of cirrhosis in patients. This dataset is useful for predicting the stage of cirrhosis based on patient attributes. The dataset can be accessed [here](#).


These datasets provide valuable features for analyzing liver cirrhosis and predicting patient outcomes, which are pivotal for the development of predictive models.

```
In [2]: df = pd.read_csv(r"C:\Users\manoj\Downloads\Liver_Cirrhosis\liver_cirrhosis.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin	Cholesterol	Albumin	Copper	Alk_F
0	2221	C	Placebo	18499	F	N	Y	N	N	0.5	149.0	4.04	227.0	5
1	1230	C	Placebo	19724	M	Y	N	Y	N	0.5	219.0	3.93	22.0	6
2	4184	C	Placebo	11839	F	N	N	N	N	0.5	320.0	3.54	51.0	12
3	2090	D	Placebo	16467	F	N	N	N	N	0.7	255.0	3.74	23.0	10
4	2105	D	Placebo	21699	F	N	Y	N	N	1.9	486.0	3.54	74.0	10



```
In [4]: df.tail()
```

Out[4]:

	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin	Cholesterol	Albumin	Cop
24995	3584	D	D- penicillamine	23612	F	N	N	N	N	0.8	231.000000	3.87	17
24996	3584	D	D- penicillamine	23612	F	N	N	N	N	0.8	231.000000	3.87	17
24997	971	D	D- penicillamine	16736	F	N	Y	Y	Y	5.1	369.510563	3.23	1
24998	3707	C	D- penicillamine	16990	F	N	Y	N	N	0.8	315.000000	4.24	1
24999	3707	C	D- penicillamine	16990	F	N	Y	N	N	0.8	315.000000	4.24	1

In [5]:

```
df.info()
```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25000 entries, 0 to 24999
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   N_Days                25000 non-null  int64  
1   Status                25000 non-null  object  
2   Drug                 25000 non-null  object  
3   Age                  25000 non-null  int64  
4   Sex                  25000 non-null  object  
5   Ascites              25000 non-null  object  
6   Hepatomegaly         25000 non-null  object  
7   Spiders              25000 non-null  object  
8   Edema                25000 non-null  object  
9   Bilirubin            25000 non-null  float64 
10  Cholesterol           25000 non-null  float64 
11  Albumin              25000 non-null  float64 
12  Copper               25000 non-null  float64 
13  Alk_Phos             25000 non-null  float64 
14  SGOT                 25000 non-null  float64 
15  Tryglicerides        25000 non-null  float64 
16  Platelets            25000 non-null  float64 
17  Prothrombin          25000 non-null  float64 
18  Stage                25000 non-null  int64  
dtypes: float64(9), int64(3), object(7)
memory usage: 3.6+ MB

```

```
In [6]: df.columns
```

```

Out[6]: Index(['N_Days', 'Status', 'Drug', 'Age', 'Sex', 'Ascites', 'Hepatomegaly',
              'Spiders', 'Edema', 'Bilirubin', 'Cholesterol', 'Albumin', 'Copper',
              'Alk_Phos', 'SGOT', 'Tryglicerides', 'Platelets', 'Prothrombin',
              'Stage'],
              dtype='object')

```

DATA CLEANING

```
In [7]: df.shape
```

```
Out[7]: (25000, 19)
```

```
In [8]: df.shape[0]
```

```
Out[8]: 25000
```

rows are 2500

```
In [10]: df.shape[1]
```

```
Out[10]: 19
```

coumns are 19

```
In [12]: df.isnull()
```

Out[12]:

	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin	Cholesterol	Albumin	Copper	All
0	False	False	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False	False	False
...
24995	False	False	False	False	False	False	False	False	False	False	False	False	False	False
24996	False	False	False	False	False	False	False	False	False	False	False	False	False	False
24997	False	False	False	False	False	False	False	False	False	False	False	False	False	False
24998	False	False	False	False	False	False	False	False	False	False	False	False	False	False
24999	False	False	False	False	False	False	False	False	False	False	False	False	False	False

25000 rows × 19 columns



In [13]:

```
df.isnull().sum()
```

```
Out[13]: N_Days      0
         Status      0
         Drug        0
         Age         0
         Sex         0
         Ascites     0
         Hepatomegaly 0
         Spiders     0
         Edema       0
         Bilirubin   0
         Cholesterol 0
         Albumin     0
         Copper      0
         Alk_Phos    0
         SGOT        0
         Tryglicerides 0
         Platelets   0
         Prothrombin 0
         Stage       0
         dtype: int64
```

```
In [14]: df = df.dropna()
```

```
In [15]: df.describe()
```

Out[15]:

	N_Days	Age	Bilirubin	Cholesterol	Albumin	Copper	Alk_Phos	SGOT	Tryglice
count	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000
mean	1887.117040	18495.877080	3.402644	372.331471	3.486578	100.184663	1995.675597	123.166345	123.820000
std	1091.690918	3737.596616	4.707491	193.668452	0.380488	73.184840	1798.885660	47.747616	52.780000
min	41.000000	9598.000000	0.300000	120.000000	1.960000	4.000000	289.000000	26.350000	33.000000
25%	1080.000000	15694.000000	0.800000	275.000000	3.290000	52.000000	1032.000000	92.000000	92.000000
50%	1680.000000	18499.000000	1.300000	369.510563	3.510000	97.648387	1828.000000	122.556346	124.700000
75%	2576.000000	20955.000000	3.400000	369.510563	3.750000	107.000000	1982.655769	134.850000	127.000000
max	4795.000000	28650.000000	28.000000	1775.000000	4.640000	588.000000	13862.400000	457.250000	598.000000



In [16]: df.columns

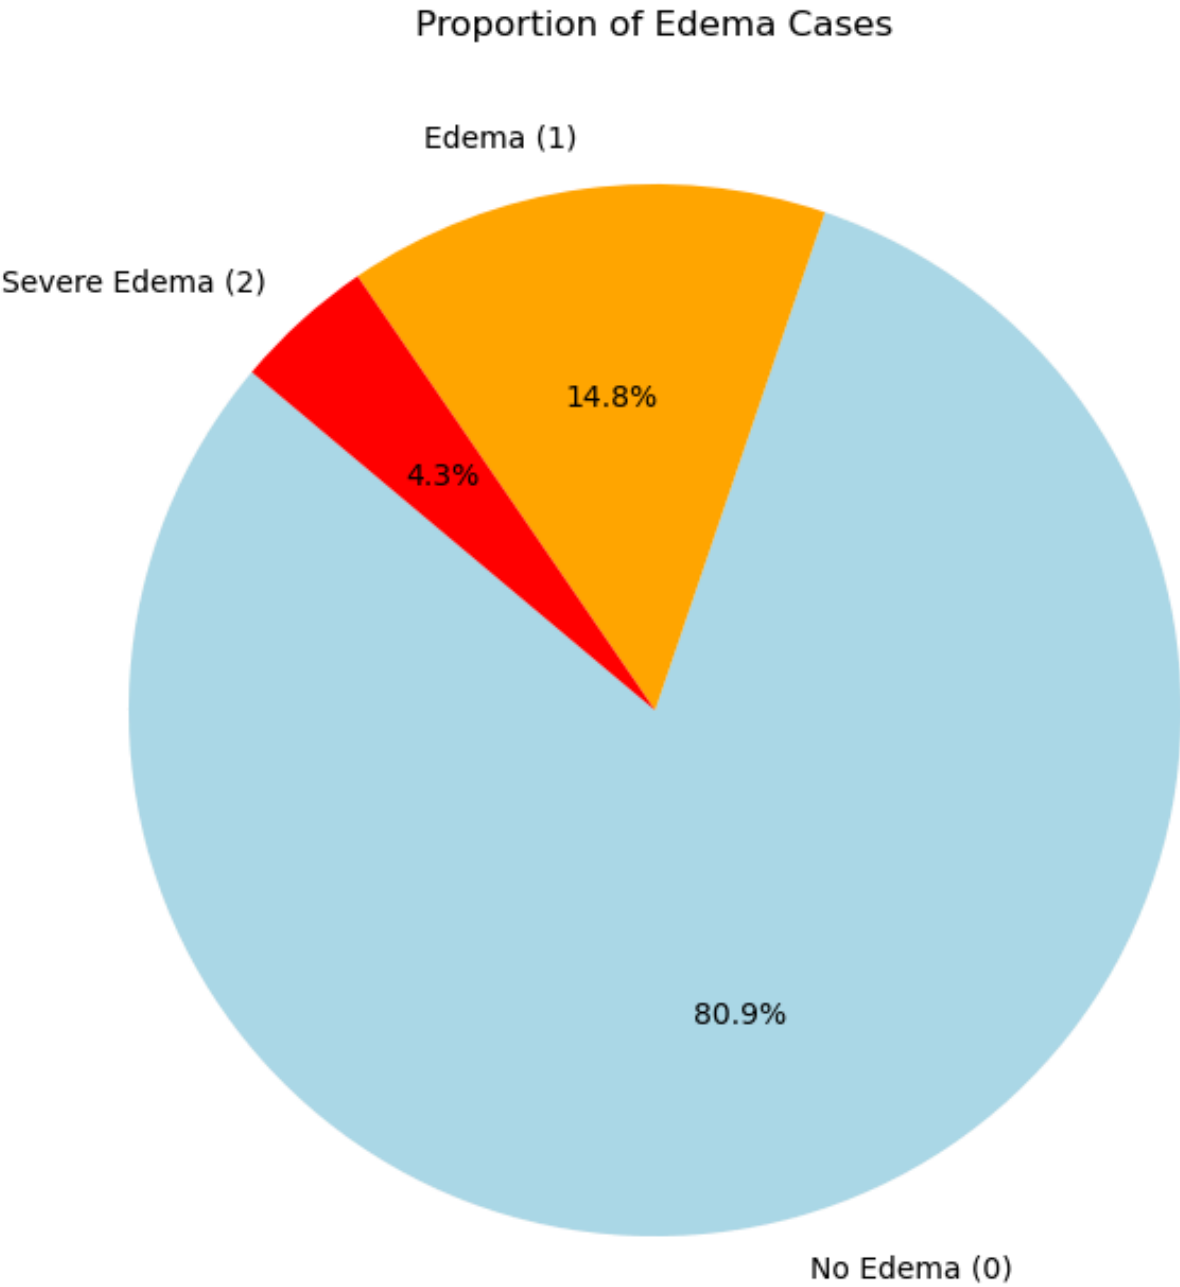
Out[16]: Index(['N_Days', 'Status', 'Drug', 'Age', 'Sex', 'Ascites', 'Hepatomegaly',
 'Spiders', 'Edema', 'Bilirubin', 'Cholesterol', 'Albumin', 'Copper',
 'Alk_Phos', 'SGOT', 'Tryglicerides', 'Platelets', 'Prothrombin',
 'Stage'],
 dtype='object')

In [17]: df.dtypes

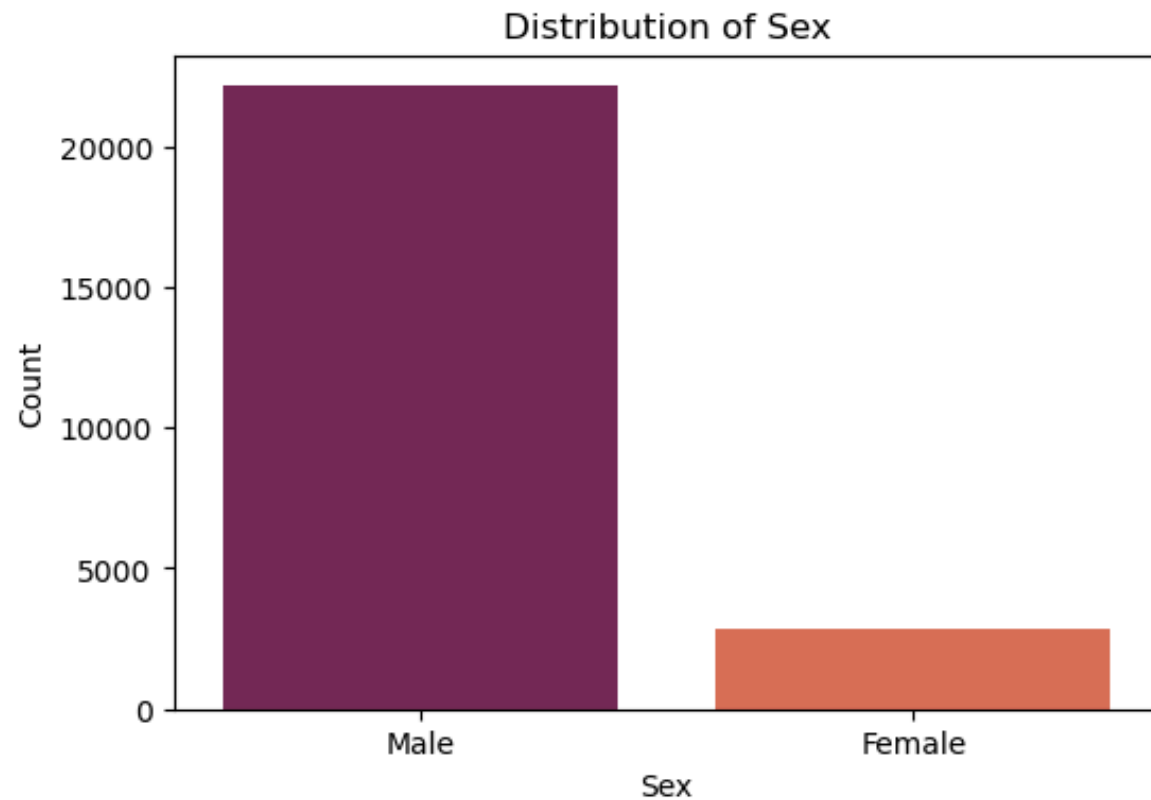
```
Out[17]: N_Days          int64
         Status         object
         Drug           object
         Age            int64
         Sex            object
         Ascites        object
         Hepatomegaly   object
         Spiders        object
         Edema          object
         Bilirubin      float64
         Cholesterol    float64
         Albumin        float64
         Copper         float64
         Alk_Phos       float64
         SGOT           float64
         Tryglicerides  float64
         Platelets      float64
         Prothrombin    float64
         Stage          int64
         dtype: object
```

EXPLORATORY DATA ANALYSIS

```
In [19]: edema_counts = df['Edema'].value_counts()
         labels = ['No Edema (0)', 'Edema (1)', 'Severe Edema (2)']
         plt.figure(figsize=(8,9))
         colors = ['lightblue', 'orange', 'red']
         plt.pie(edema_counts, labels=labels, autopct='%1.1f%%', colors=colors, startangle=140)
         plt.title("Proportion of Edema Cases")
         plt.show()
```



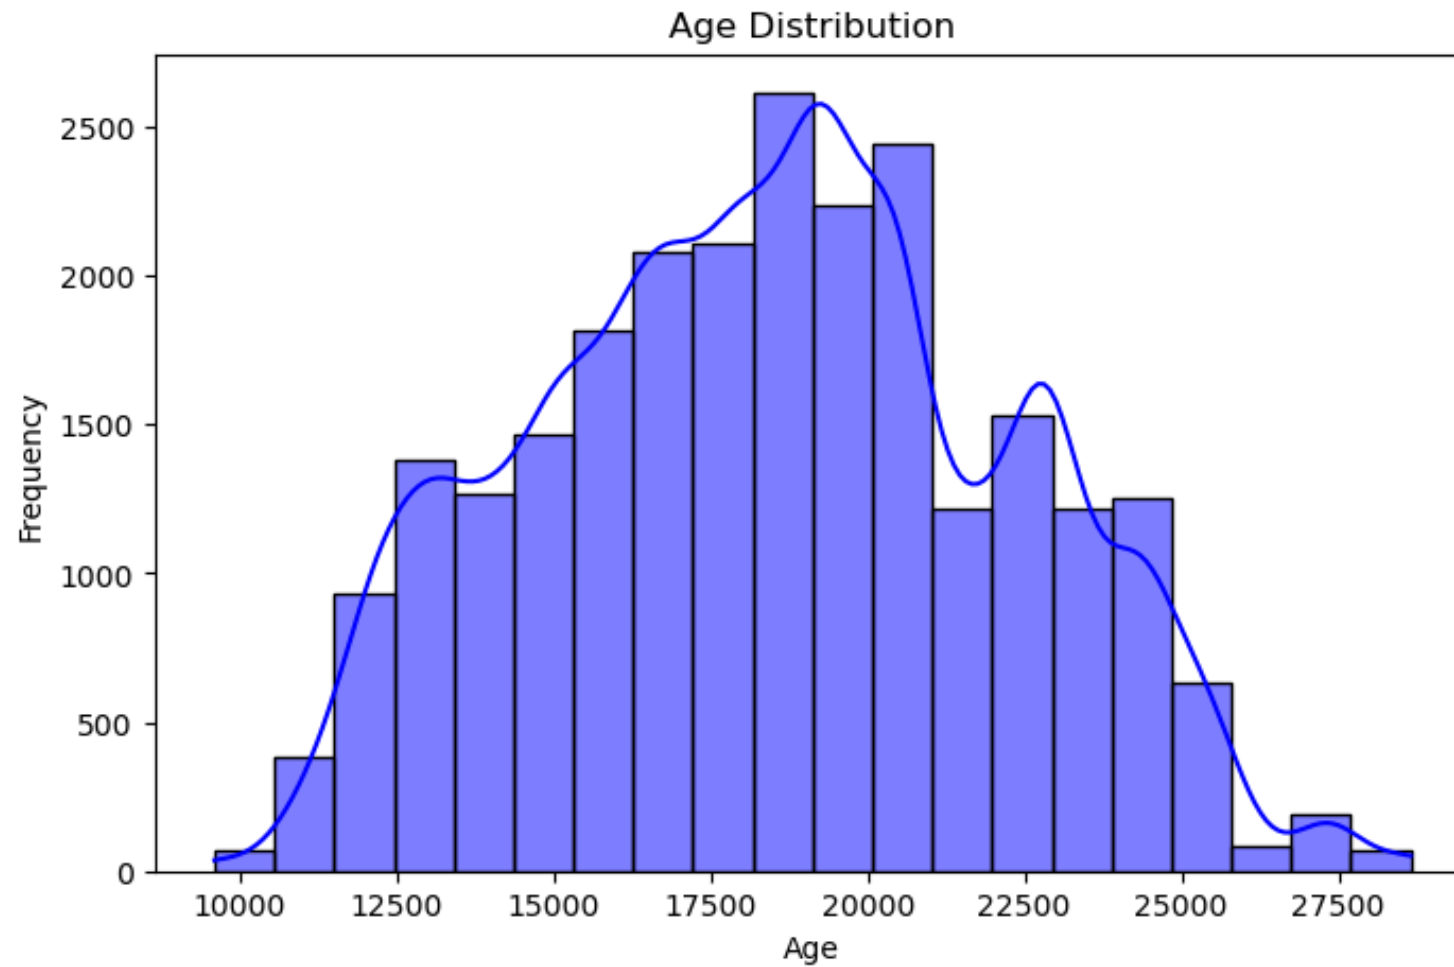
```
In [20]: plt.figure(figsize=(6, 4))
sns.countplot(x='Sex', data=df, palette="rocket")
plt.xticks(ticks=[0, 1], labels=["Male", "Female"])
plt.title("Distribution of Sex")
plt.xlabel("Sex")
plt.ylabel("Count")
plt.show()
```



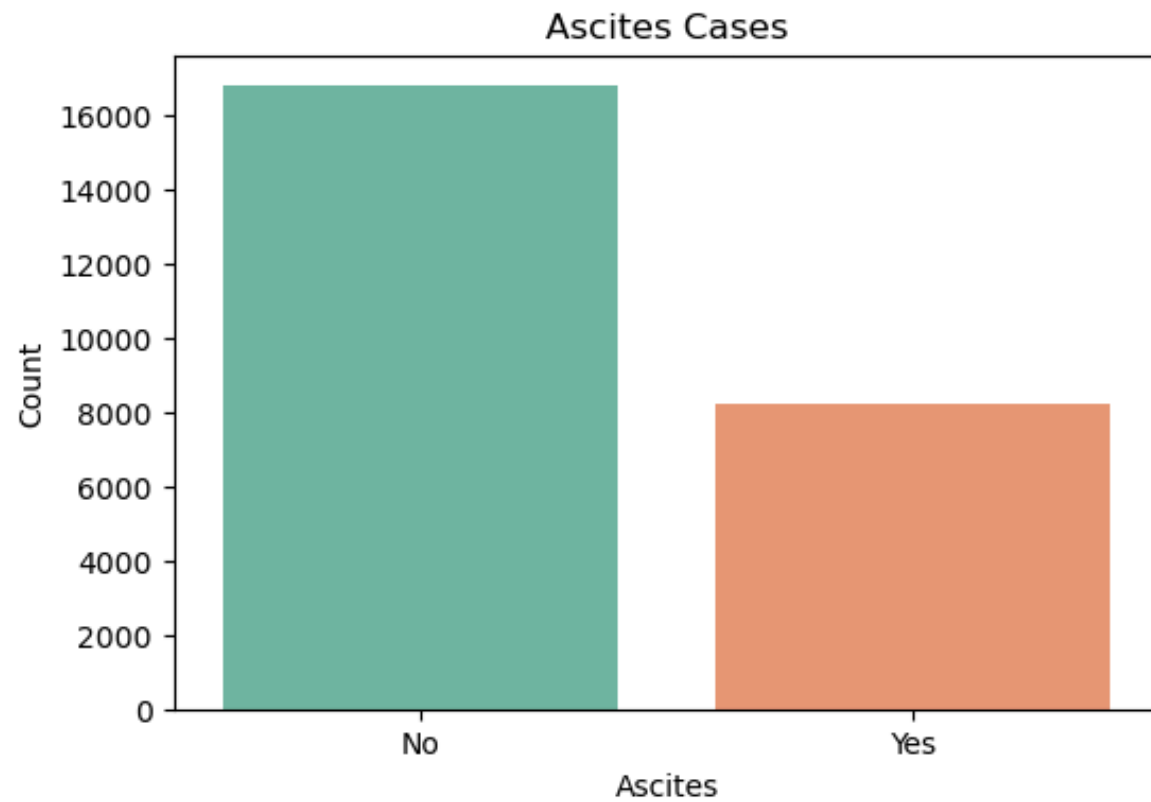
```
In [21]: plt.figure(figsize=(8, 5))
sns.histplot(df['Age'], bins=20, kde=True, color="blue")
plt.title("Age Distribution")
plt.xlabel("Age")
```



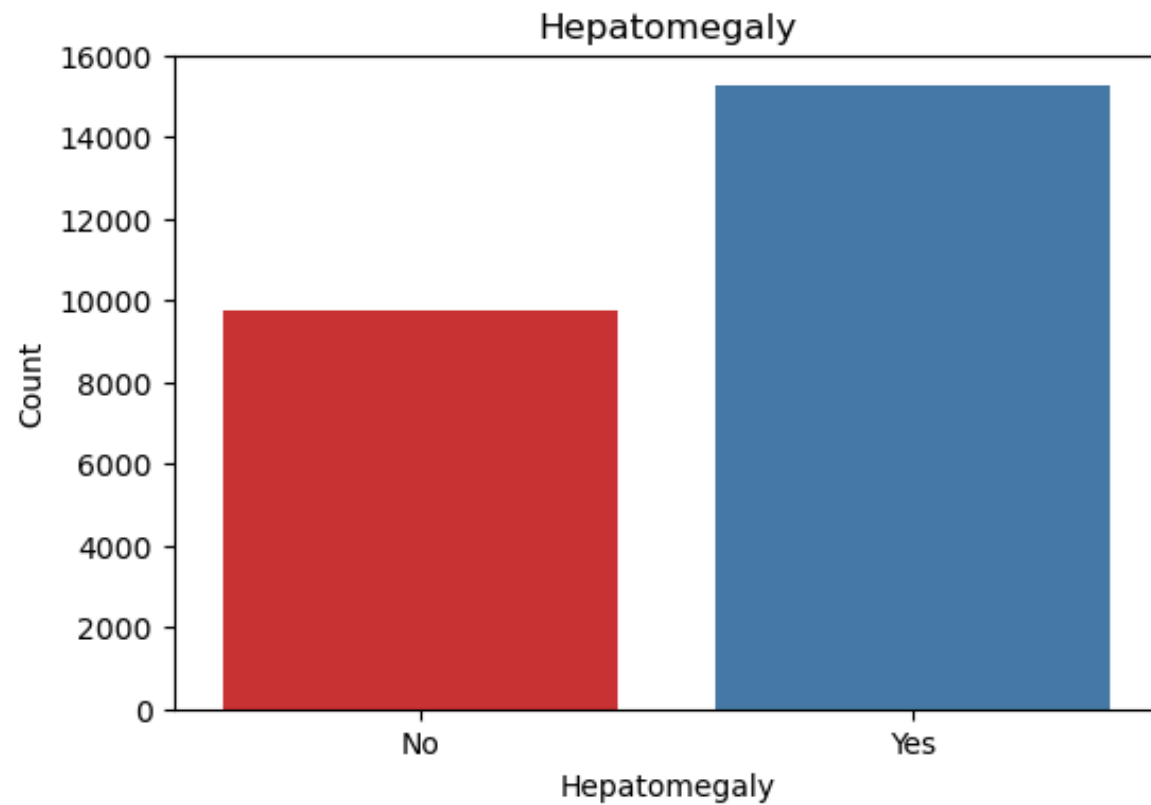
```
plt.ylabel("Frequency")  
plt.show()
```



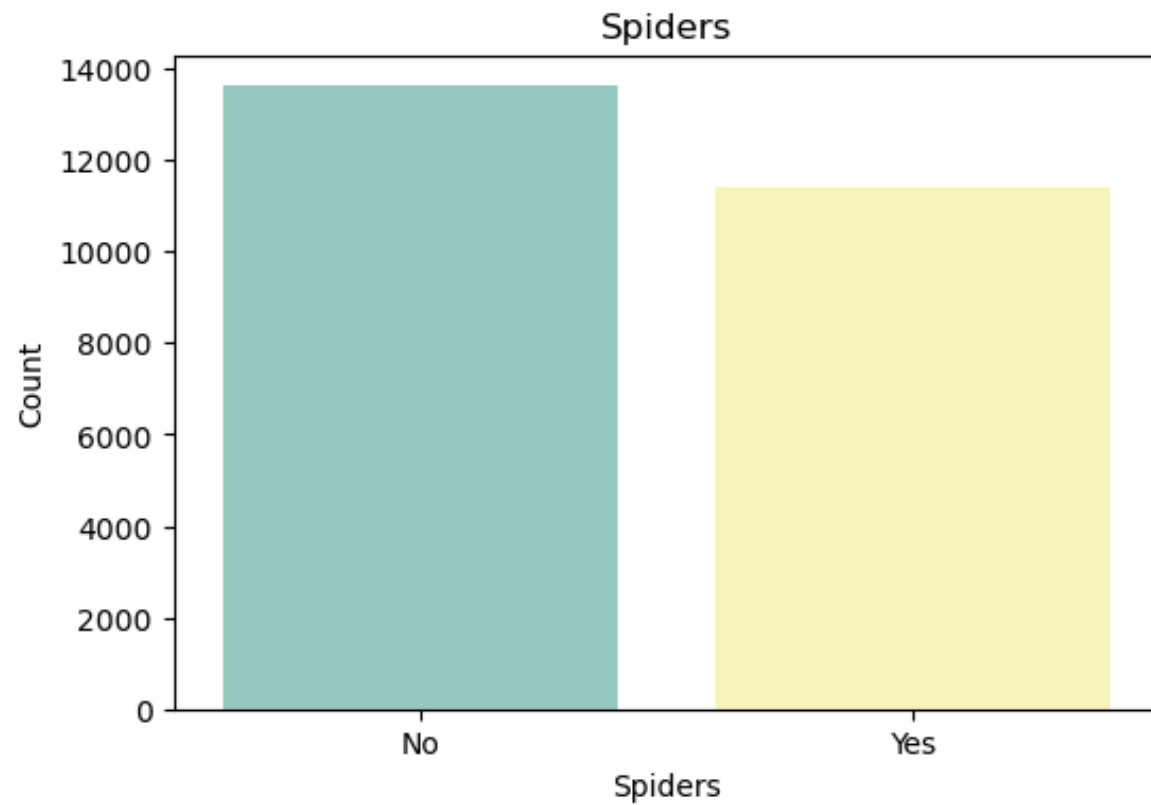
```
In [22]: plt.figure(figsize=(6, 4))  
sns.countplot(x='Ascites', data=df, palette="Set2")  
plt.xticks(ticks=[0, 1], labels=["No", "Yes"])  
plt.title("Ascites Cases")  
plt.xlabel("Ascites")  
plt.ylabel("Count")  
plt.show()
```



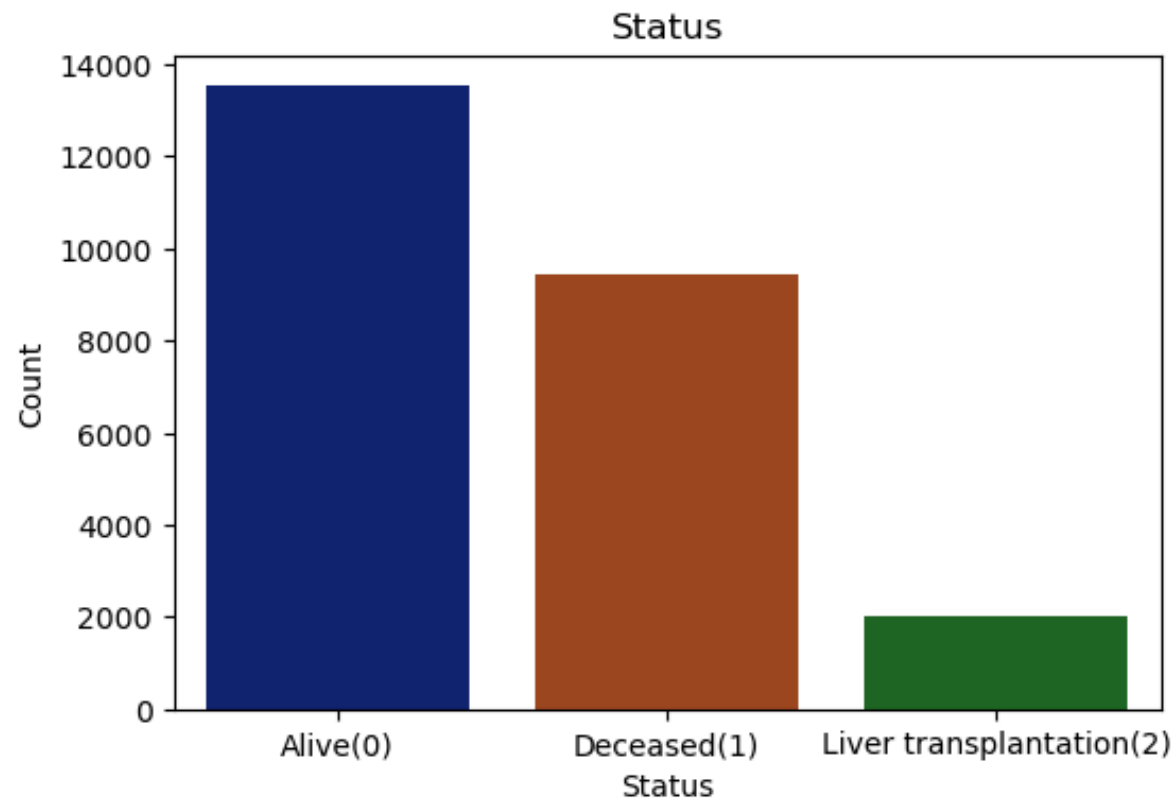
```
In [23]: plt.figure(figsize=(6, 4))
sns.countplot(x='Hepatomegaly', data=df, palette="Set1")
plt.xticks(ticks=[0, 1], labels=["No", "Yes"])
plt.title("Hepatomegaly")
plt.xlabel("Hepatomegaly")
plt.ylabel("Count")
plt.show()
```



```
In [24]: plt.figure(figsize=(6, 4))
sns.countplot(x='Spiders', data=df, palette="Set3")
plt.xticks(ticks=[0, 1], labels=["No", "Yes"])
plt.title("Spiders")
plt.xlabel("Spiders")
plt.ylabel("Count")
plt.show()
```



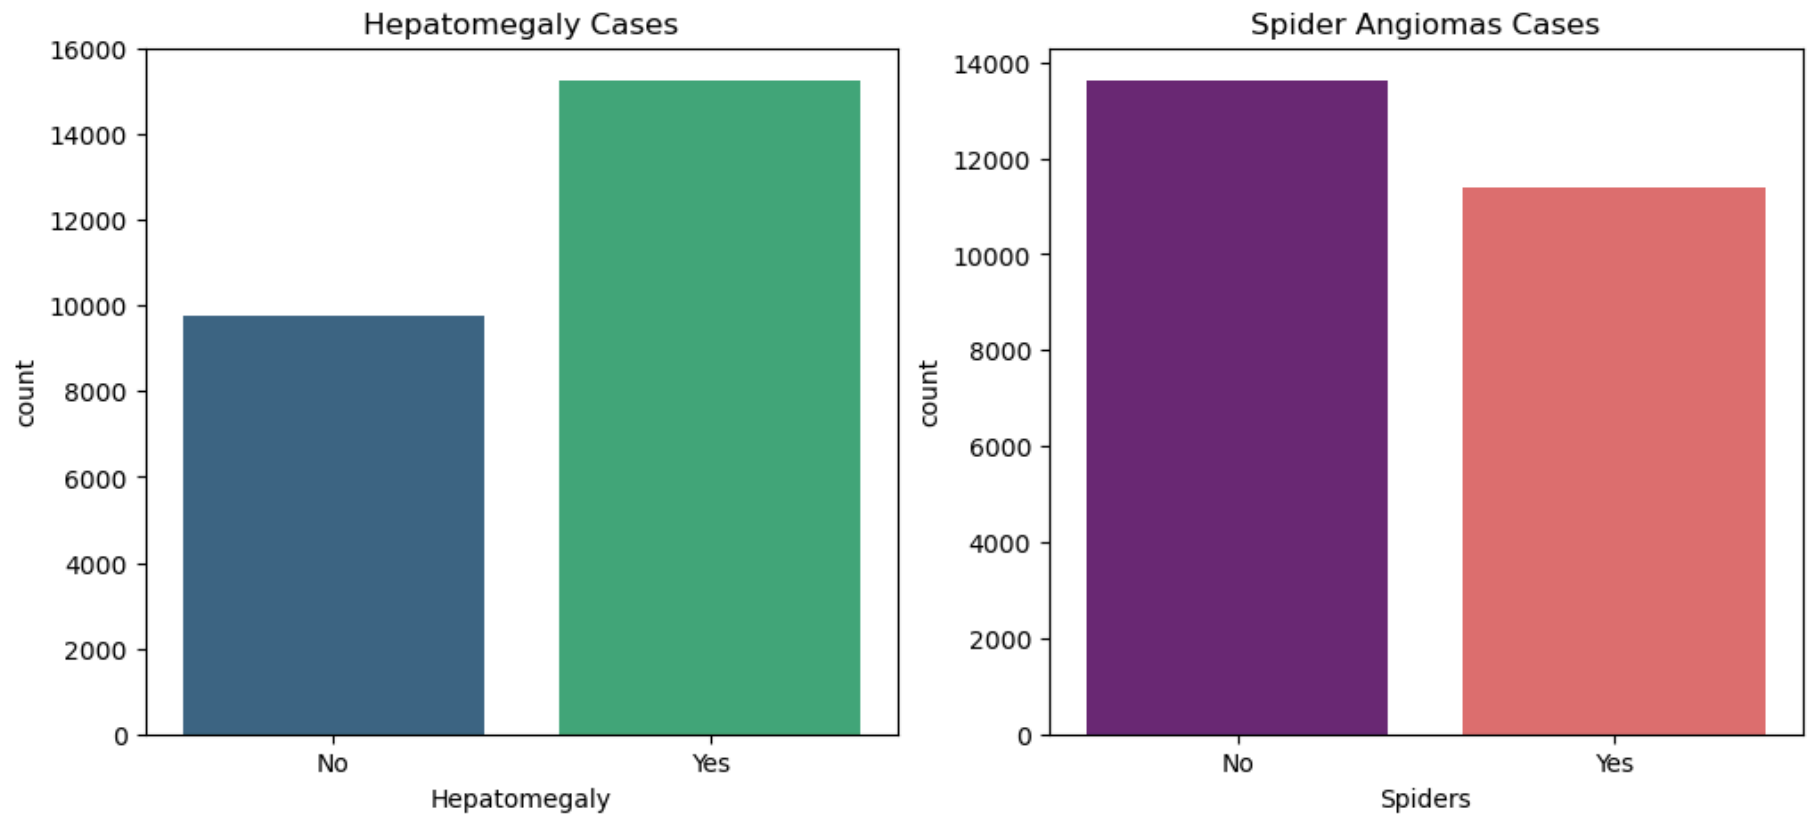
```
In [25]: status_counts = df['Status'].value_counts()
plt.figure(figsize=(6, 4))
sns.countplot(x='Status', data=df, palette="dark")
plt.xticks(ticks=[0, 1, 2], labels=["Alive(0)", "Deceased(1)", "Liver transplantation(2)"])
plt.title("Status")
plt.xlabel("Status")
plt.ylabel("Count")
plt.show()
```



```
In [26]: fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.countplot(x='Hepatomegaly', data=df, ax=axes[0], palette="viridis")
axes[0].set_xticks([0, 1])
axes[0].set_xticklabels(["No", "Yes"])
axes[0].set_title("Hepatomegaly Cases")

sns.countplot(x='Spiders', data=df, ax=axes[1], palette="magma")
axes[1].set_xticks([0, 1])
axes[1].set_xticklabels(["No", "Yes"])
axes[1].set_title("Spider Angiomas Cases")
```

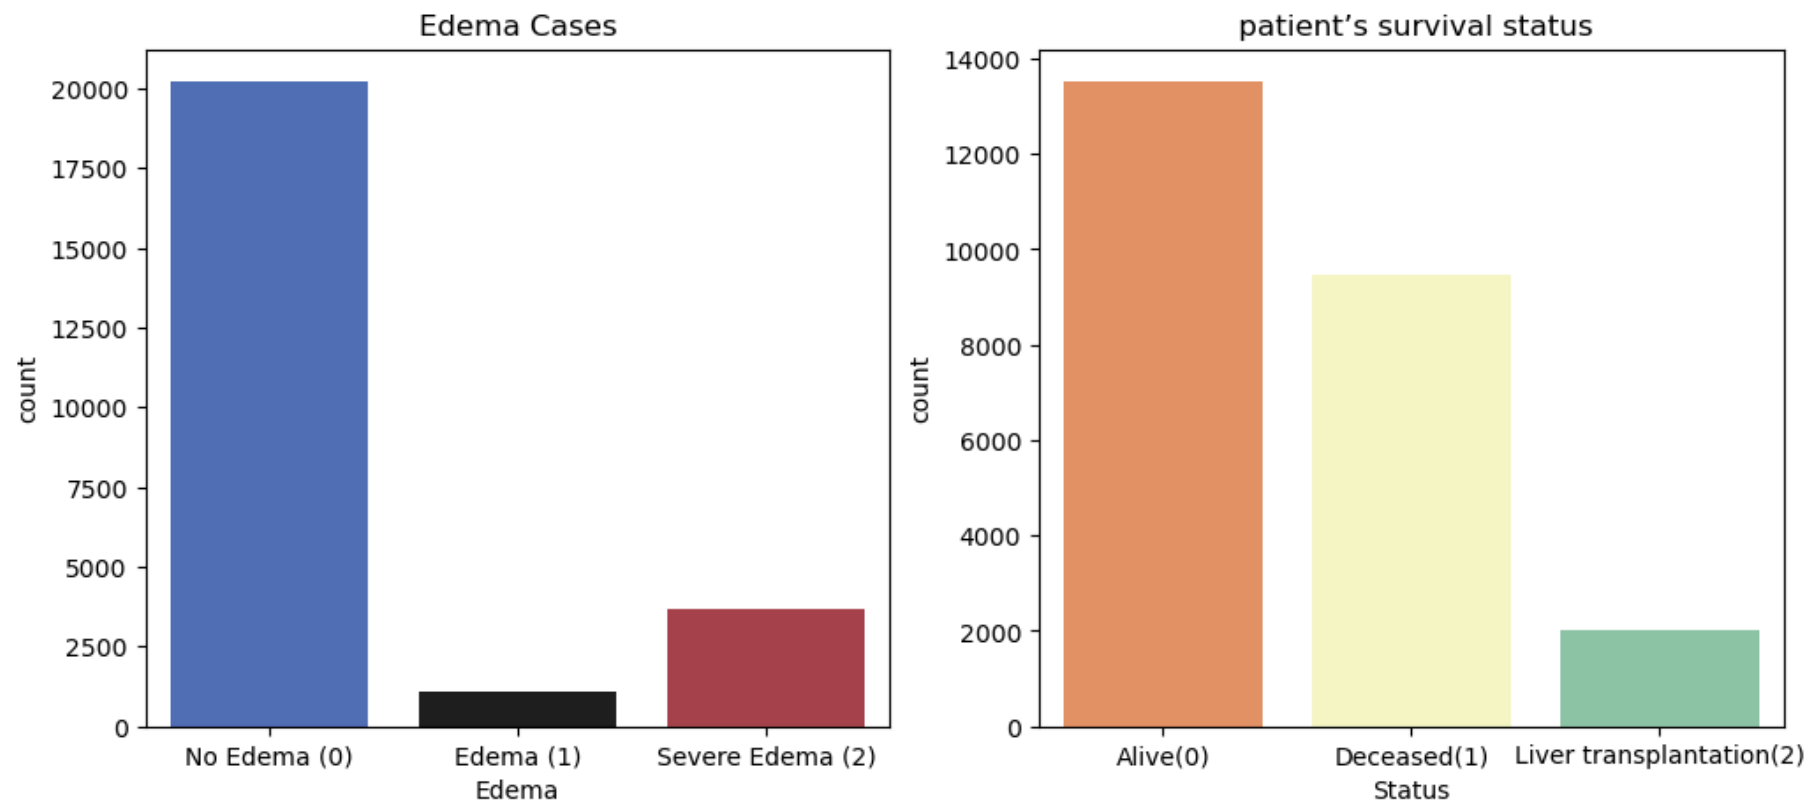
```
Out[26]: Text(0.5, 1.0, 'Spider Angiomas Cases')
```



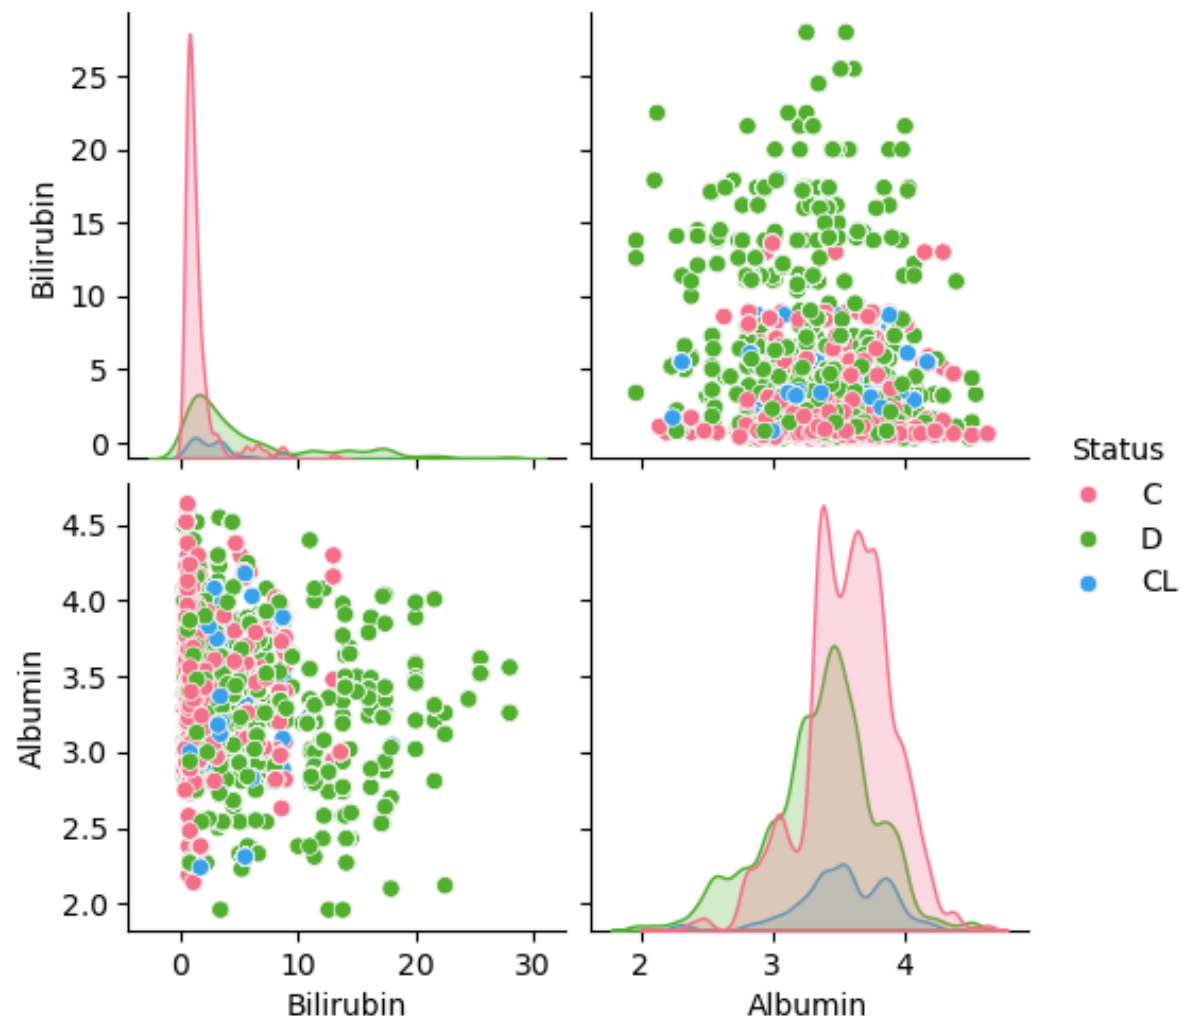
```
In [27]: fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.countplot(x='Edema', data=df, ax=axes[0], palette="icefire")
axes[0].set_xticks([0, 1, 2])
axes[0].set_xticklabels(['No Edema (0)', 'Edema (1)', 'Severe Edema (2)'])
axes[0].set_title("Edema Cases")

sns.countplot(x='Status', data=df, ax=axes[1], palette="Spectral")
axes[1].set_xticks([0, 1, 2])
axes[1].set_xticklabels(["Alive(0)", "Deceased(1)", "Liver transplantation(2)"])
axes[1].set_title(" patient's survival status")
```

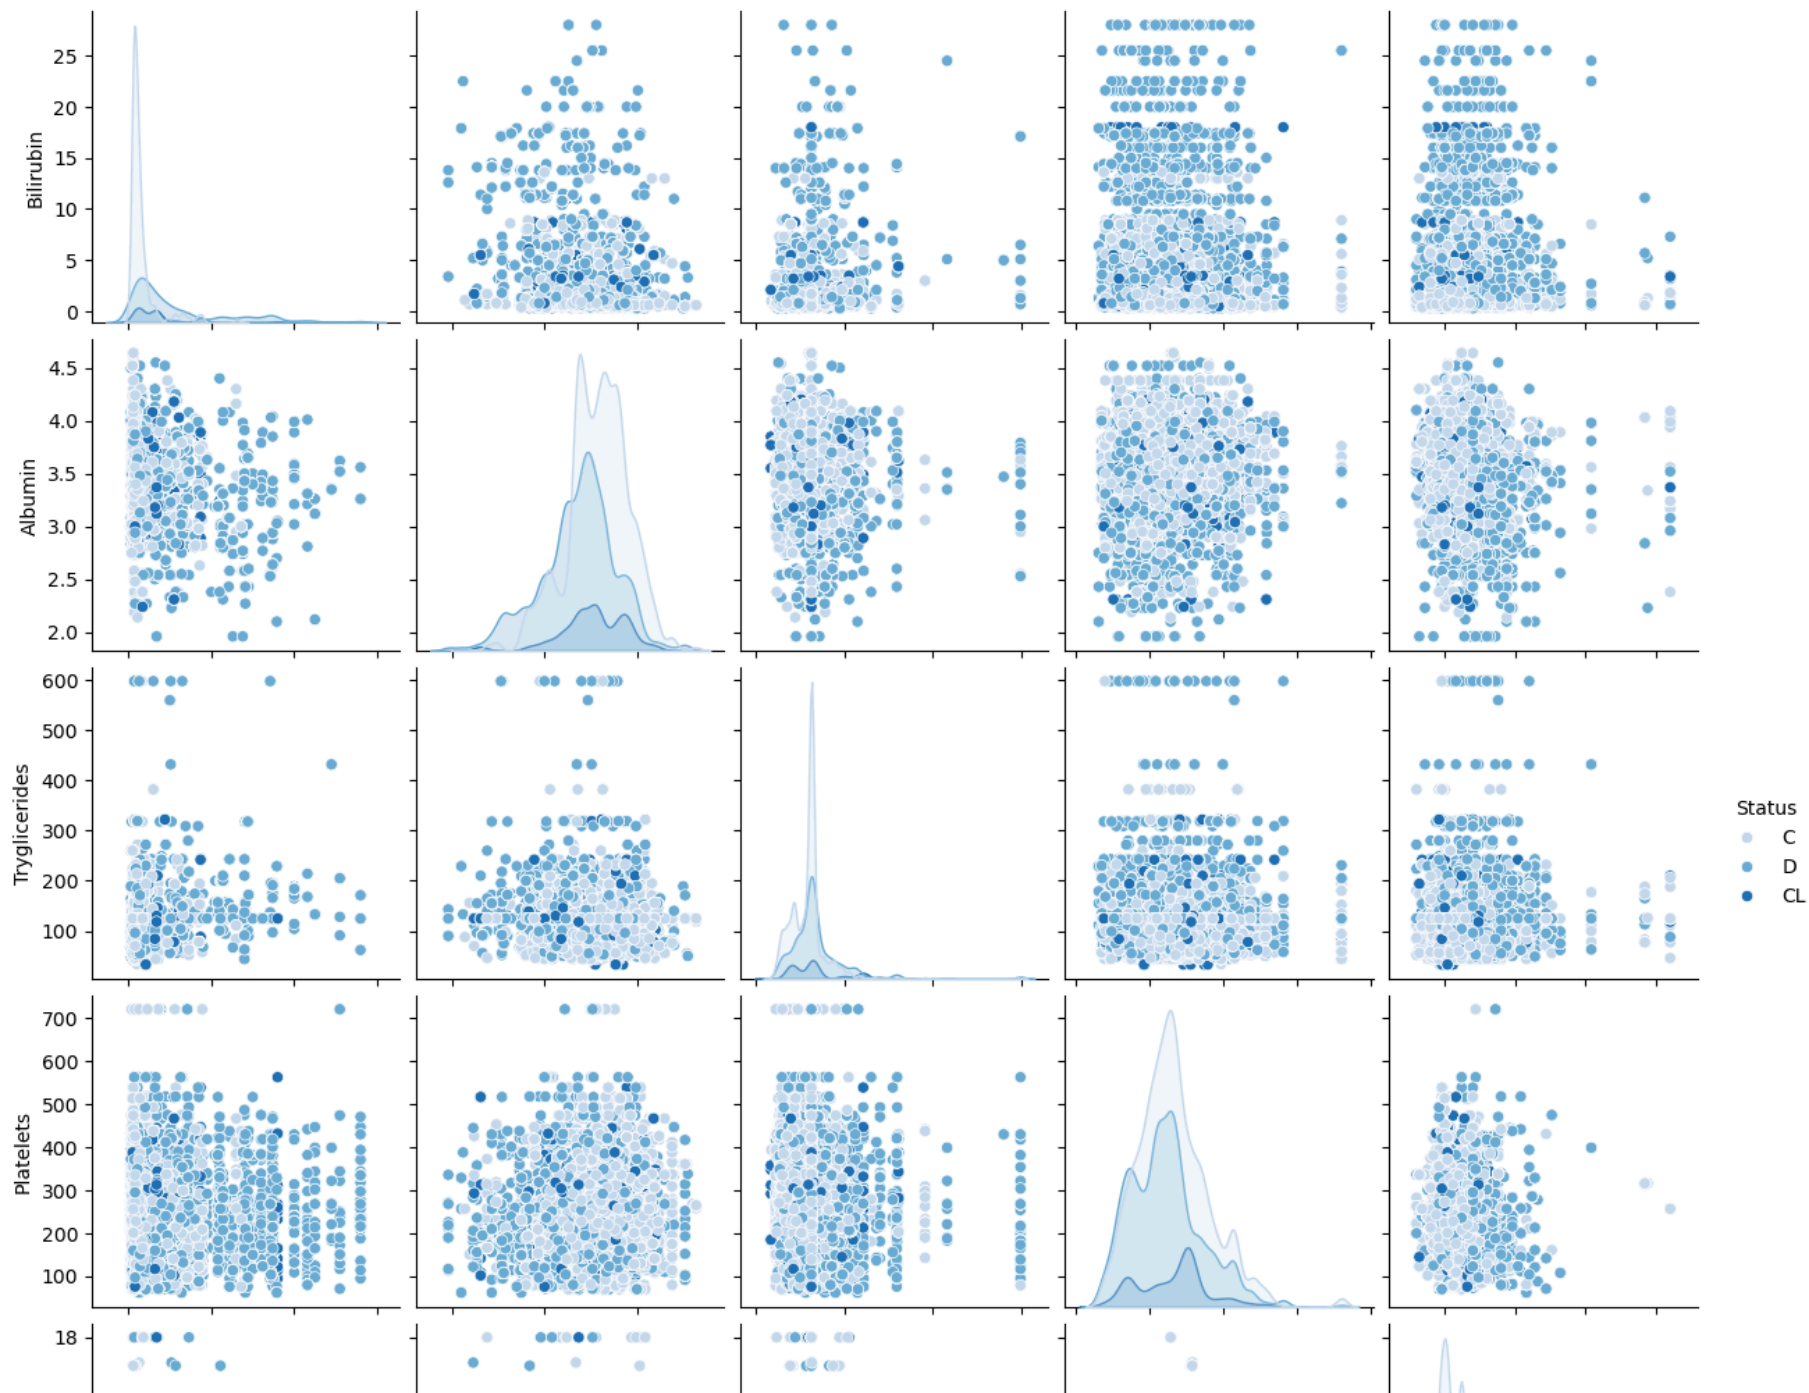
```
Out[27]: Text(0.5, 1.0, ' patient's survival status')
```

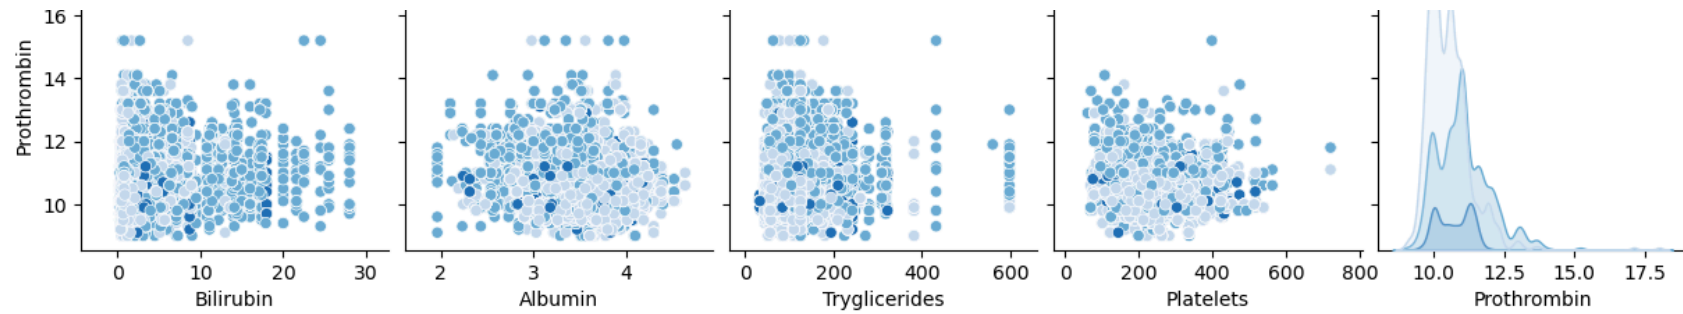


```
In [28]: sns.pairplot(df, vars=['Bilirubin', 'Albumin'], hue="Status", palette="husl")  
plt.show()
```

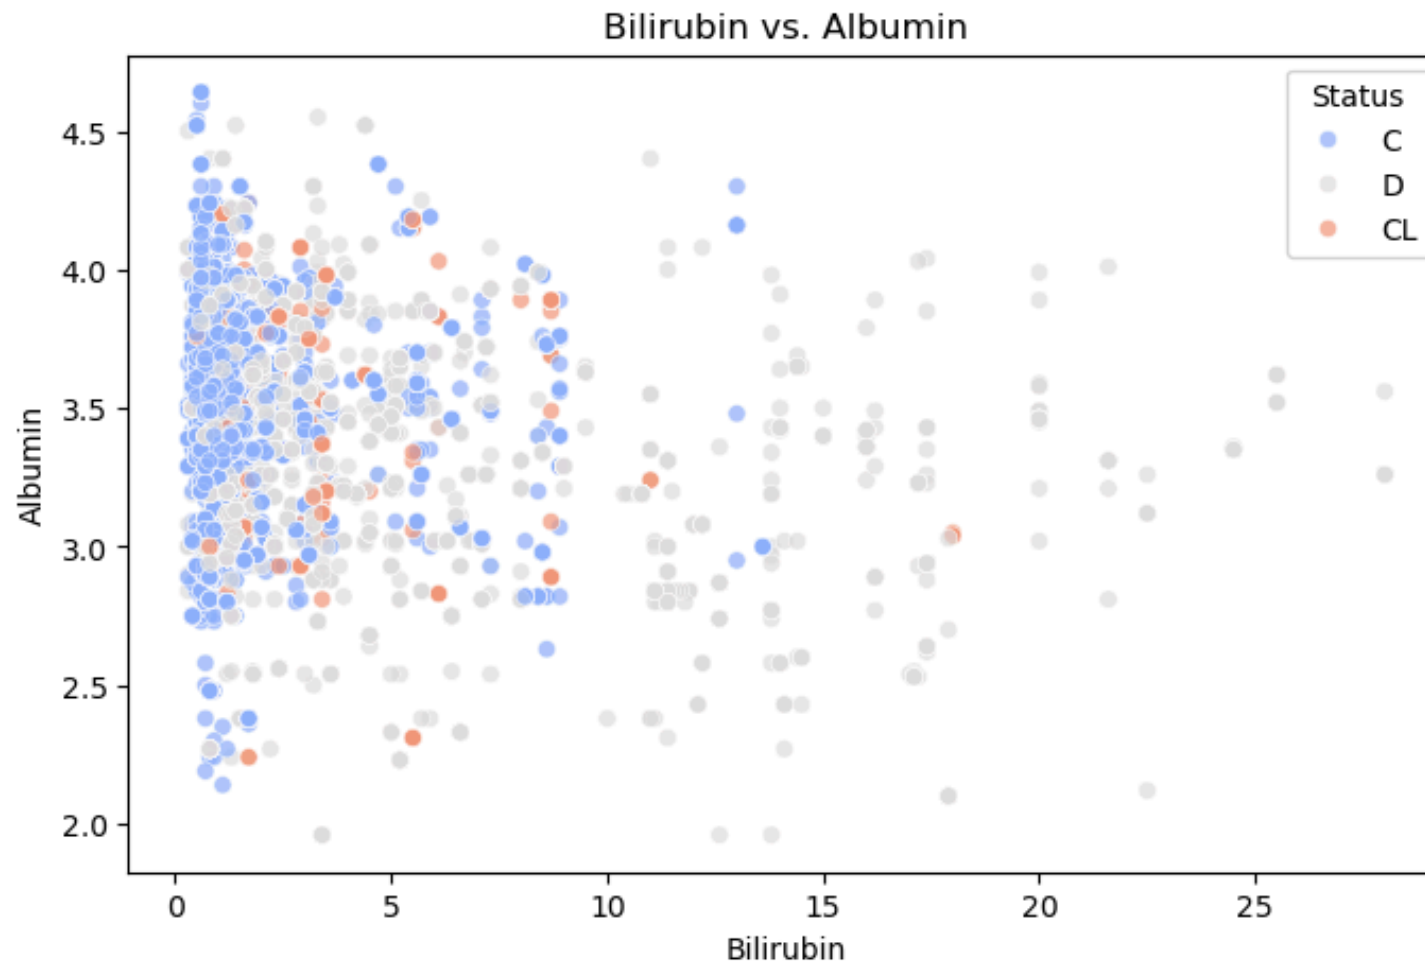


```
In [29]: sns.pairplot(df, vars=['Bilirubin', 'Albumin', 'Tryglicerides', 'Platelets', 'Prothrombin'],  
                hue="Status", palette="Blues")  
plt.show()
```

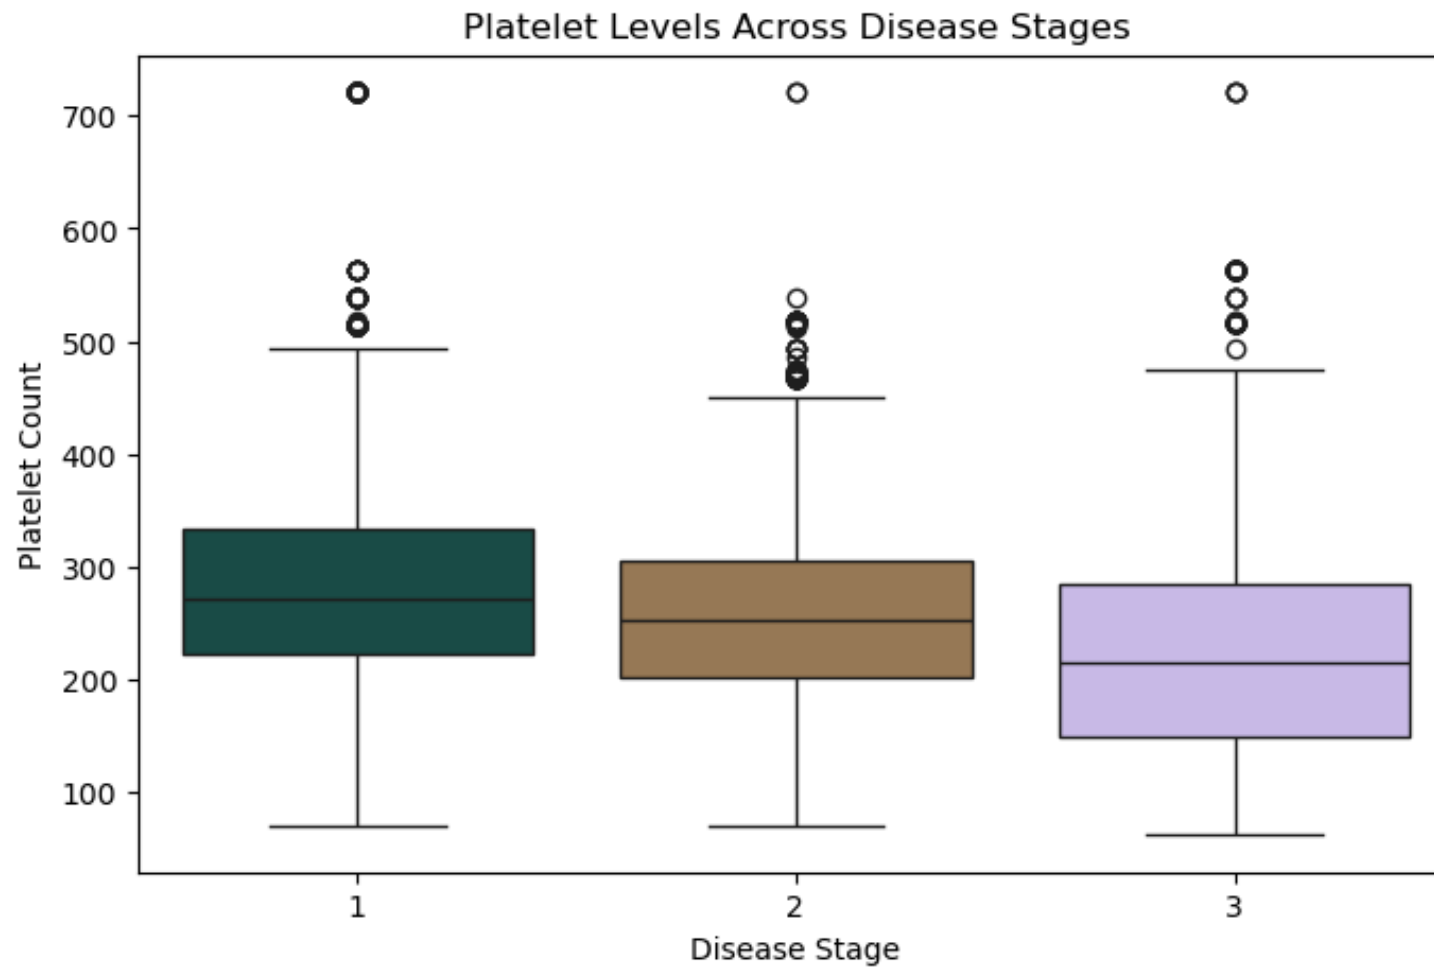





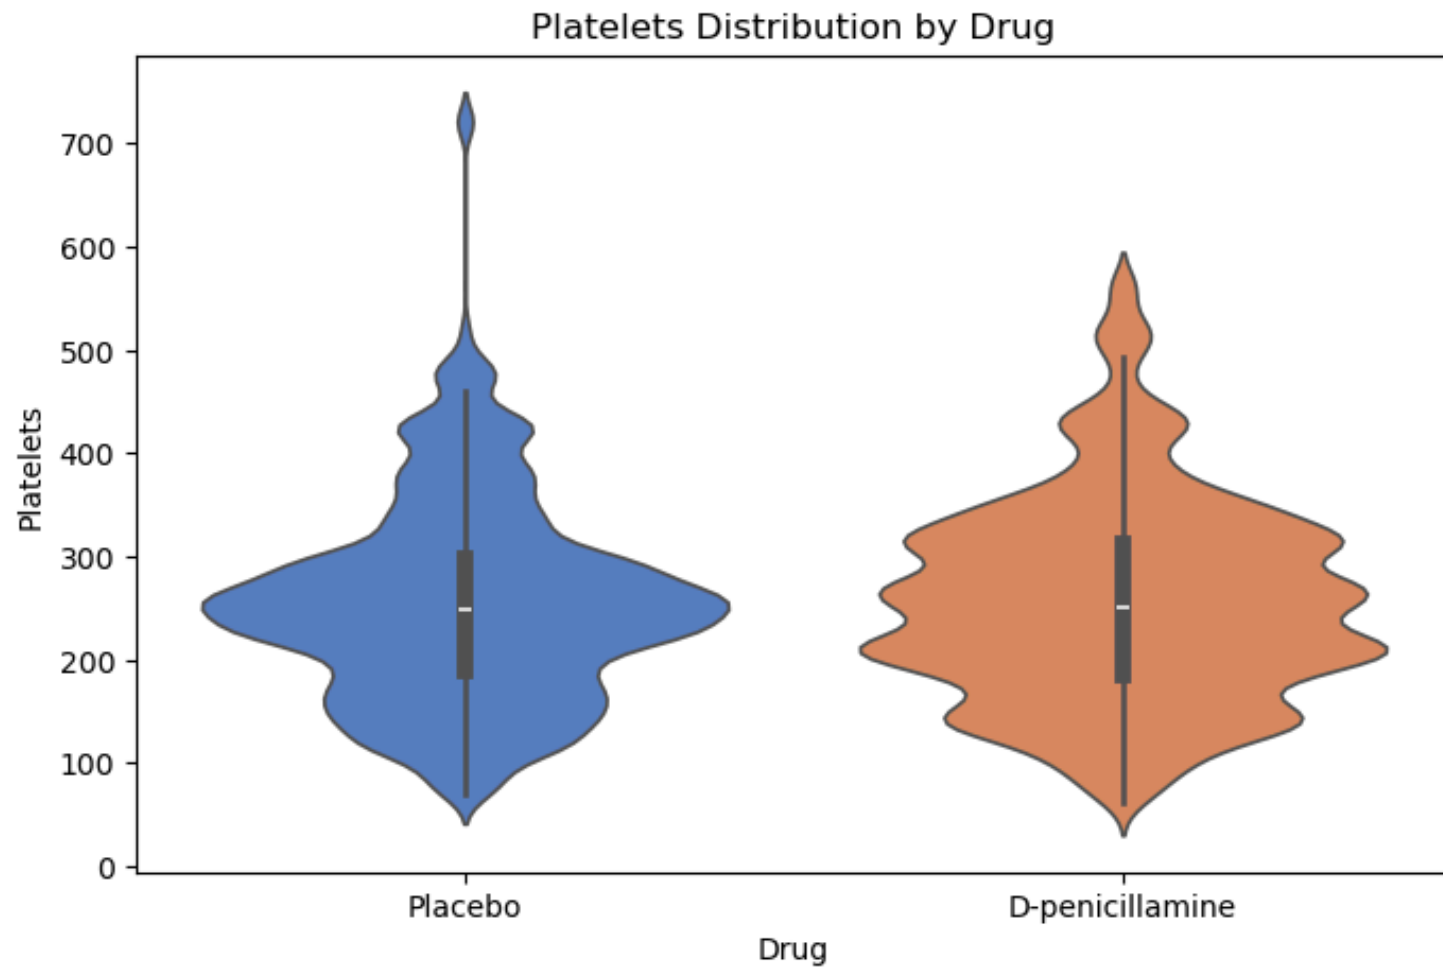
```
In [30]: plt.figure(figsize=(8, 5))
sns.scatterplot(x='Bilirubin', y='Albumin', data=df, hue='Status', palette="coolwarm", alpha=0.7)
plt.title("Bilirubin vs. Albumin")
plt.xlabel("Bilirubin")
plt.ylabel("Albumin")
plt.show()
```



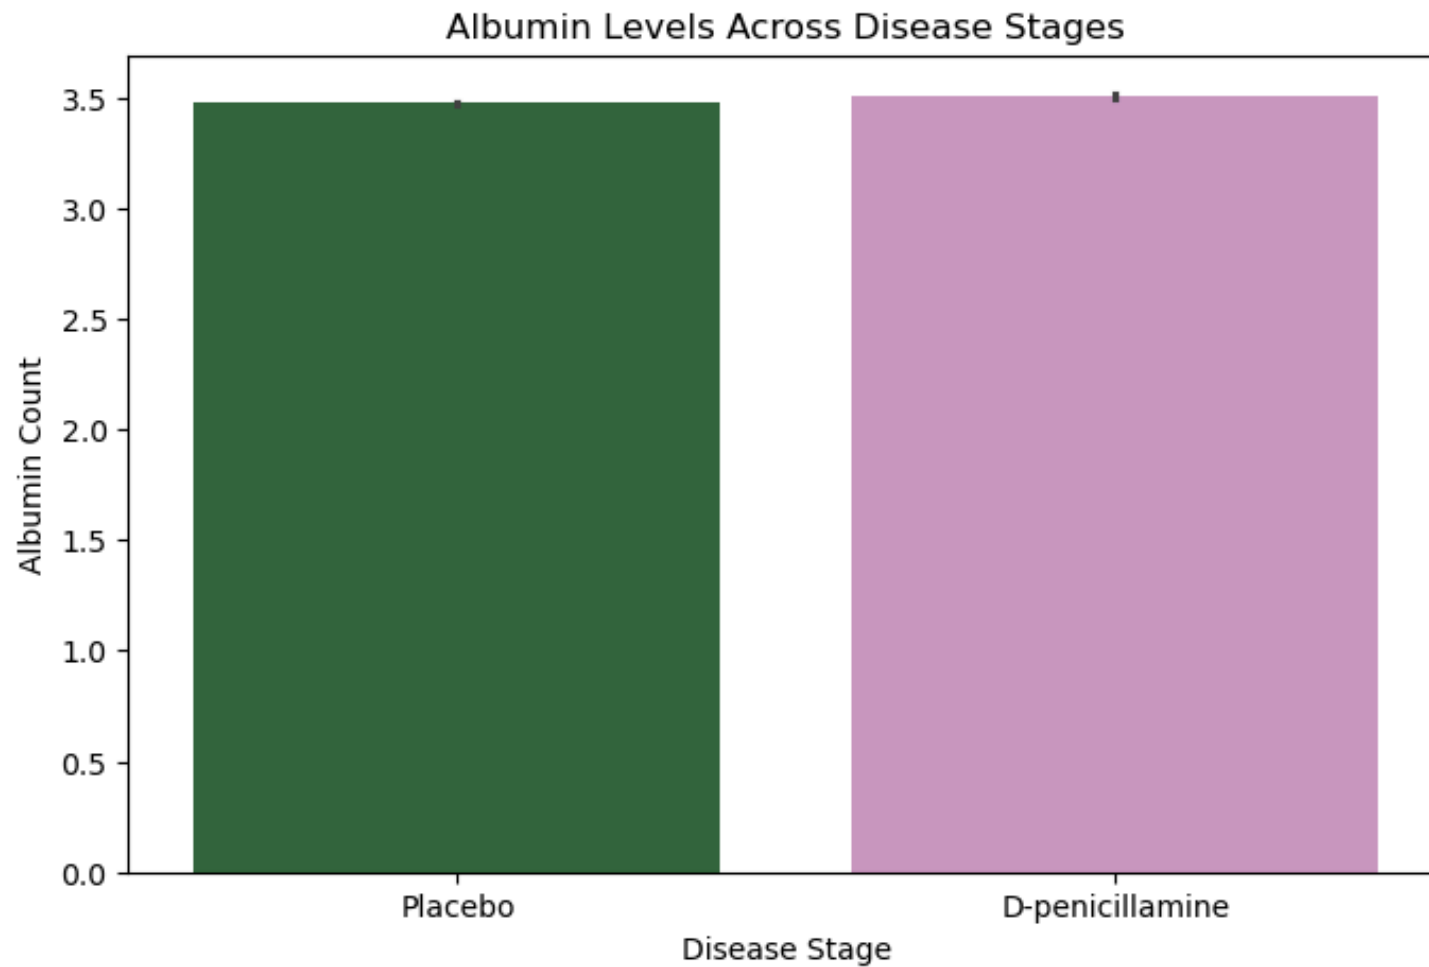
```
In [31]: plt.figure(figsize=(8, 5))
sns.boxplot(x='Stage', y='Platelets', data=df, palette="cubehelix")
plt.title("Platelet Levels Across Disease Stages")
plt.xlabel("Disease Stage")
plt.ylabel("Platelet Count")
plt.show()
```



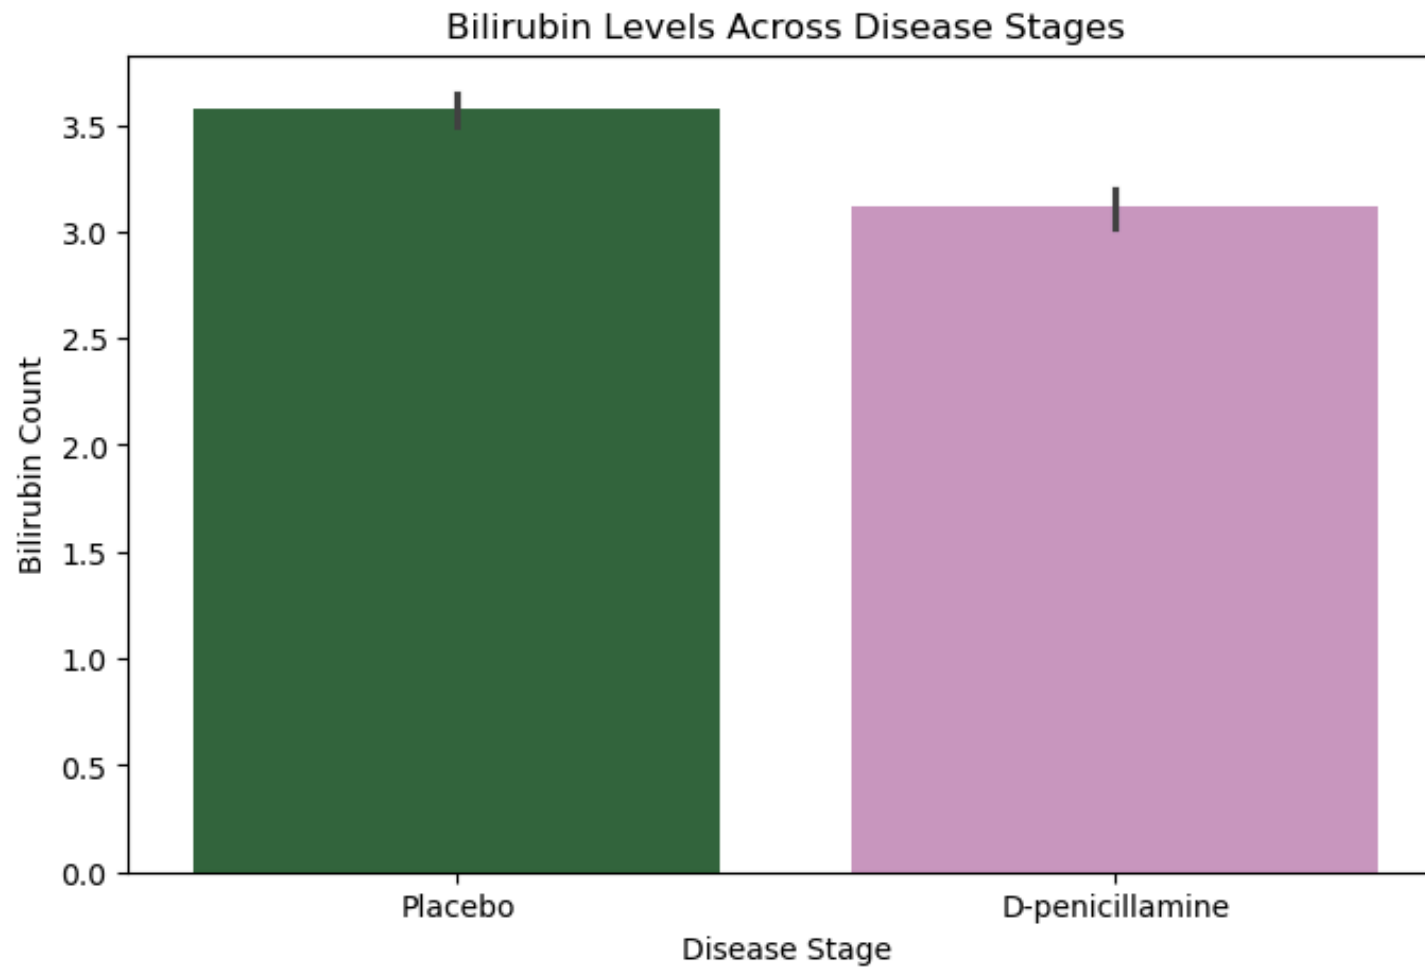
```
In [32]: plt.figure(figsize=(8, 5))
sns.violinplot(x='Drug', y='Platelets', data=df, palette="muted")
plt.title("Platelets Distribution by Drug")
plt.xlabel("Drug")
plt.ylabel("Platelets")
plt.show()
```



```
In [33]: plt.figure(figsize=(8, 5))
sns.barplot(x='Drug', y='Albumin', data=df, palette="cubehelix")
plt.title("Albumin Levels Across Disease Stages")
plt.xlabel("Disease Stage")
plt.ylabel("Albumin Count")
plt.show()
```

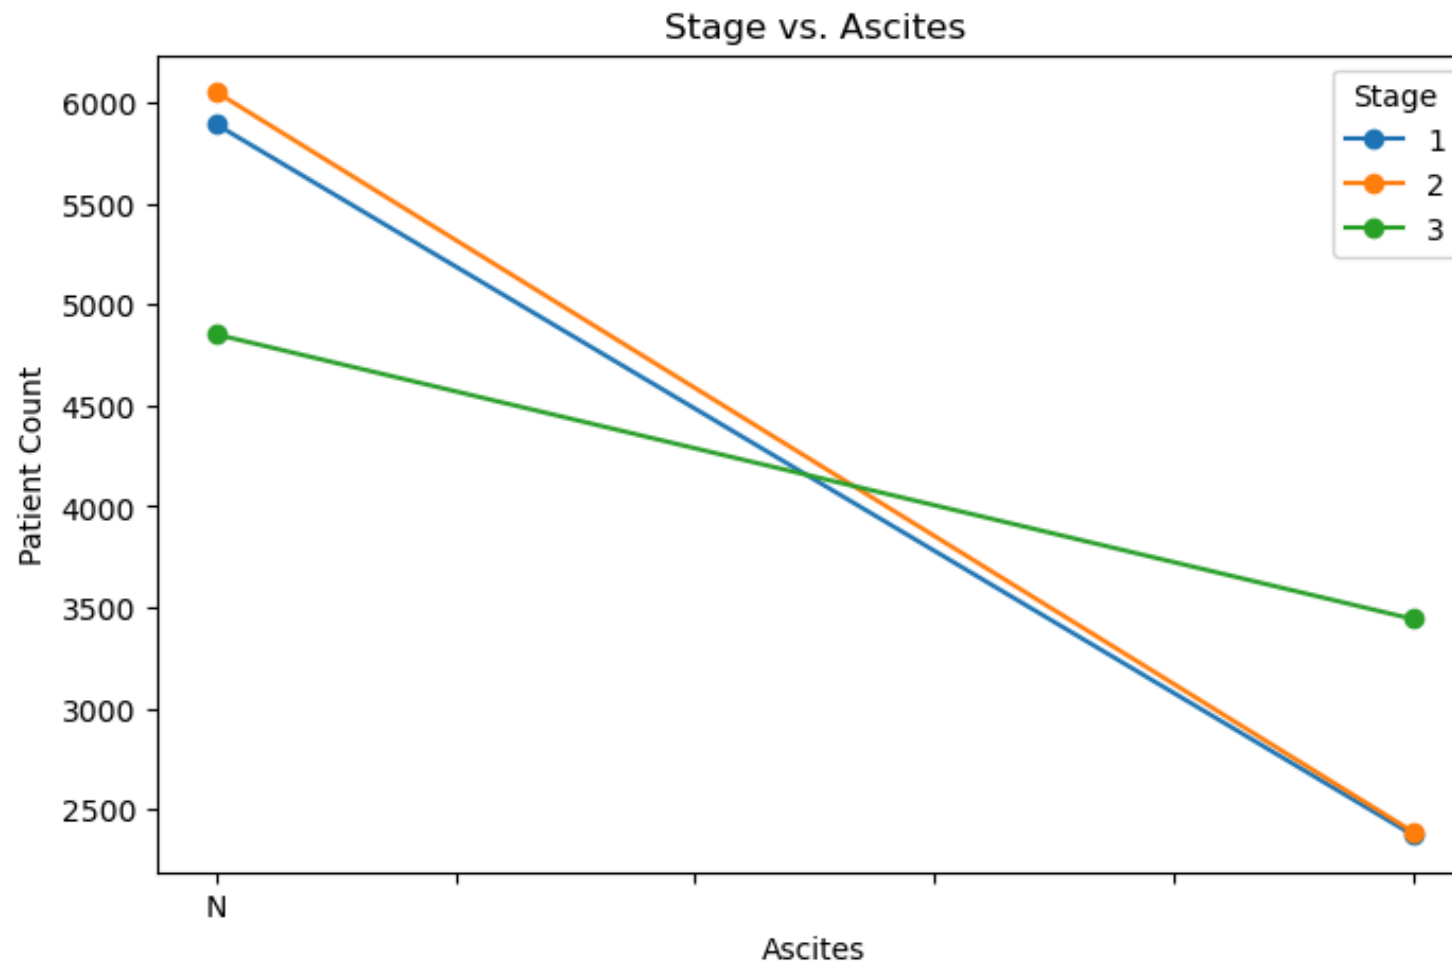


```
In [34]: plt.figure(figsize=(8, 5))
sns.barplot(x='Drug', y='Bilirubin', data=df, palette="cubehelix")
plt.title("Bilirubin Levels Across Disease Stages")
plt.xlabel("Disease Stage")
plt.ylabel("Bilirubin Count")
plt.show()
```



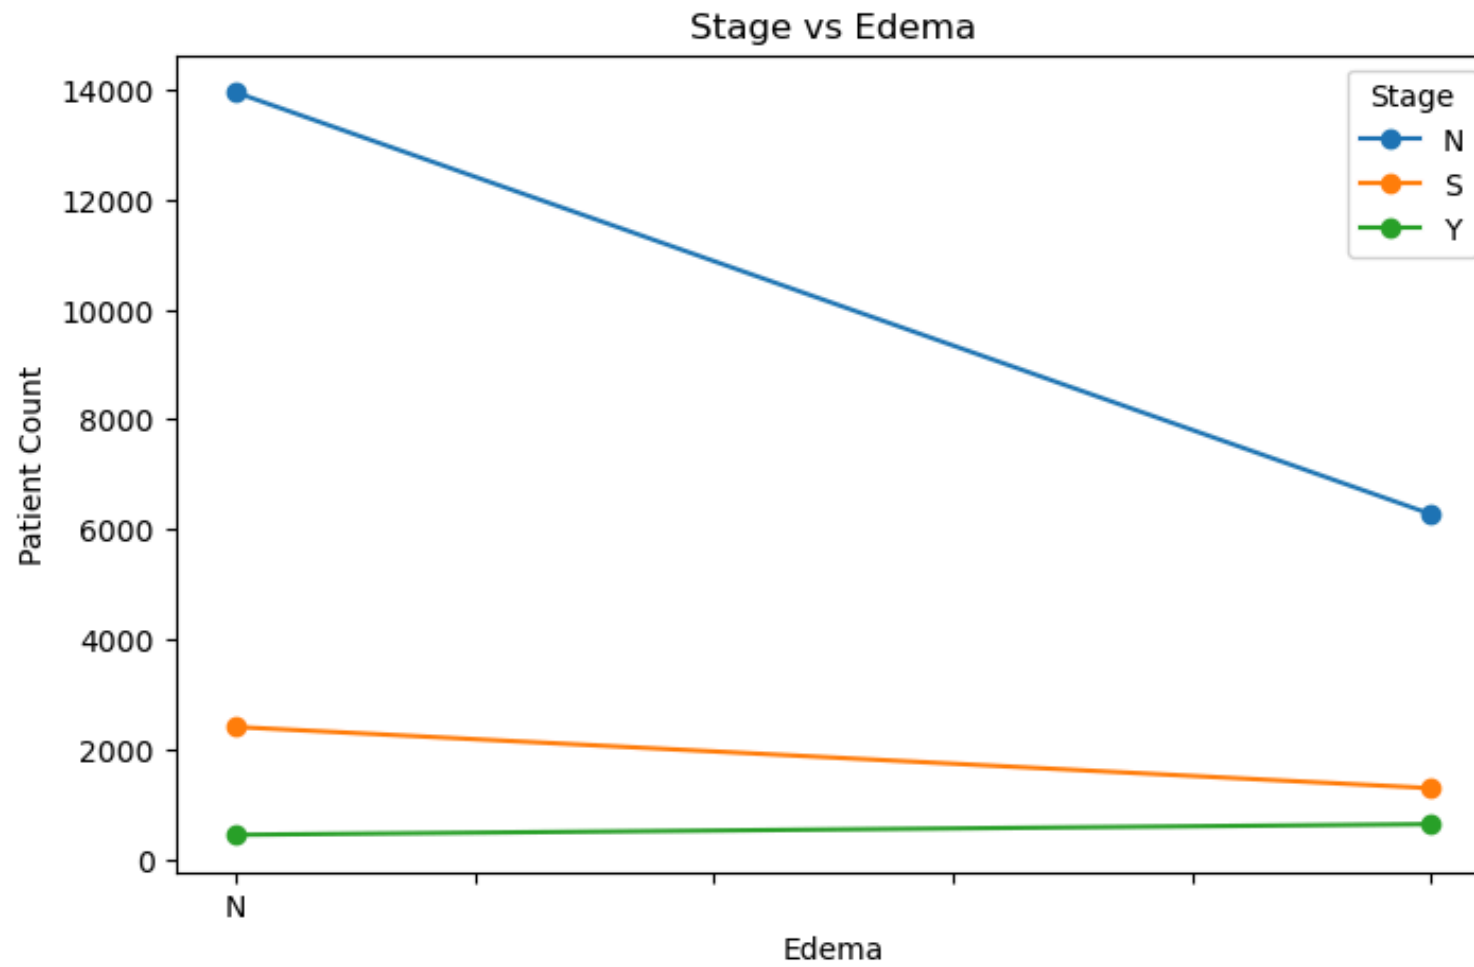
```
In [35]: plt.figure(figsize=(8, 5))
status_drug_counts = df.groupby(['Ascites', 'Stage']).size().unstack()
status_drug_counts.plot(kind="line", marker="o", figsize=(8, 5))
plt.title("Stage vs. Ascites")
plt.xlabel("Ascites")
plt.ylabel("Patient Count")
plt.legend(title="Stage")
plt.show()
```

<Figure size 800x500 with 0 Axes>



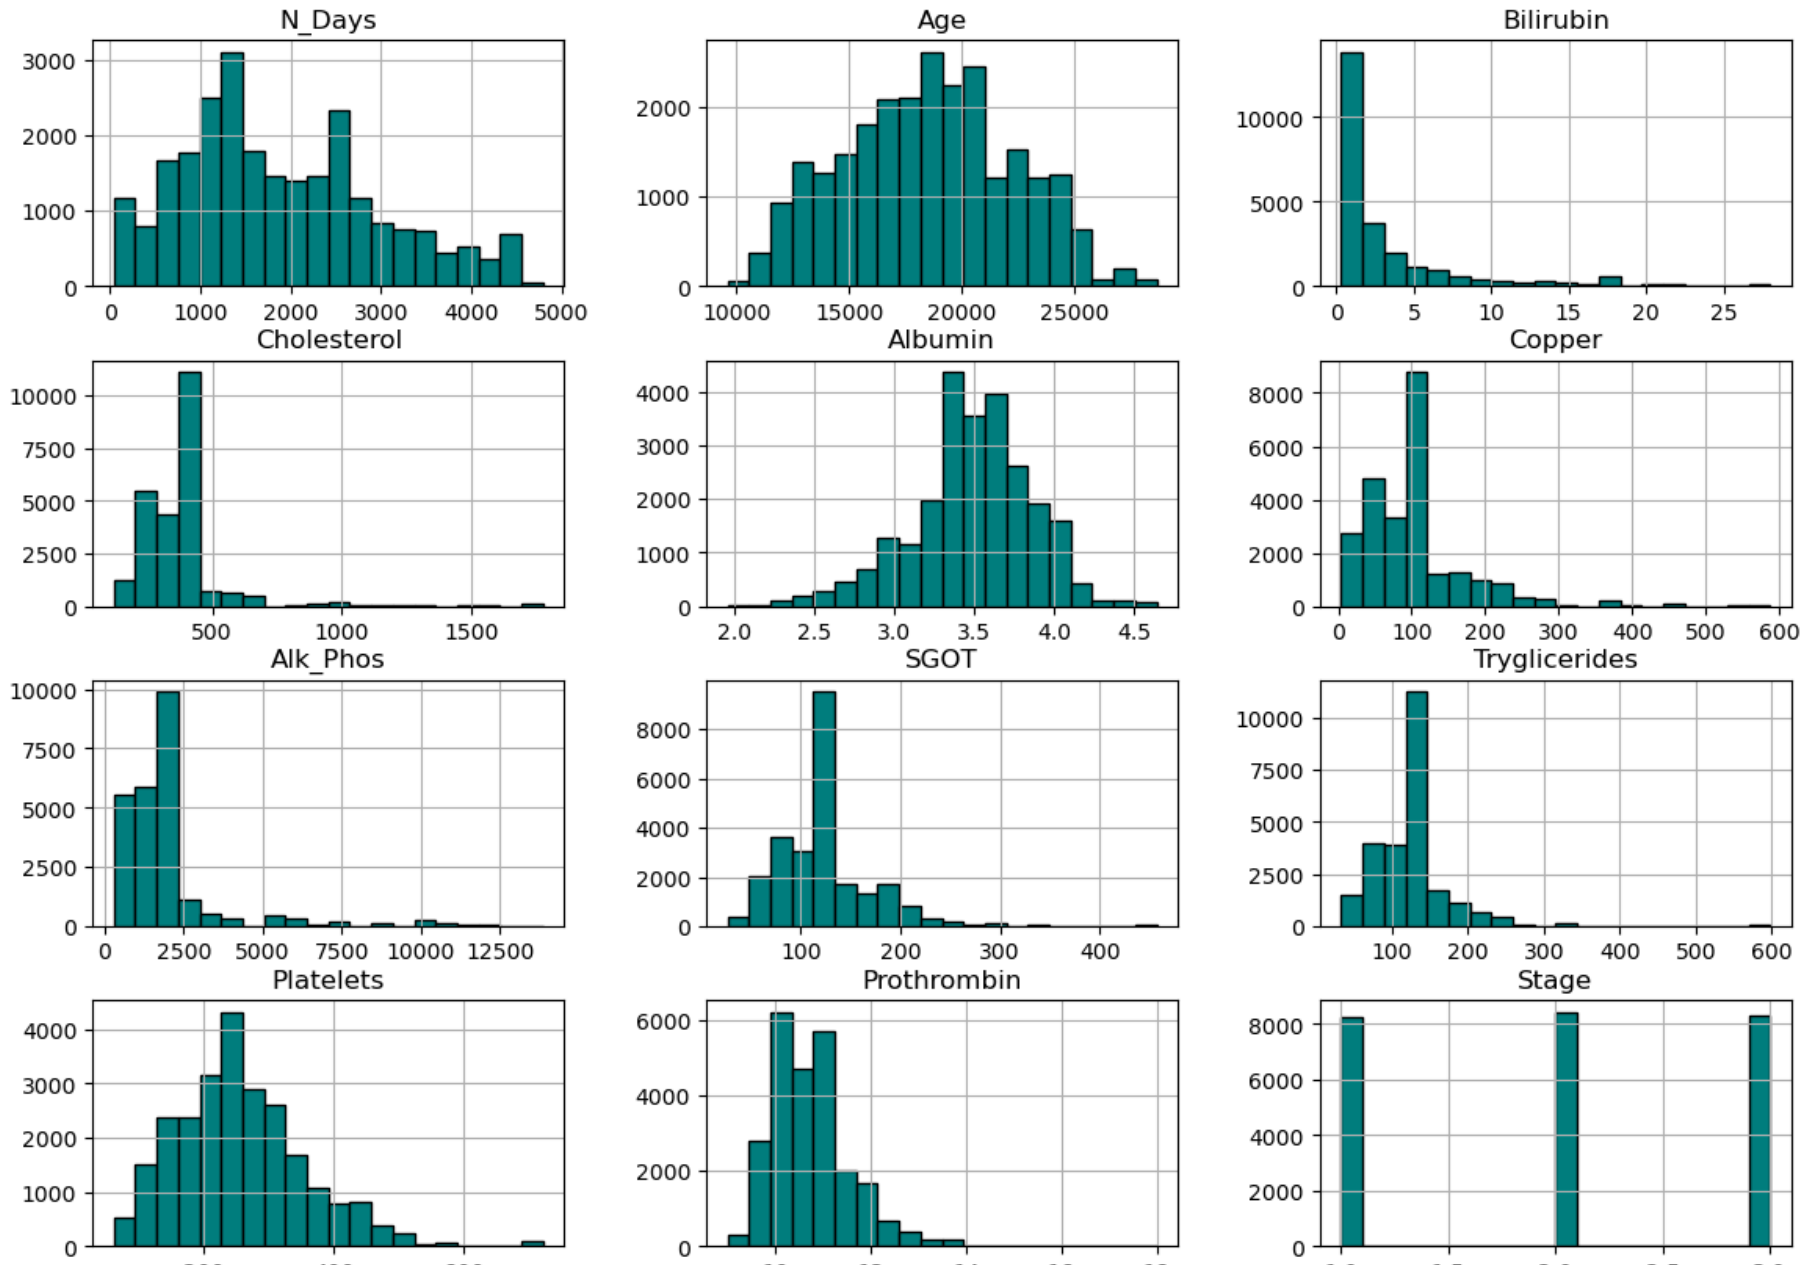
```
In [36]: plt.figure(figsize=(8, 5))
status_drug_counts = df.groupby(['Ascites', 'Edema']).size().unstack()
status_drug_counts.plot(kind="line", marker="o", figsize=(8, 5))
plt.title("Stage vs Edema")
plt.xlabel("Edema")
plt.ylabel("Patient Count")
plt.legend(title="Stage")
plt.show()
```

<Figure size 800x500 with 0 Axes>



```
In [37]: df.hist(figsize=(14, 10), bins=20, color="teal", edgecolor="black")
plt.suptitle("Histograms of All Numerical Columns")
plt.show()
```

Histograms of All Numerical Columns



DATA PREPROCESSING

Feature Encoding in Machine Learning

Feature encoding is the process of converting categorical data into numerical form so that machine learning models can understand and process it. Since many models (like linear regression, SVM, and neural networks) work with numerical values, encoding is crucial when dealing with categorical variables.

```
In [93]: from sklearn.preprocessing import LabelEncoder

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [95]: Cat_Col = []
        Num_Col = []
```

```
In [97]: for col in df.columns:
        if df[col].dtype == 'object':
            Cat_Col.append(col)
        else:
            Num_Col.append(col)
```

```
In [103... Cat_Col
```

```
Out[103... ['Status', 'Drug', 'Sex', 'Ascites', 'Hepatomegaly', 'Spiders', 'Edema']
```

```
In [105... Num_Col
```

```
Out[105... ['N_Days',  
            'Age',  
            'Bilirubin',  
            'Cholesterol',  
            'Albumin',  
            'Copper',  
            'Alk_Phos',  
            'SGOT',  
            'Tryglicerides',  
            'Platelets',  
            'Prothrombin',  
            'Stage']
```

Label Encoding (LE)

Label Encoding is a technique used to convert categorical values into numerical values. It assigns a unique integer (0, 1, 2, ...) to each category in a column.

Machine Learning Algorithms Need Numbers: Most models can't work with text data directly. They require numerical inputs.

Simple & Efficient: It replaces categories with numbers efficiently, taking up less memory than one-hot encoding.

Useful for Ordered Data: If your categorical values have a meaningful order (e.g., "Low", "Medium", "High"), then label encoding makes sense.

```
In [109... LE = LabelEncoder()
```

```
In [113... LE = {}
```

```
In [138... for col in Cat_Col:  
    encoder = LabelEncoder()  
    df[col] = encoder.fit_transform(df[col])
```

```
LE[f"{col}_Encoder"] = encoder  
df[col].value_counts()
```

In [140... LE

```
Out[140... {'Status_Encoder': LabelEncoder(),  
            'Drug_Encoder': LabelEncoder(),  
            'Sex_Encoder': LabelEncoder(),  
            'Ascites_Encoder': LabelEncoder(),  
            'Hepatomegaly_Encoder': LabelEncoder(),  
            'Spiders_Encoder': LabelEncoder(),  
            'Edema_Encoder': LabelEncoder()}
```

In [142... df.columns

```
Out[142... Index(['N_Days', 'Status', 'Drug', 'Age', 'Sex', 'Ascites', 'Hepatomegaly',  
        'Spiders', 'Edema', 'Bilirubin', 'Cholesterol', 'Albumin', 'Copper',  
        'Alk_Phos', 'SGOT', 'Tryglicerides', 'Platelets', 'Prothrombin',  
        'Stage'],  
      dtype='object')
```

In [144... df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25000 entries, 0 to 24999
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   N_Days                25000 non-null  int64
1   Status                25000 non-null  int64
2   Drug                 25000 non-null  int64
3   Age                  25000 non-null  int64
4   Sex                  25000 non-null  int64
5   Ascites              25000 non-null  int64
6   Hepatomegaly         25000 non-null  int64
7   Spiders              25000 non-null  int64
8   Edema                25000 non-null  int64
9   Bilirubin            25000 non-null  float64
10  Cholesterol           25000 non-null  float64
11  Albumin               25000 non-null  float64
12  Copper               25000 non-null  float64
13  Alk_Phos             25000 non-null  float64
14  SGOT                 25000 non-null  float64
15  Tryglicerides        25000 non-null  float64
16  Platelets            25000 non-null  float64
17  Prothrombin          25000 non-null  float64
18  Stage                25000 non-null  int64
dtypes: float64(9), int64(10)
memory usage: 3.6 MB
```

In [146... `df.head()`

Out[146...

	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin	Cholesterol	Albumin	Copper	Alk_Phc
0	2221	0	1	18499	0	0	1	0	0	0.5	149.0	4.04	227.0	598
1	1230	0	1	19724	1	1	0	1	0	0.5	219.0	3.93	22.0	663
2	4184	0	1	11839	0	0	0	0	0	0.5	320.0	3.54	51.0	1243
3	2090	2	1	16467	0	0	0	0	0	0.7	255.0	3.74	23.0	1024
4	2105	2	1	21699	0	0	1	0	0	1.9	486.0	3.54	74.0	1052

In [150...

```
df.describe()
```

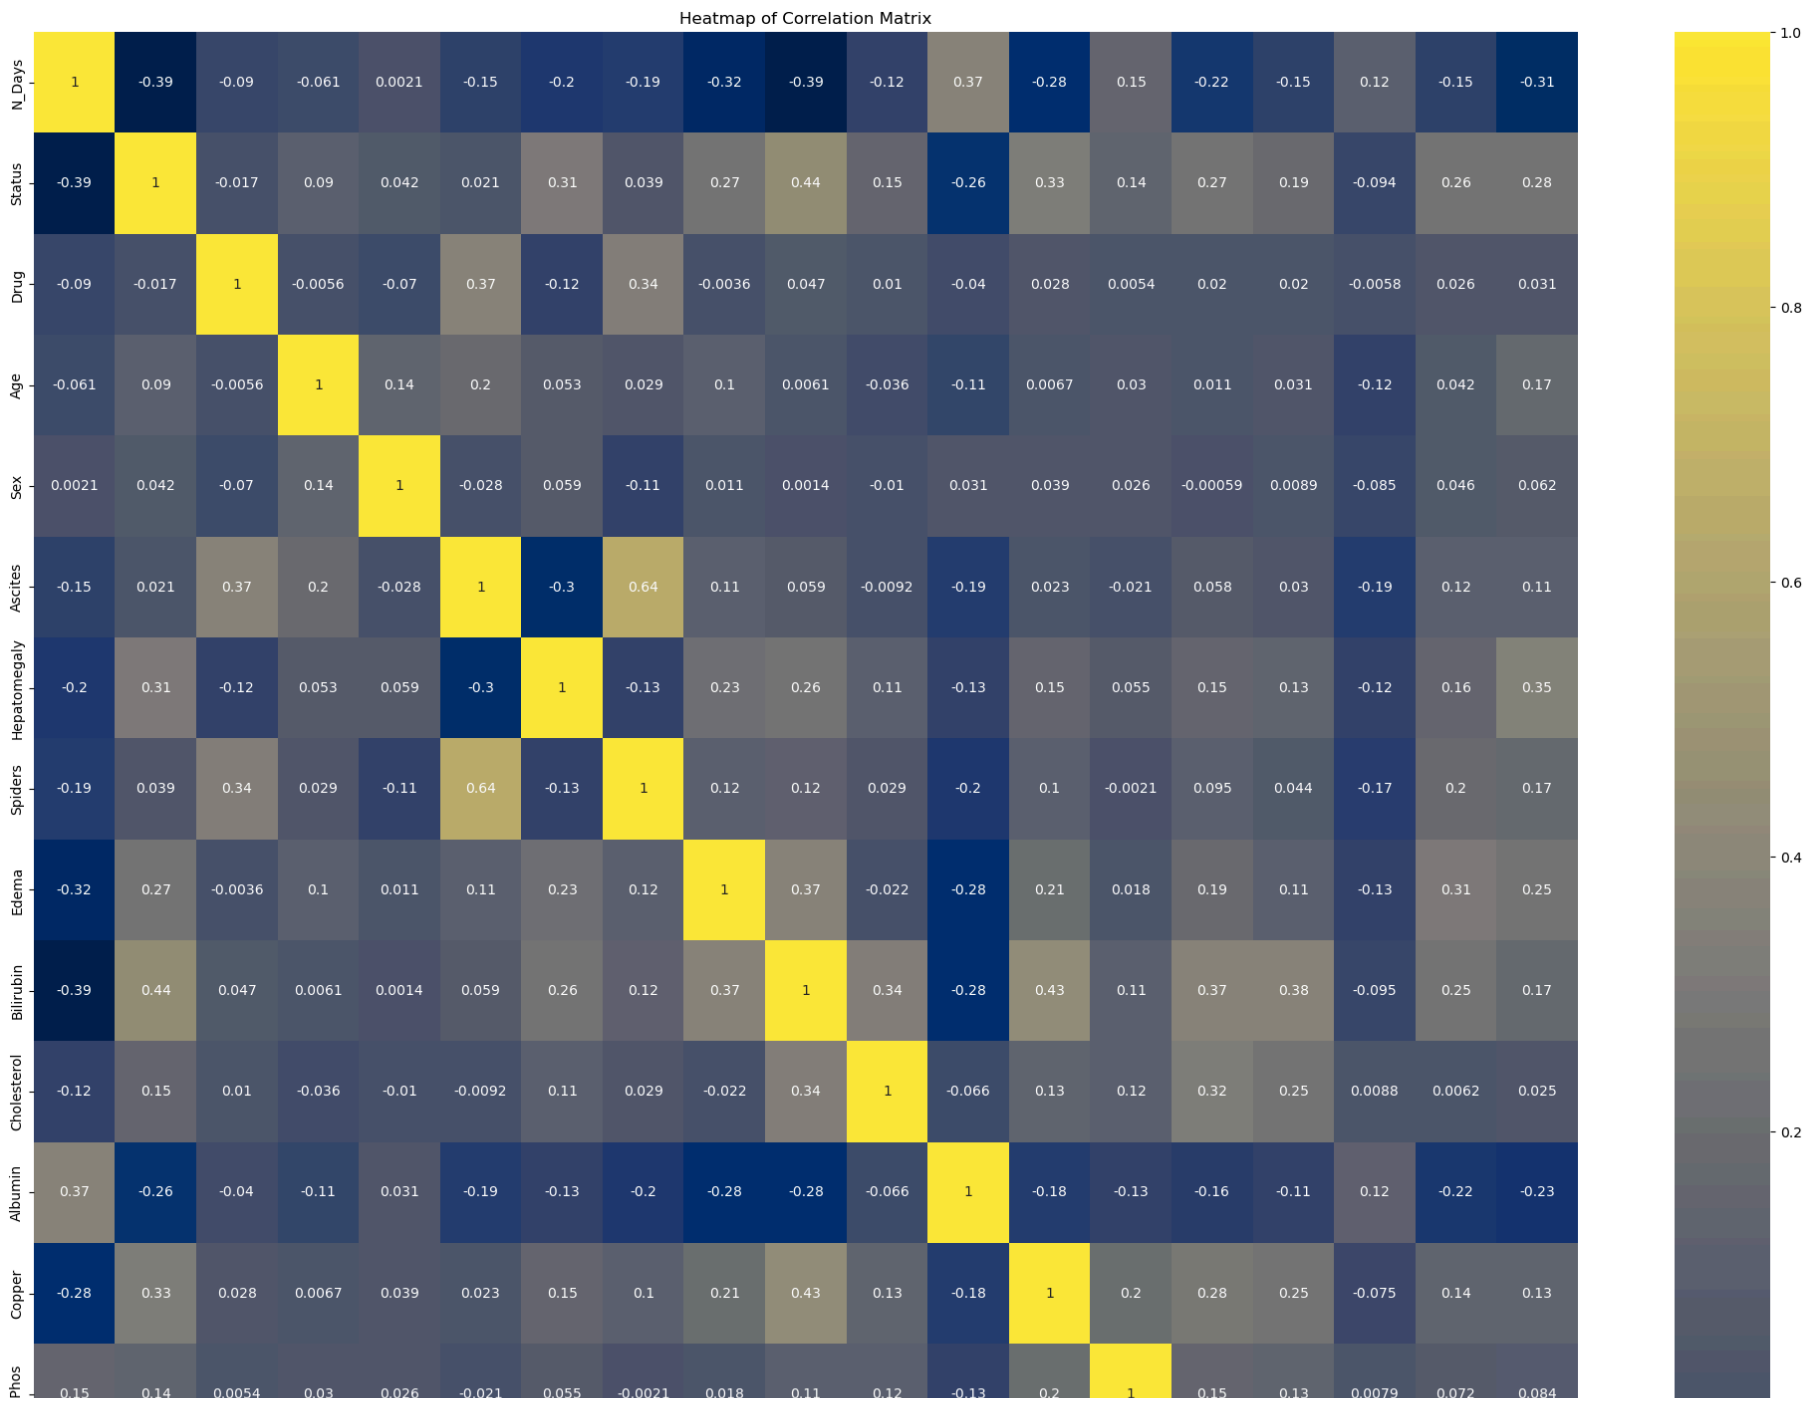
Out[150...

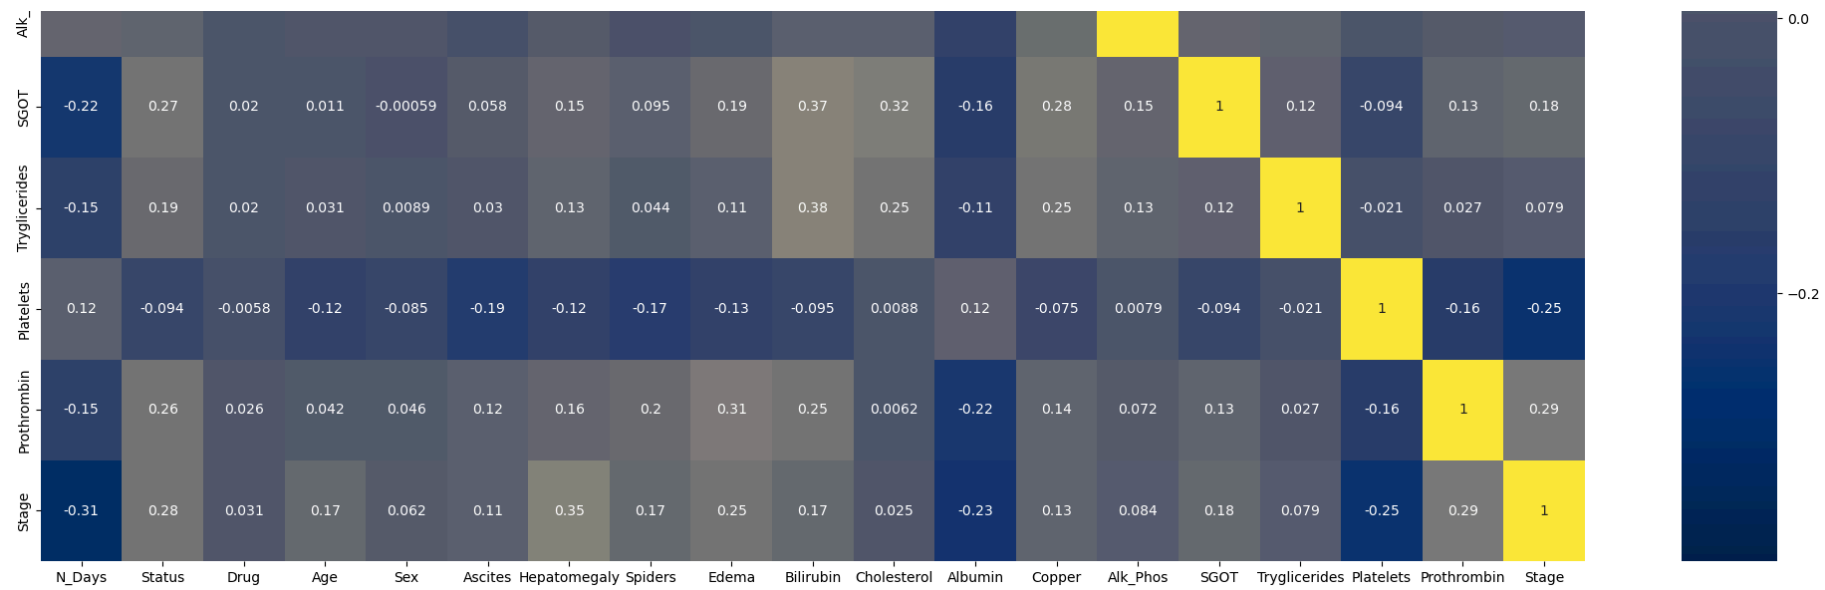
	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Ed
count	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000
mean	1887.117040	0.837600	0.633080	18495.877080	0.114520	0.328080	0.390280	0.45544	0.230000
std	1091.690918	0.944744	0.481974	3737.596616	0.318448	0.469524	0.487823	0.49802	0.500000
min	41.000000	0.000000	0.000000	9598.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1080.000000	0.000000	0.000000	15694.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	1680.000000	0.000000	1.000000	18499.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	2576.000000	2.000000	1.000000	20955.000000	0.000000	1.000000	1.000000	1.000000	0.000000
max	4795.000000	2.000000	1.000000	28650.000000	1.000000	1.000000	1.000000	1.000000	2.000000

In [196...

```
plt.figure(figsize=(25,25))
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='cividis')
```

```
plt.title('Heatmap of Correlation Matrix')  
plt.show()
```

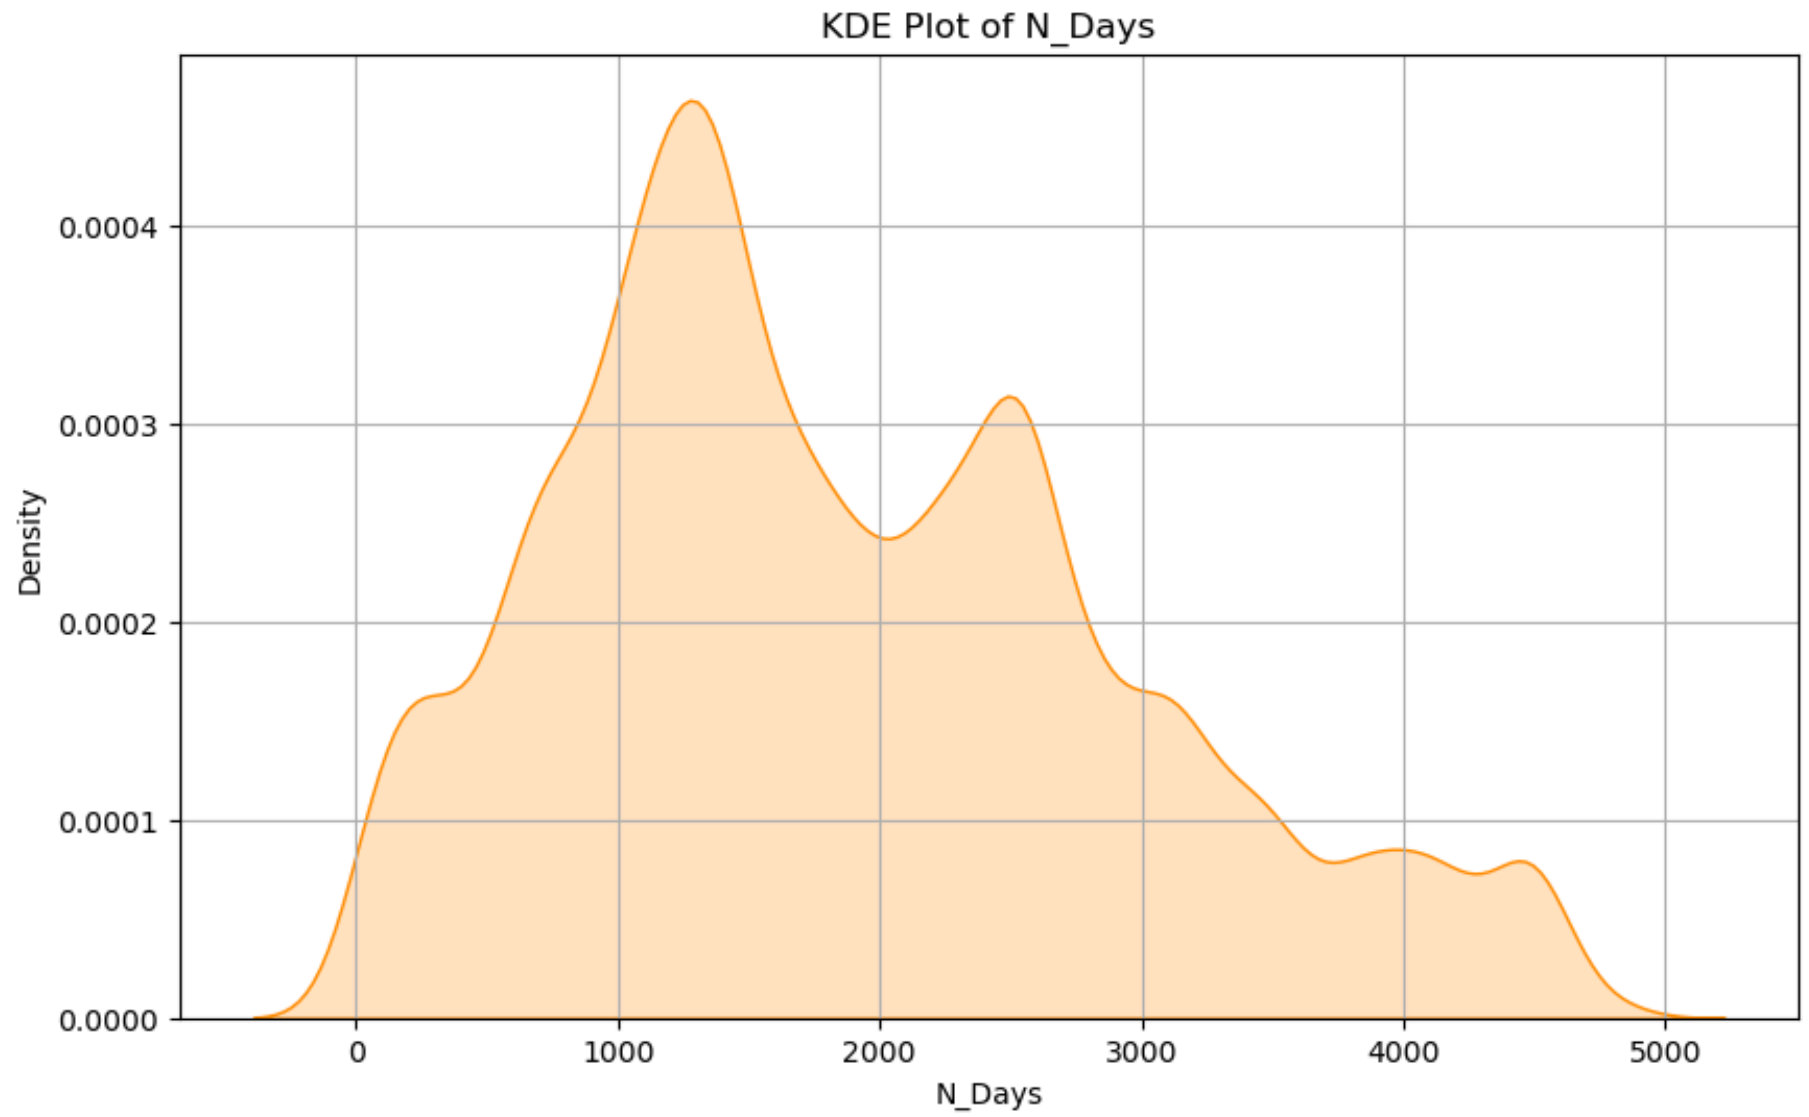


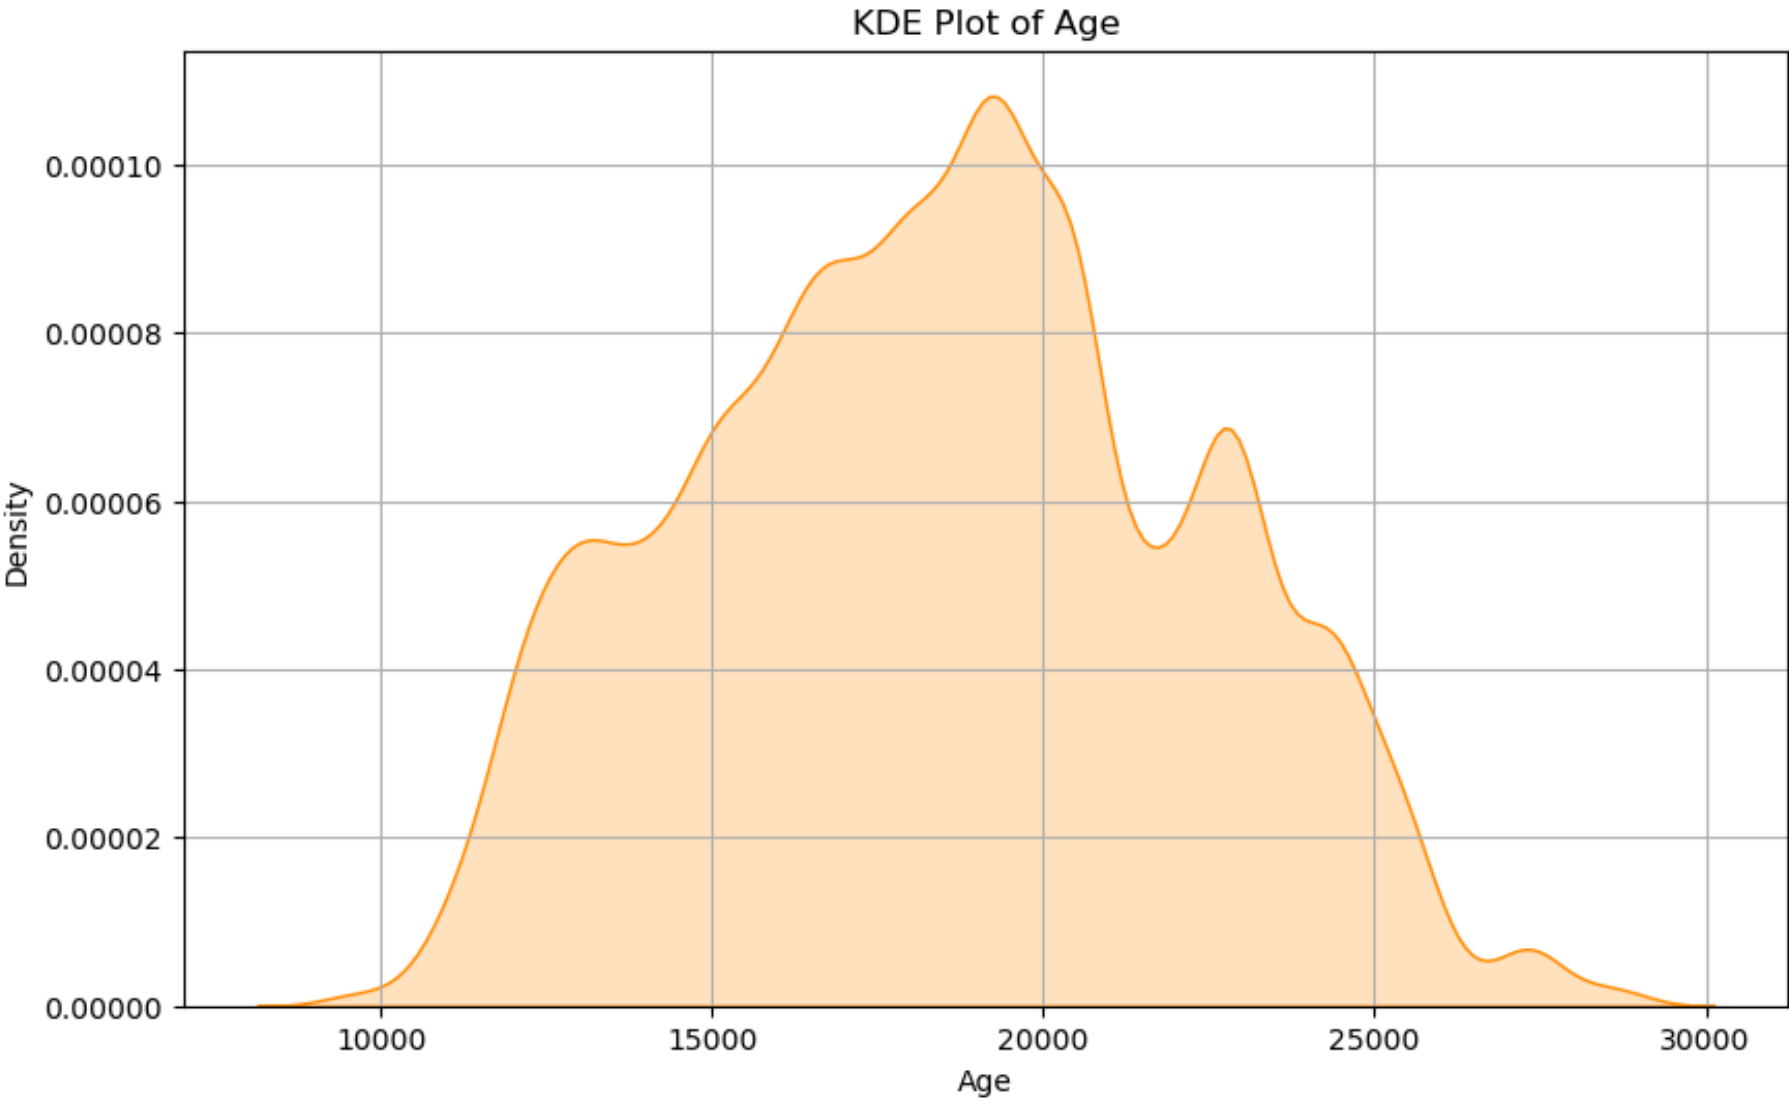
```
In [156... print(df.dtypes)
```

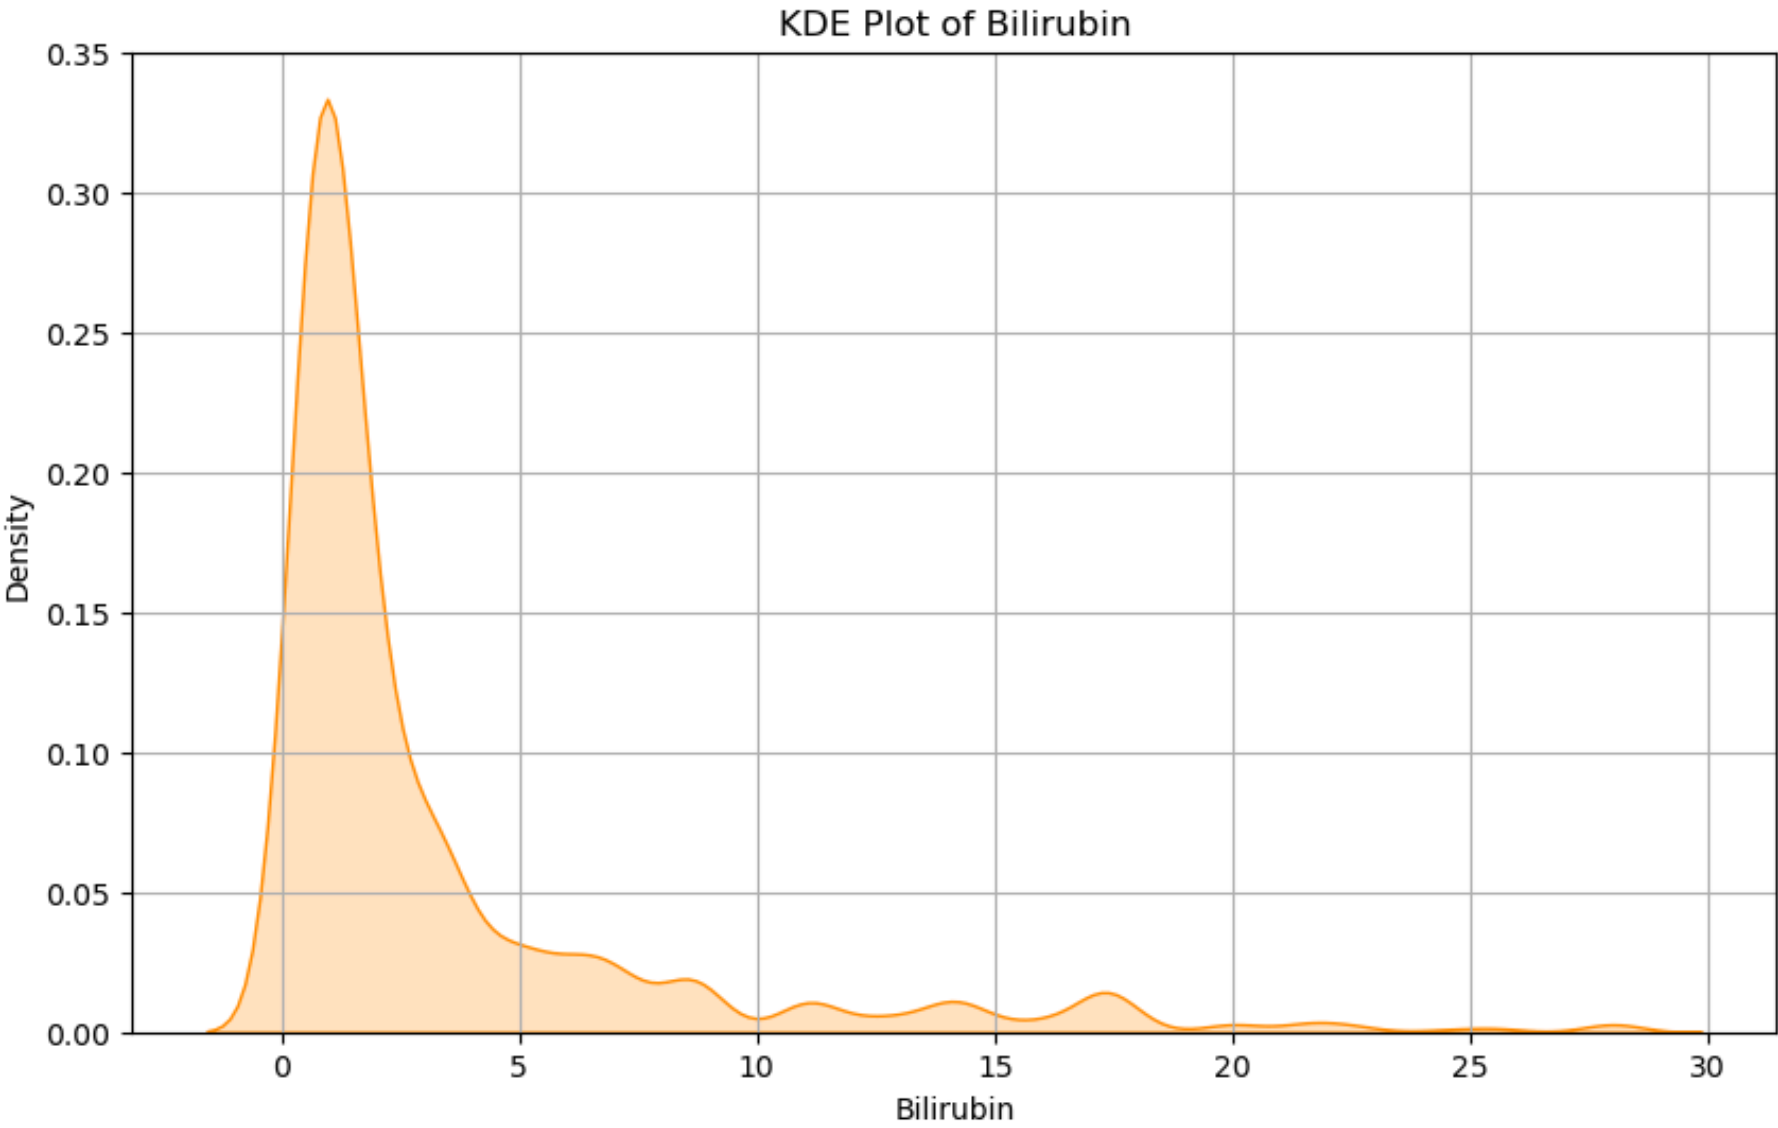
N_Days	int64
Status	int64
Drug	int64
Age	int64
Sex	int64
Ascites	int64
Hepatomegaly	int64
Spiders	int64
Edema	int64
Bilirubin	float64
Cholesterol	float64
Albumin	float64
Copper	float64
Alk_Phos	float64
SGOT	float64
Tryglicerides	float64
Platelets	float64
Prothrombin	float64
Stage	int64

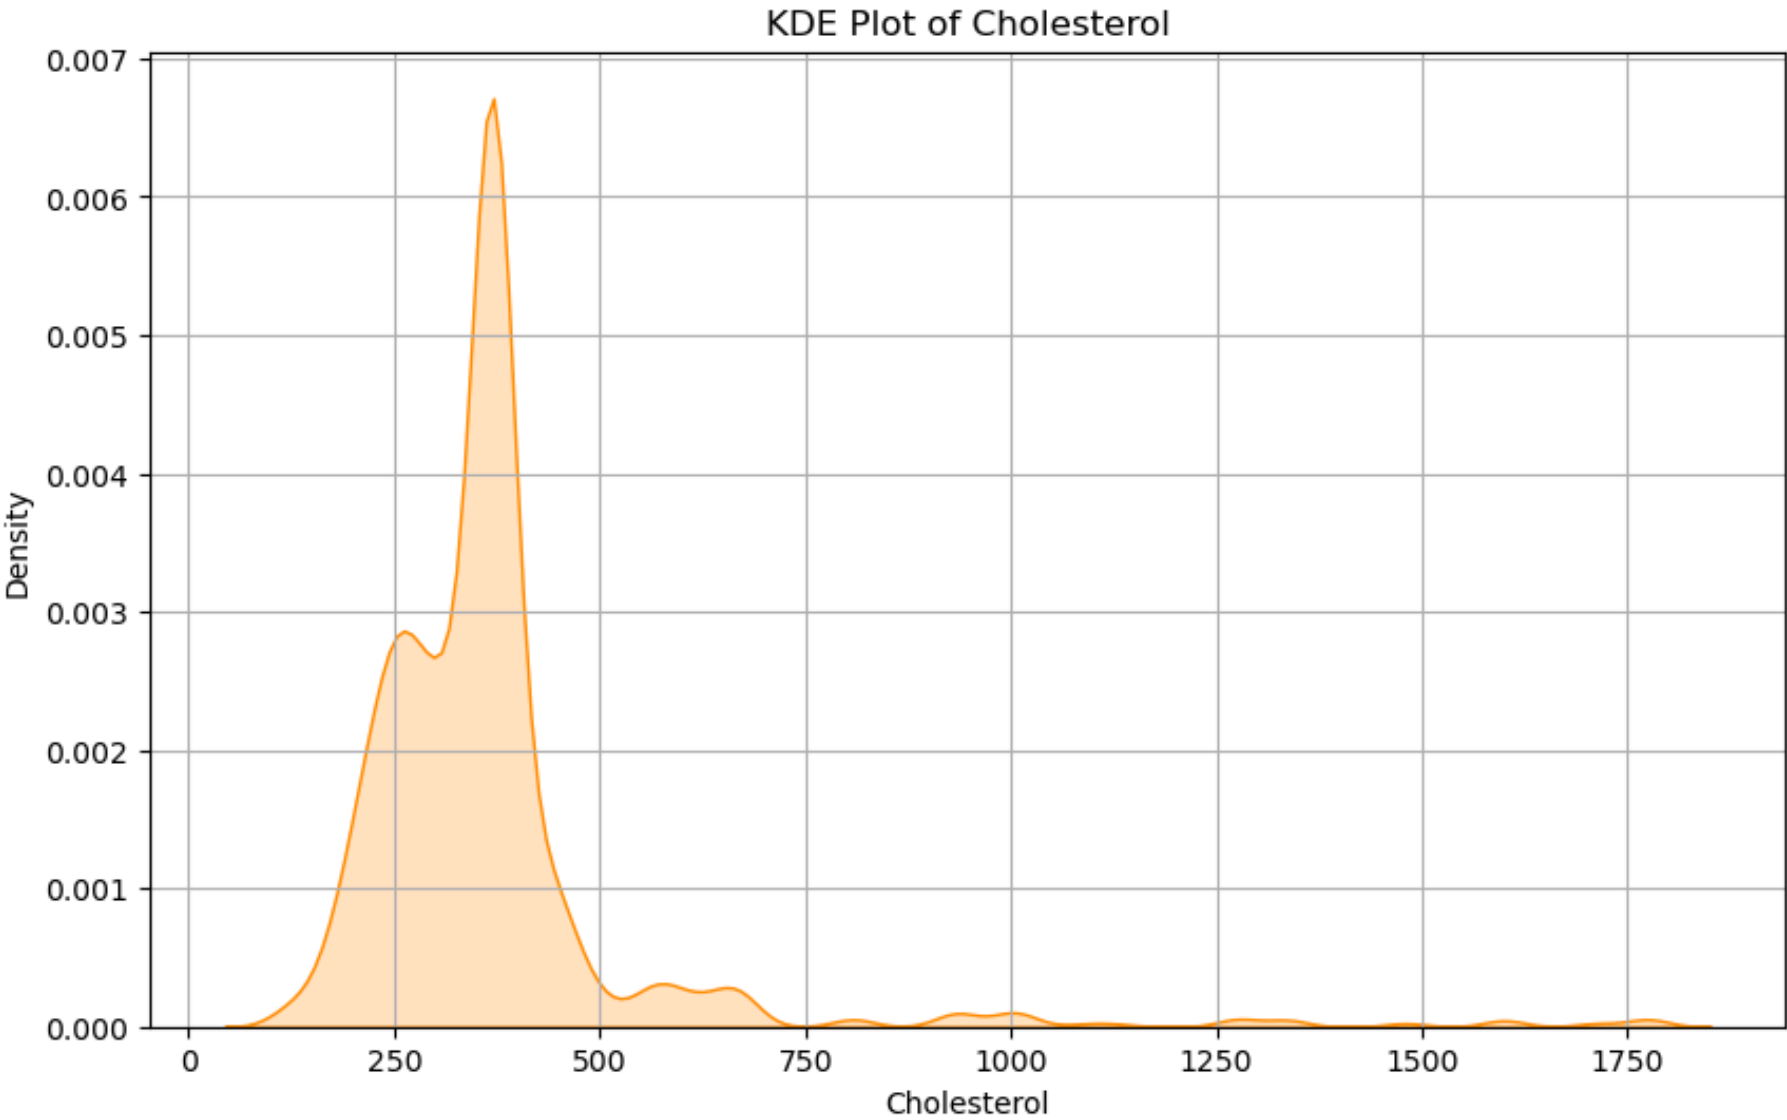
dtype: object

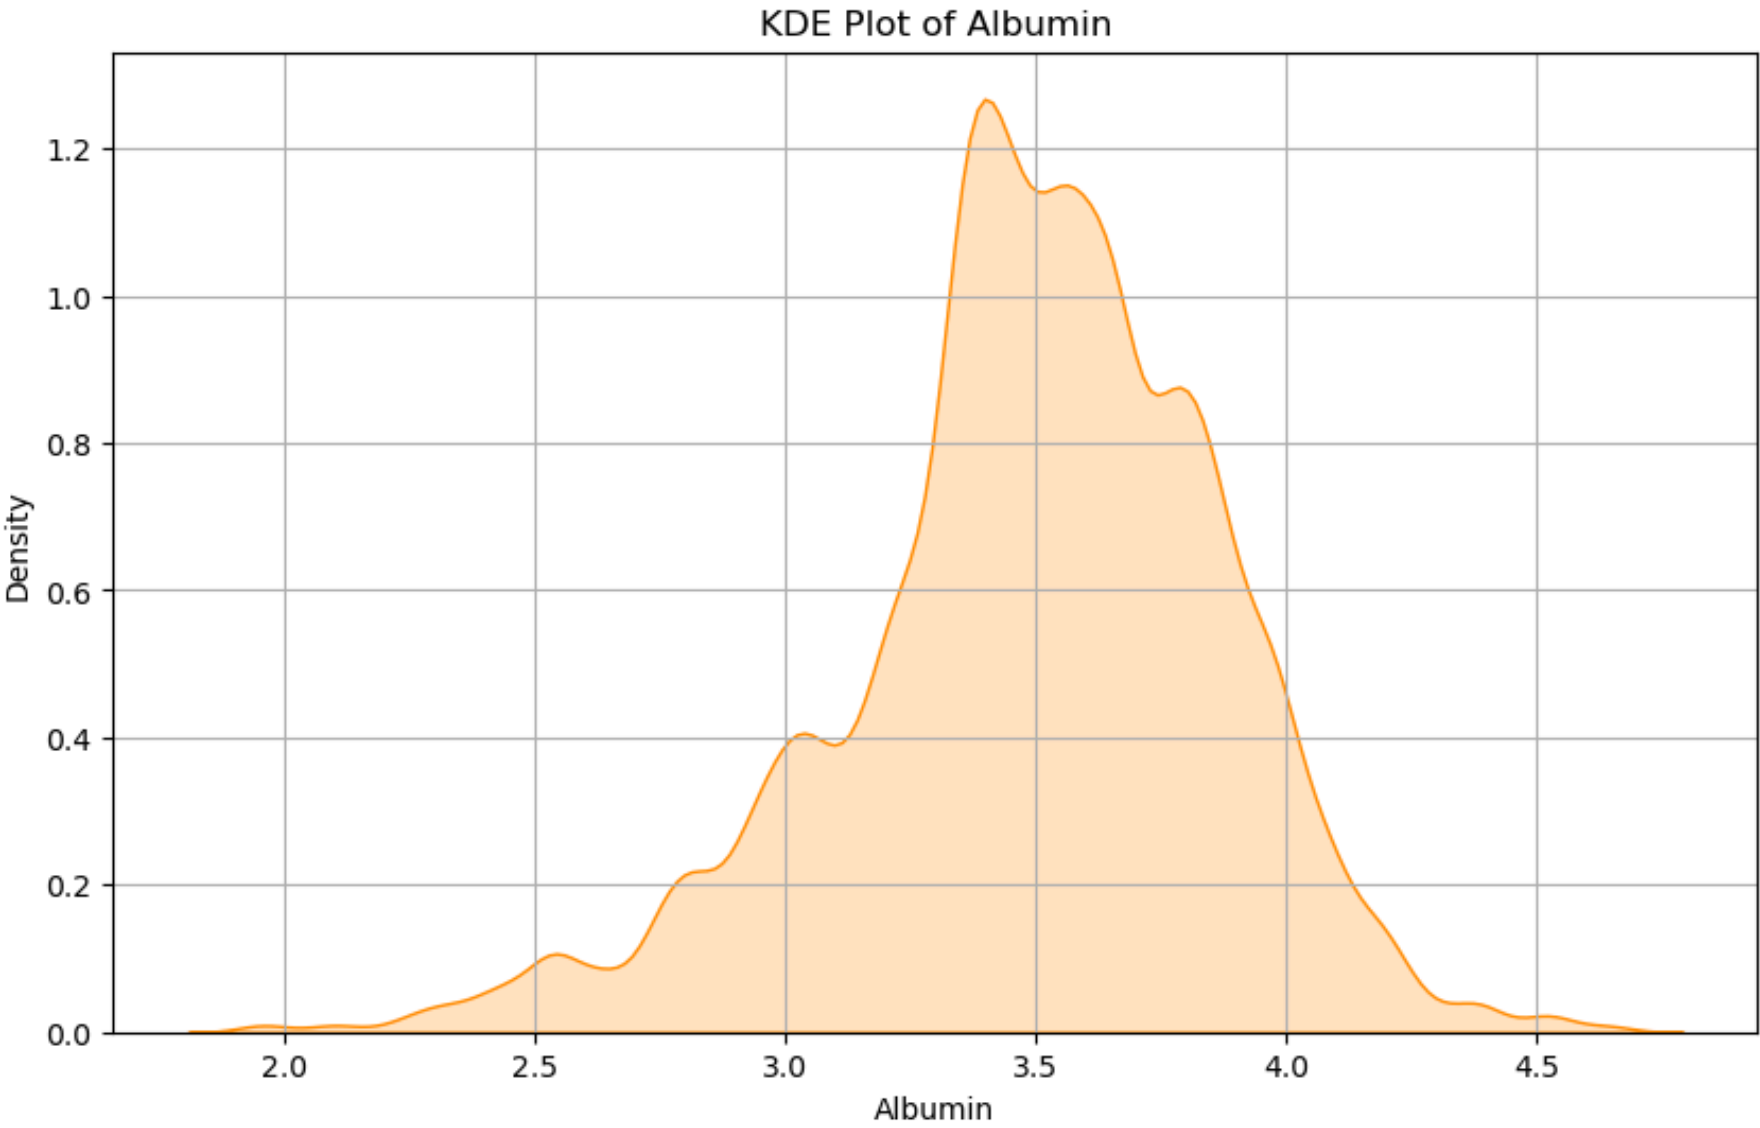
```
In [220... for col in Num_Col:
    plt.figure(figsize=(10, 6))
    sns.kdeplot(data=df, x=col, fill=True, color='darkorange', bw_adjust=1)
    plt.title(f'KDE Plot of {col}')
    plt.xlabel(col)
    plt.ylabel('Density')
    plt.grid(True)
    plt.show()
```

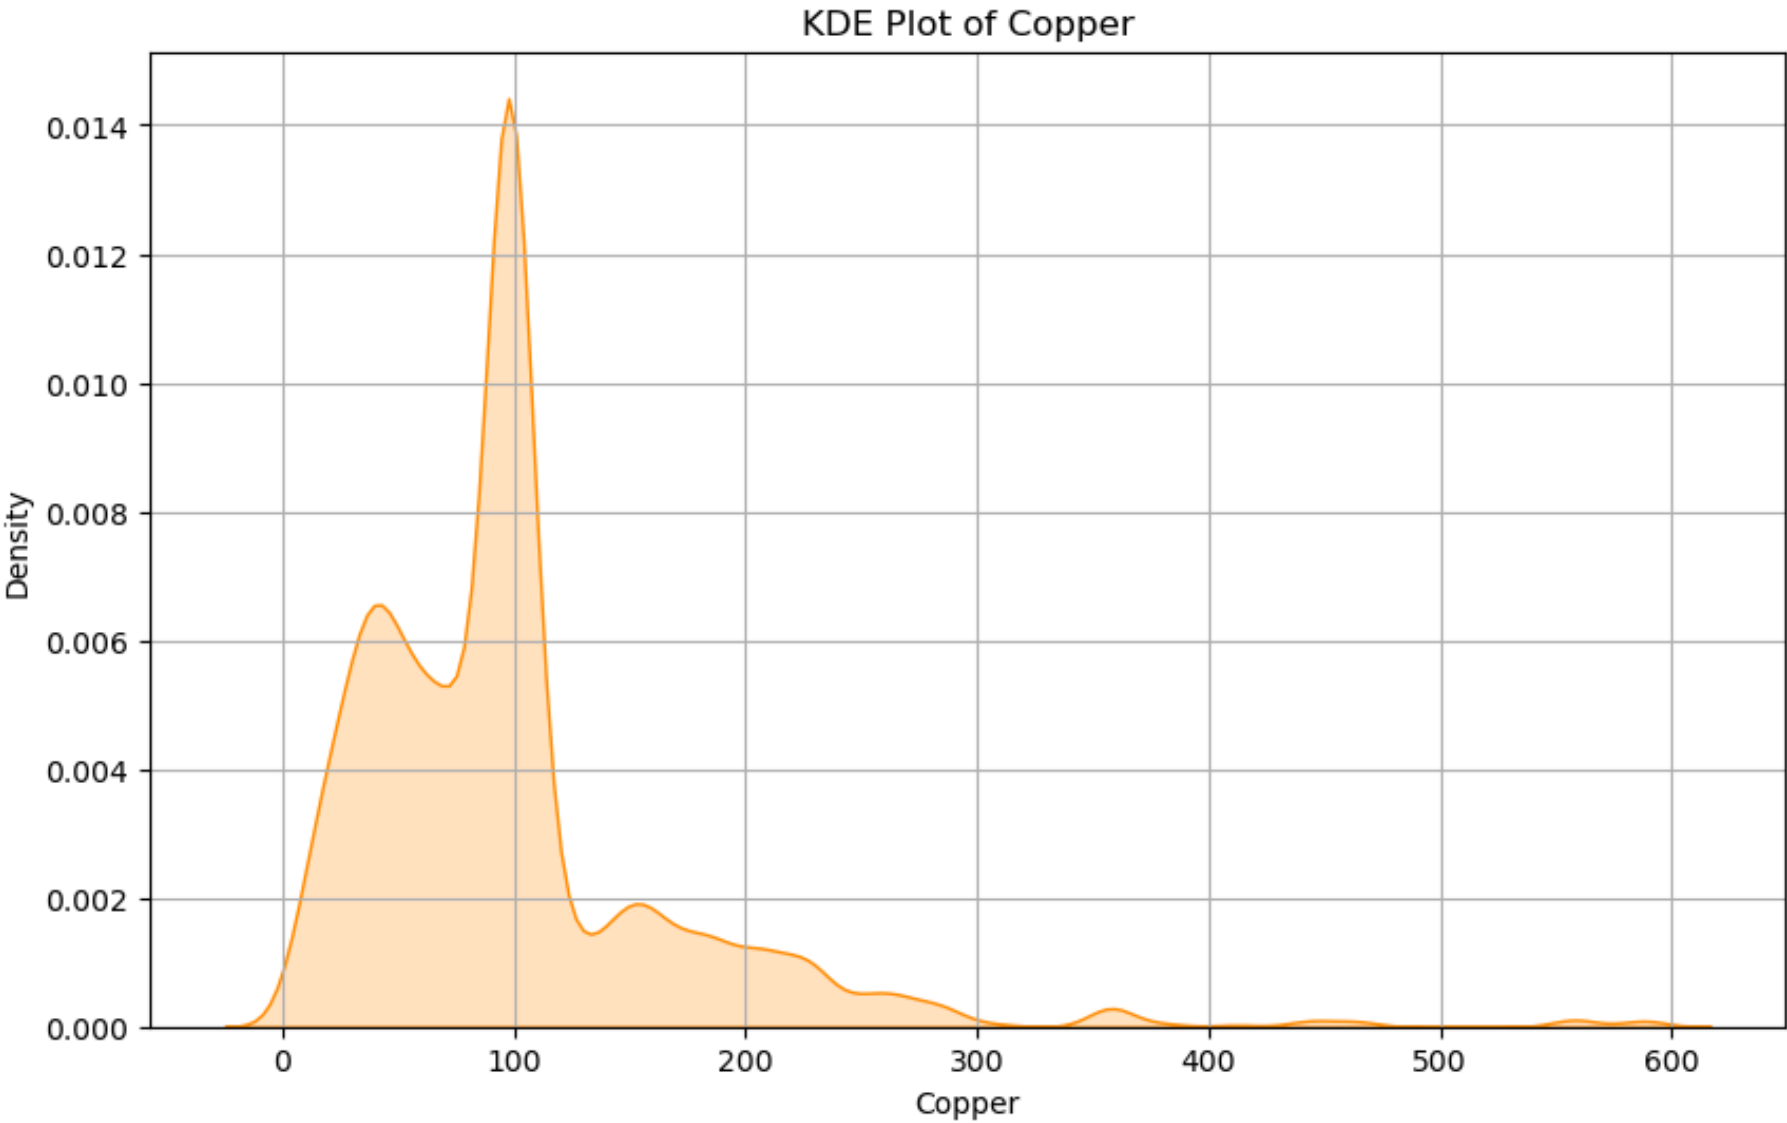


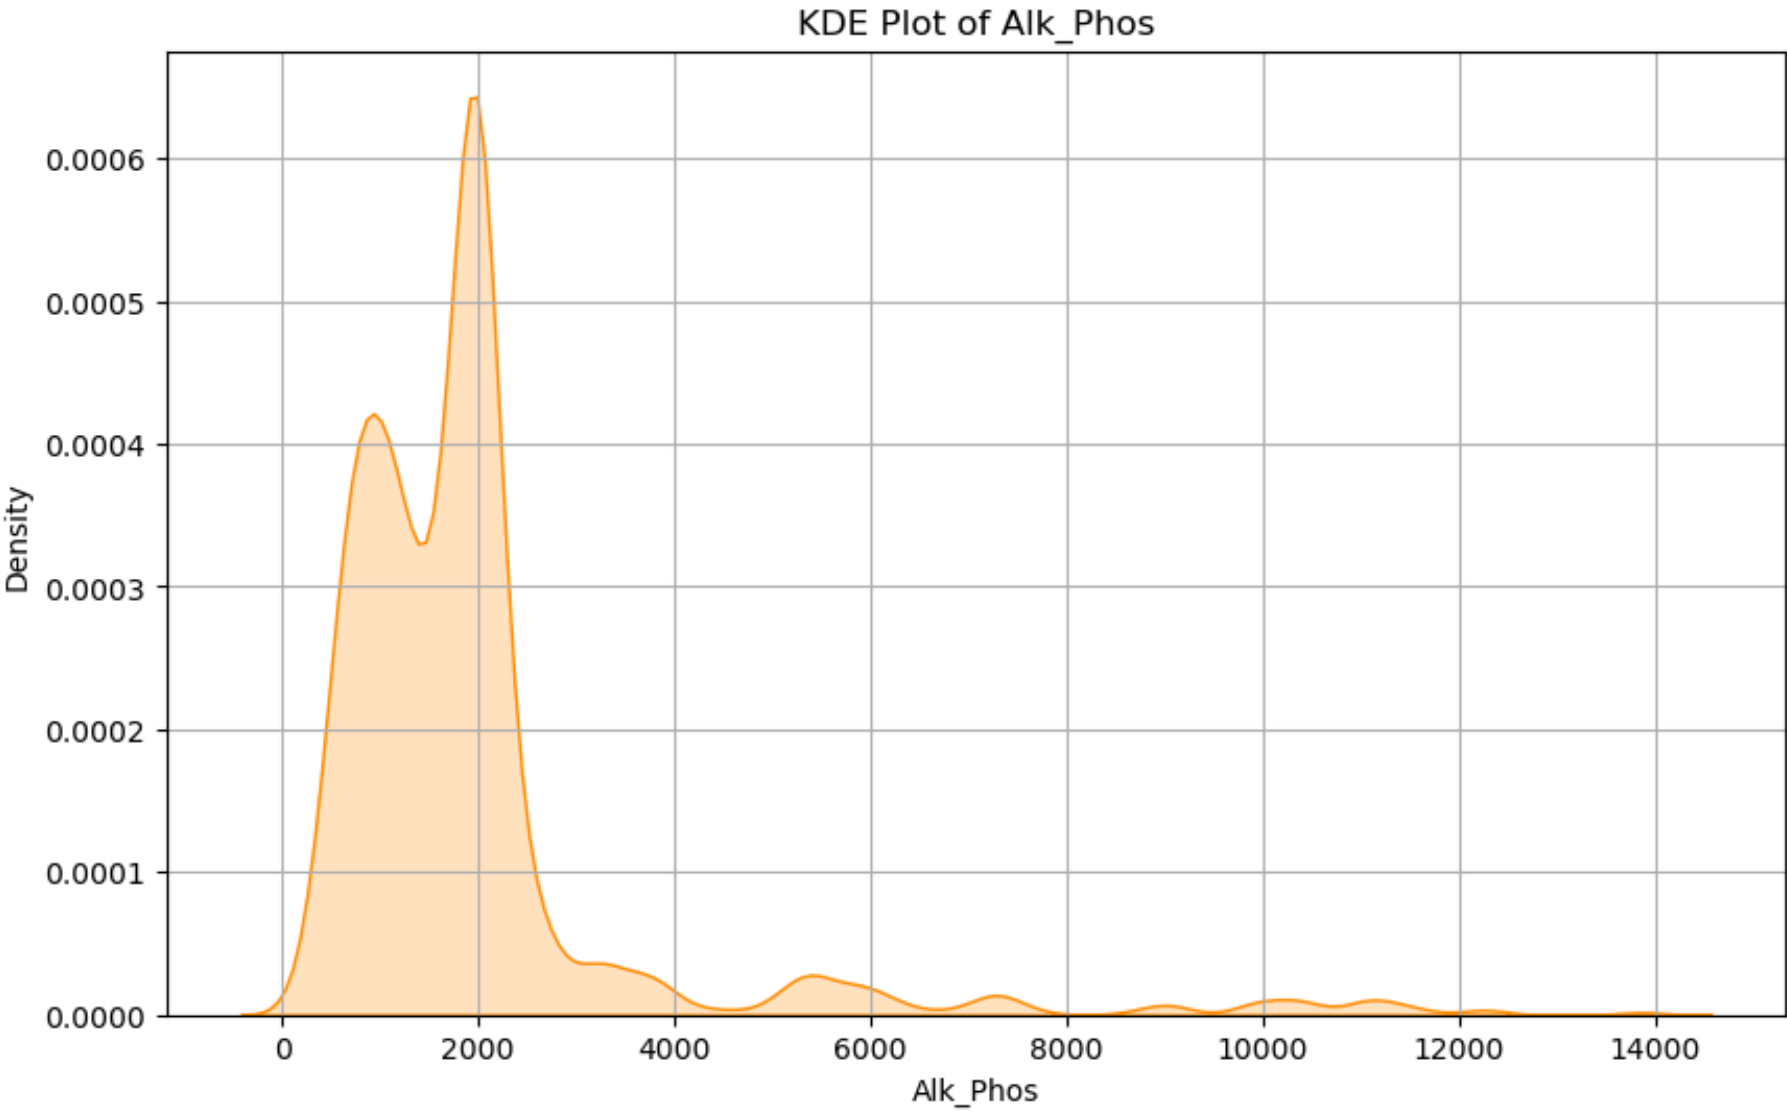


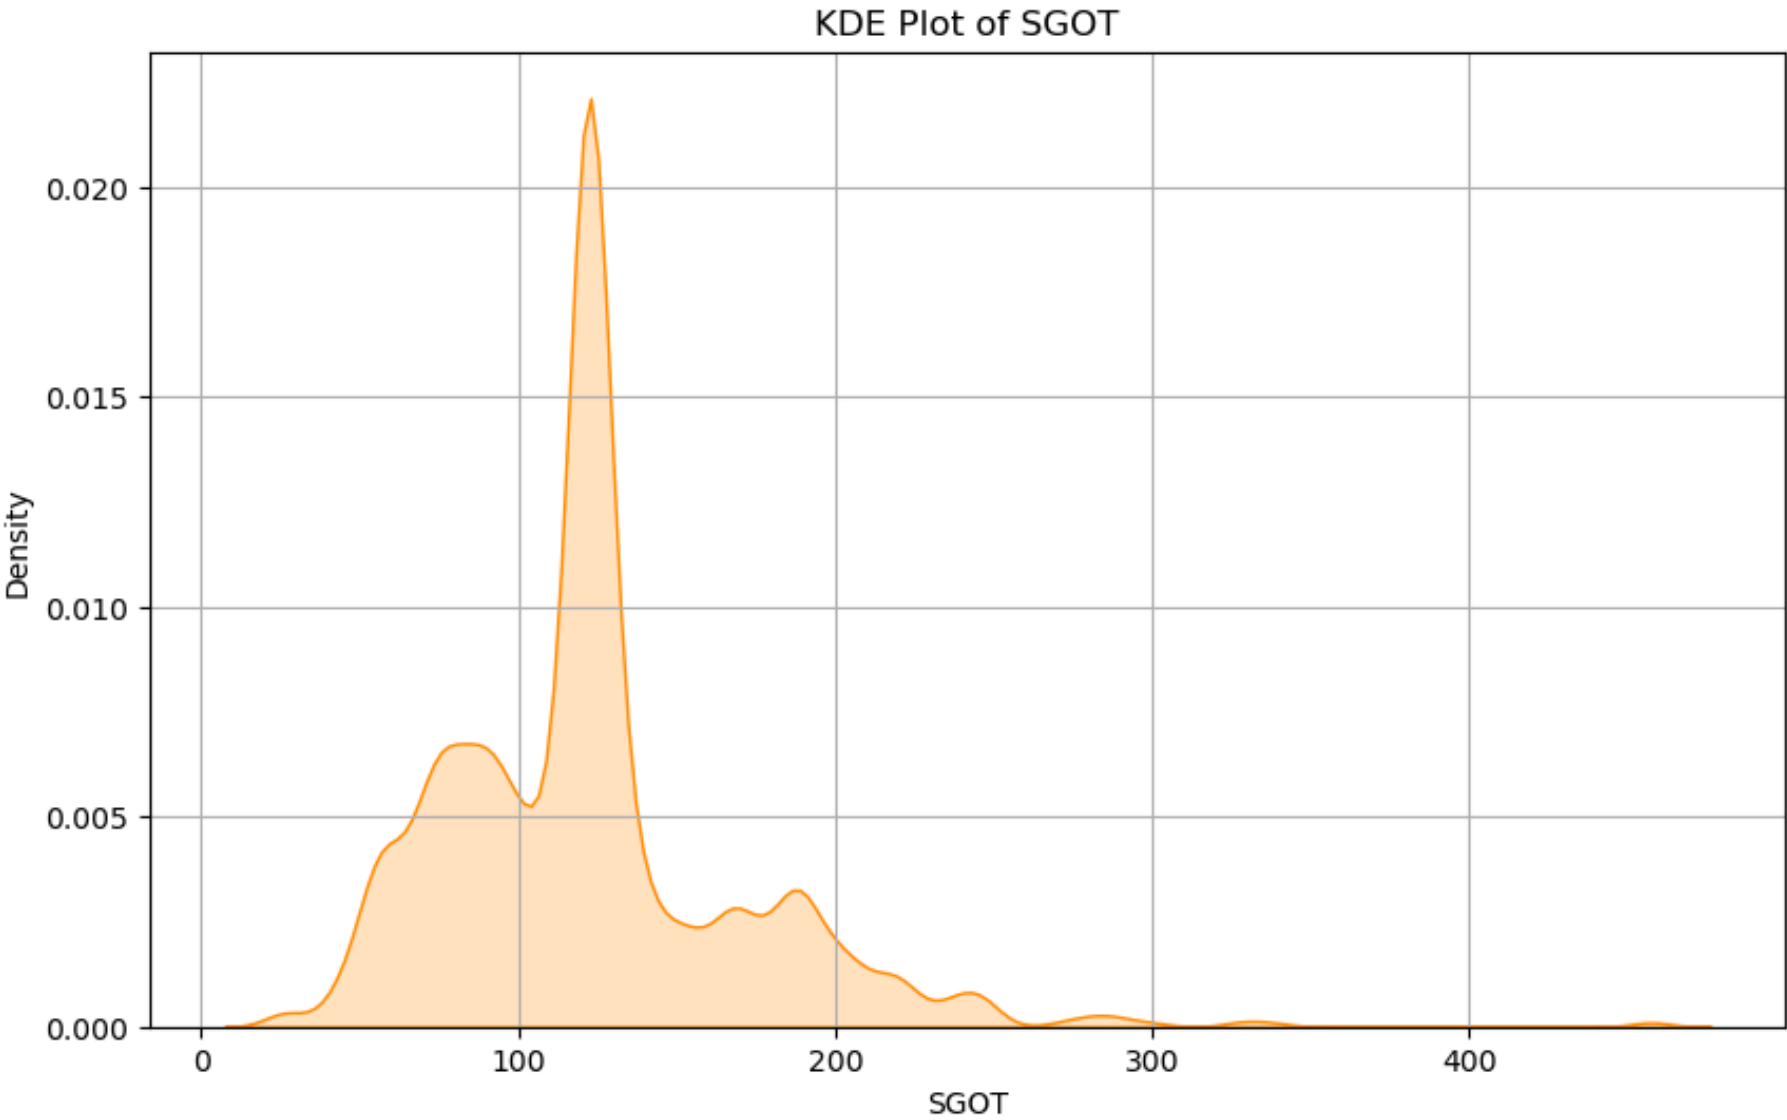


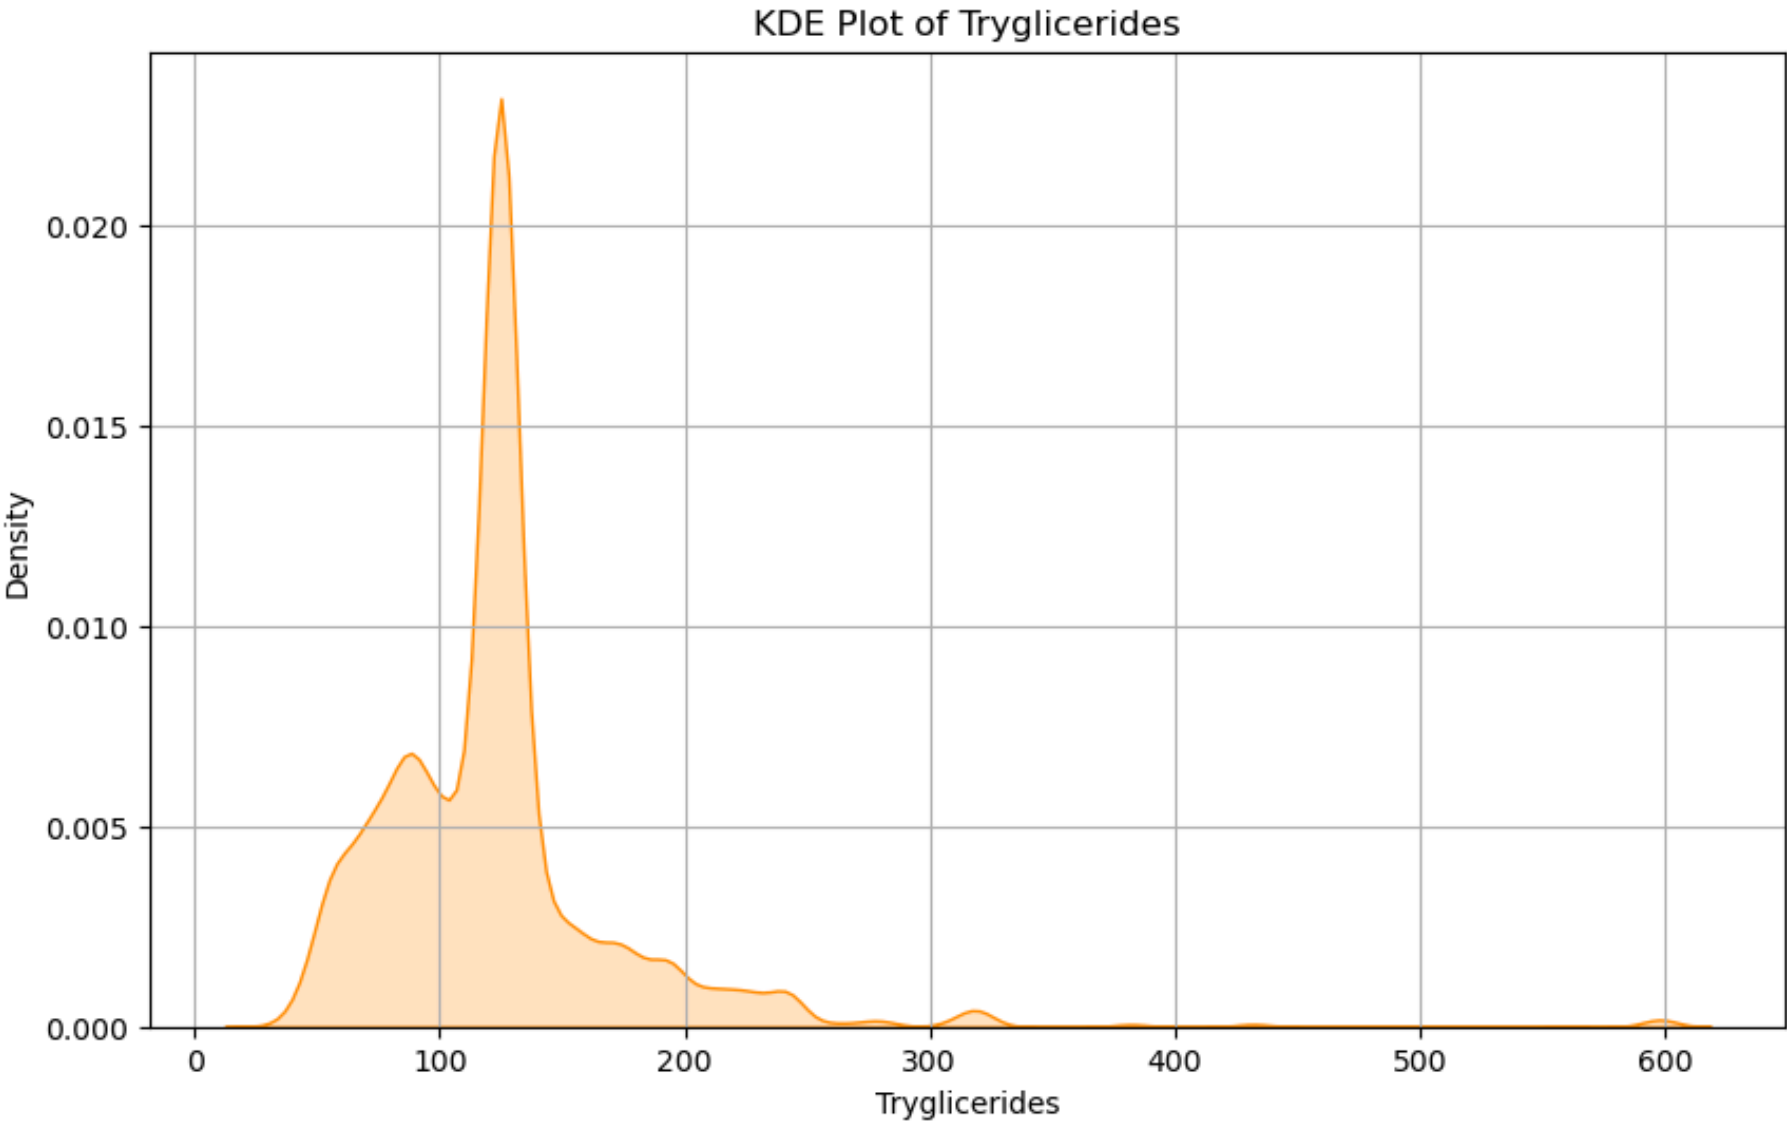


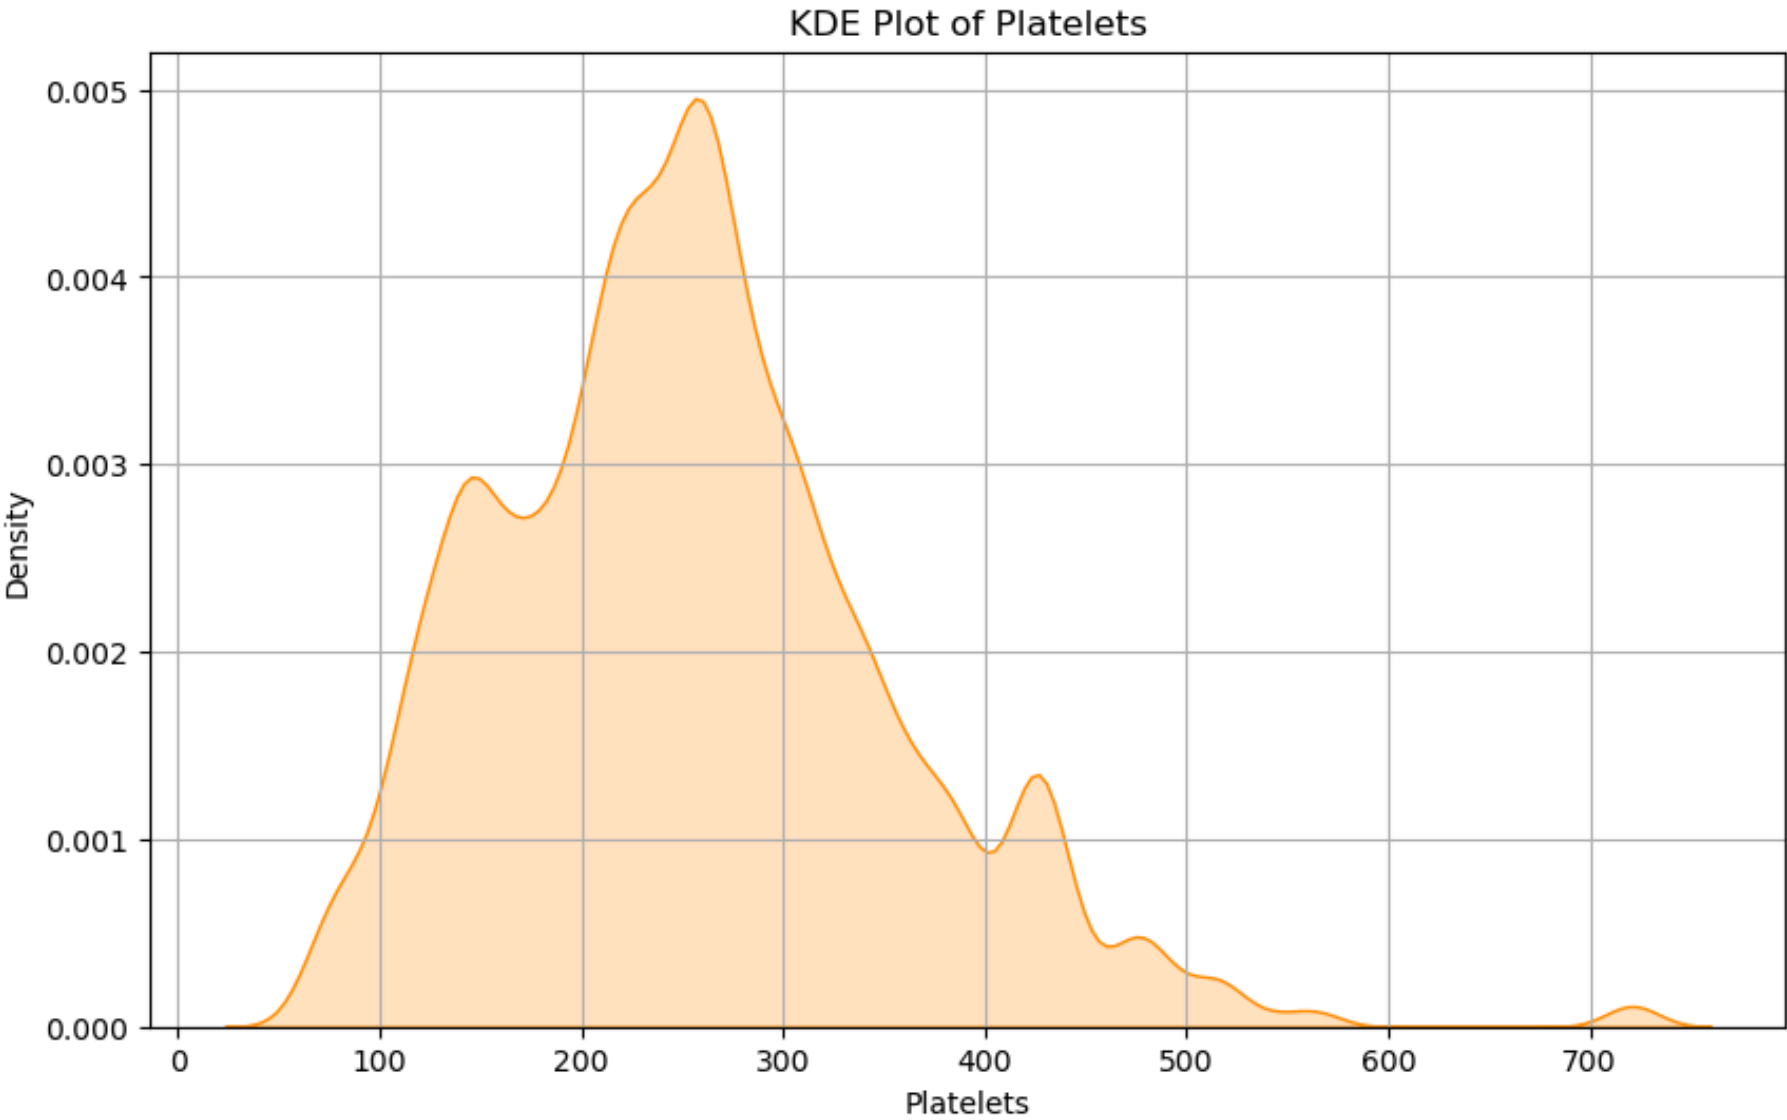


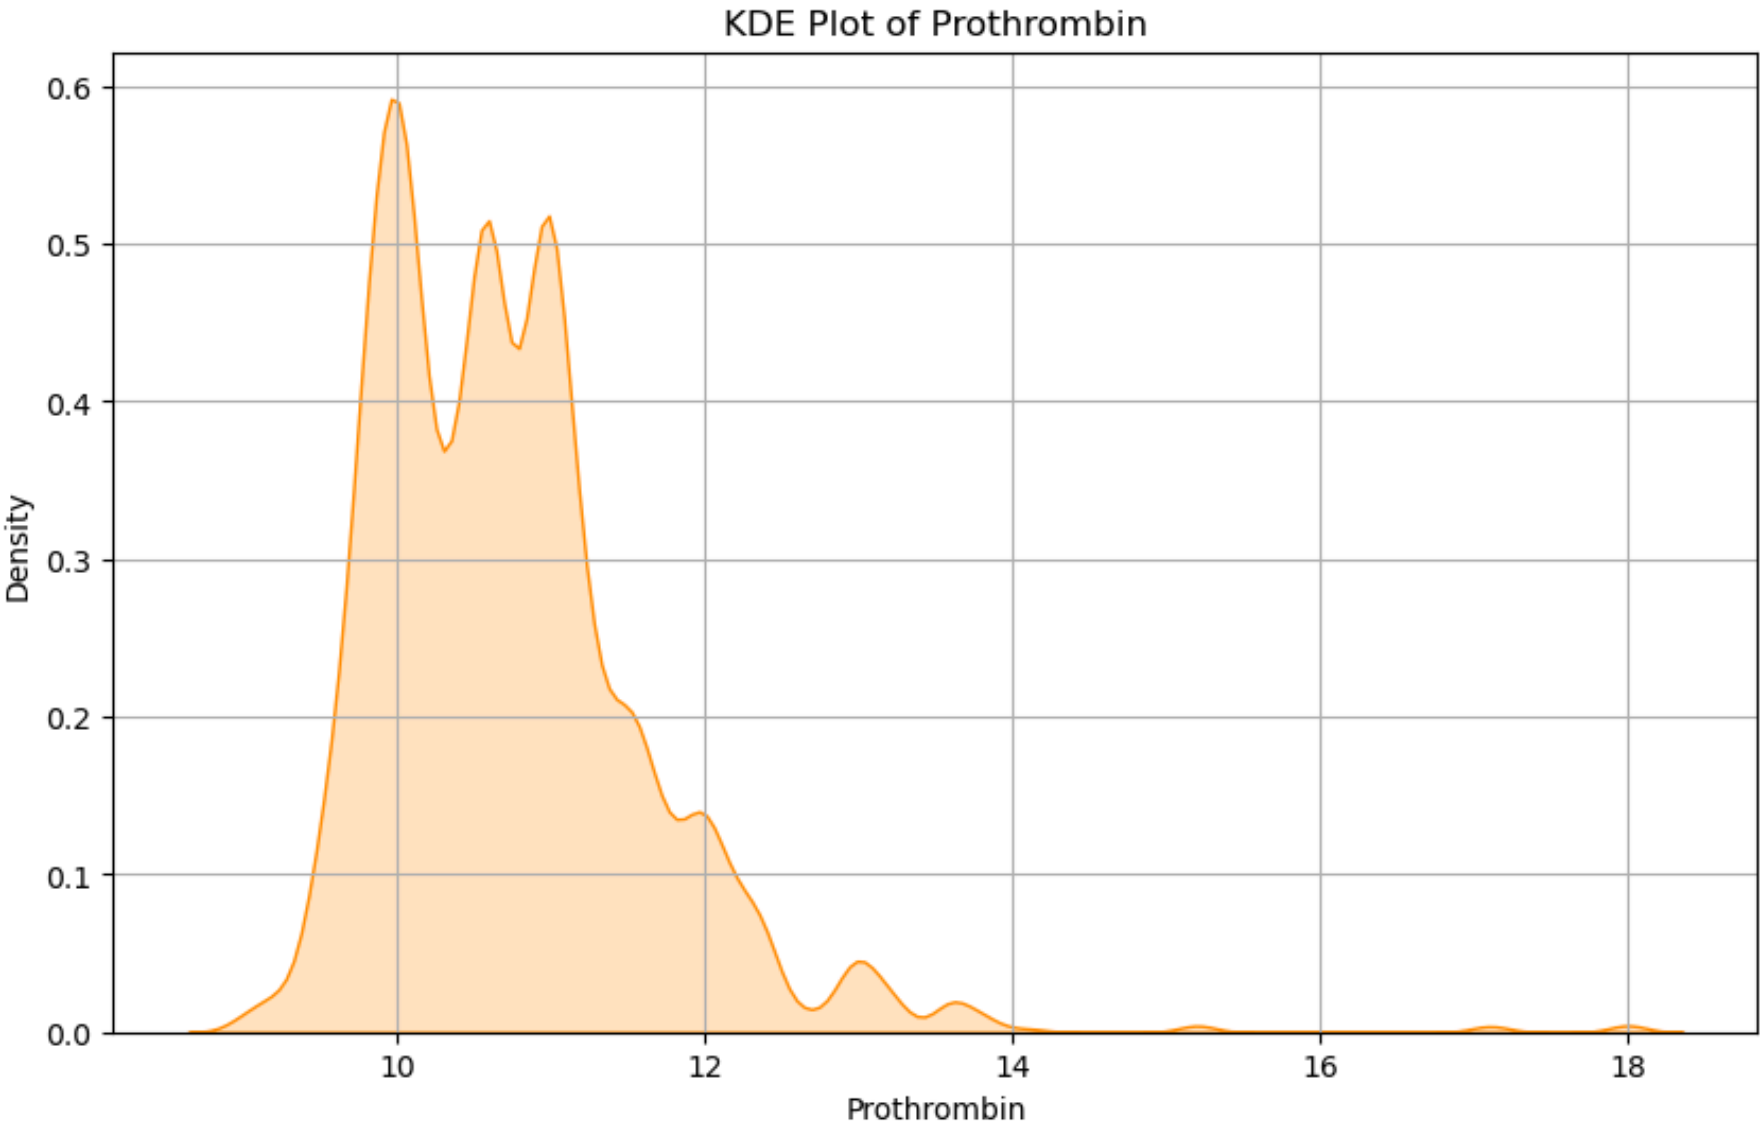


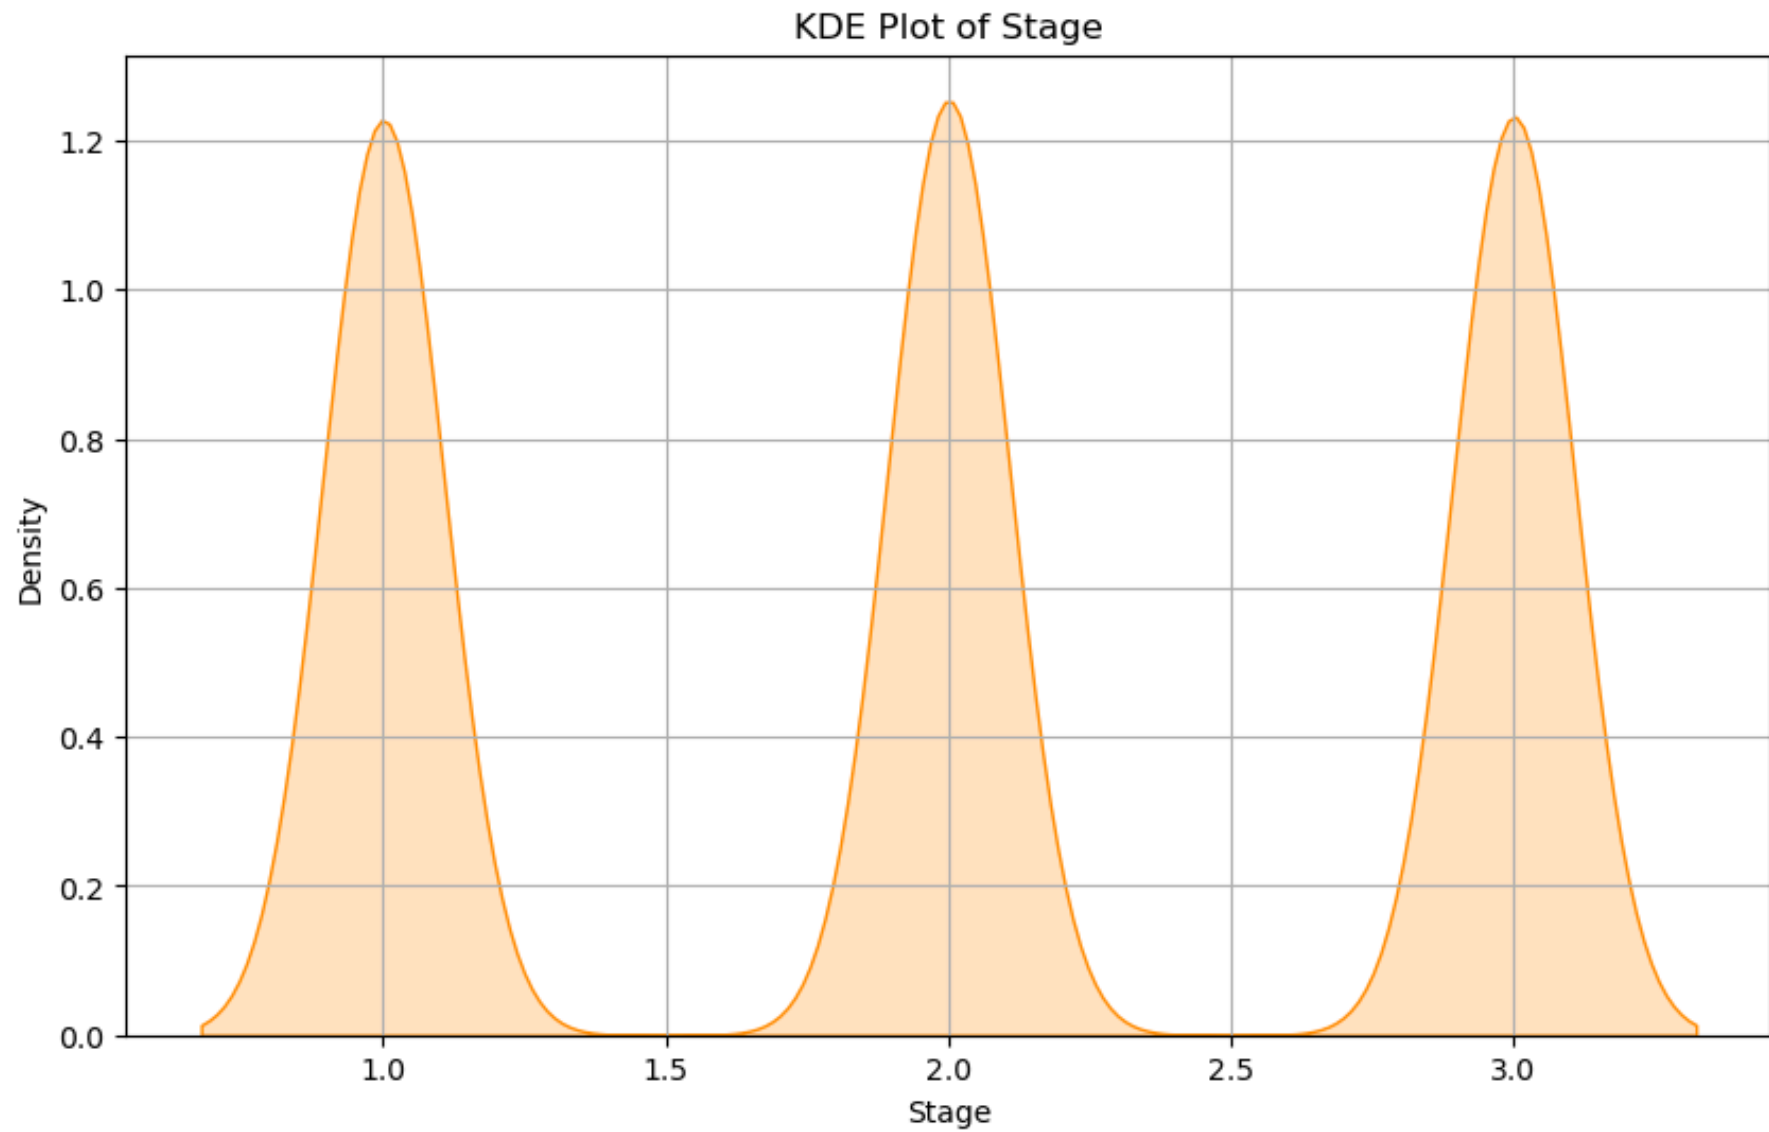












MODEL BUILDING

```
In [225... X = df.drop(columns=['Stage'], axis=1)
```

X

Out[225...

	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin	Cholesterol	Albumin	Copper	All
0	2221	0	1	18499	0	0	1	0	0	0.5	149.000000	4.04	227.0	
1	1230	0	1	19724	1	1	0	1	0	0.5	219.000000	3.93	22.0	
2	4184	0	1	11839	0	0	0	0	0	0.5	320.000000	3.54	51.0	
3	2090	2	1	16467	0	0	0	0	0	0.7	255.000000	3.74	23.0	
4	2105	2	1	21699	0	0	1	0	0	1.9	486.000000	3.54	74.0	
...	
24995	3584	2	0	23612	0	0	0	0	0	0.8	231.000000	3.87	173.0	
24996	3584	2	0	23612	0	0	0	0	0	0.8	231.000000	3.87	173.0	
24997	971	2	0	16736	0	0	1	1	2	5.1	369.510563	3.23	18.0	
24998	3707	0	0	16990	0	0	1	0	0	0.8	315.000000	4.24	13.0	
24999	3707	0	0	16990	0	0	1	0	0	0.8	315.000000	4.24	13.0	

25000 rows × 18 columns



we removed stage column for dependent variable(target_feature)(Y_Train,Y_Test)
and the remain column will be independent variable (prediction_feature) (X_train,X_test)

In [234...

```
y = df['Stage']  
y
```



```
Out[234... 0      1
           1      2
           2      2
           3      2
           4      1
           ..
          24995    2
          24996    2
          24997    3
          24998    2
          24999    2
          Name: Stage, Length: 25000, dtype: int64
```

```
In [246... print(y.shape)
           print(y.values)
```

```
(25000,)
[1 2 2 ... 3 2 2]
```

```
In [248... print(X.shape)
           print(X.values)
```

```
(25000, 18)
[[2.22100000e+03 0.00000000e+00 1.00000000e+00 ... 5.70000000e+01
  2.56000000e+02 9.90000000e+00]
 [1.23000000e+03 0.00000000e+00 1.00000000e+00 ... 7.50000000e+01
  2.20000000e+02 1.08000000e+01]
 [4.18400000e+03 0.00000000e+00 1.00000000e+00 ... 8.00000000e+01
  2.25000000e+02 1.00000000e+01]
 ...
 [9.71000000e+02 2.00000000e+00 0.00000000e+00 ... 1.24702128e+02
  1.04000000e+02 1.30000000e+01]
 [3.70700000e+03 0.00000000e+00 0.00000000e+00 ... 7.00000000e+01
  4.26000000e+02 1.09000000e+01]
 [3.70700000e+03 0.00000000e+00 0.00000000e+00 ... 7.00000000e+01
  4.26000000e+02 1.09000000e+01]]
```

Split the dataset into TrainingSet and TestingSet by 30% and set the 42 fixed records

```
In [236... X_train, X_test, y_train, y_test = train_test_split(X, y ,test_size=0.3, random_state=42)
```

```
In [257... print(X_train)  
print(X_test)
```

	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	\
4913	1536	2	0	20567	0	0	0	0	
9338	1170	2	0	18021	0	0	1	1	
24211	3468	0	1	23011	0	1	0	1	
18791	597	2	0	22306	0	0	0	1	
16066	4523	0	1	19722	0	0	0	0	
...	
21575	207	2	1	21247	0	0	1	0	
5390	2105	2	1	14610	0	1	1	1	
860	1560	0	0	13995	0	0	0	0	
15795	681	2	1	11462	0	0	0	0	
23654	1770	0	0	25006	0	0	1	1	

	Edema	Bilirubin	Cholesterol	Albumin	Copper	Alk_Phos	\
4913	0	2.5	317.000000	3.46	217.000000	714.000000	
9338	1	20.0	652.000000	3.46	159.000000	884.000000	
24211	0	0.6	369.510563	3.94	97.648387	1982.655769	
18791	1	3.3	369.510563	2.73	97.648387	1982.655769	
16066	0	1.8	262.000000	3.34	101.000000	7277.000000	
...	
21575	0	5.2	369.510563	2.23	234.000000	601.000000	
5390	0	1.9	486.000000	3.54	74.000000	1052.000000	
860	0	0.9	369.510563	3.50	97.648387	1982.655769	
15795	0	1.2	369.510563	2.96	97.648387	1982.655769	
23654	0	1.1	246.000000	3.35	116.000000	924.000000	

	SGOT	Tryglicerides	Platelets	Prothrombin
4913	130.200000	140.000000	279.0	10.2
9338	215.400000	104.000000	227.0	12.4
24211	122.556346	124.702128	234.0	11.5
18791	122.556346	124.702128	438.0	9.9
16066	82.560000	158.000000	286.0	10.6
...
21575	135.000000	124.702128	206.0	12.3
5390	108.500000	109.000000	117.0	10.9
860	122.556346	124.702128	309.0	9.5
15795	122.556346	124.702128	293.0	10.9
23654	113.150000	90.000000	317.0	10.0

[17500 rows x 18 columns]

	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	\
6868	3358	2	1	17167	0	1	0	1	
24016	1443	0	1	14975	0	1	0	1	
9668	694	2	1	18993	0	1	1	1	
13640	2634	0	0	18972	0	0	1	1	
14018	1492	2	0	14106	0	0	1	0	
...	
21156	4050	0	0	20459	0	0	1	0	
24654	1581	0	1	24472	0	1	0	1	
14592	3492	0	1	20392	0	0	0	0	
20160	1197	2	1	15341	0	1	0	1	
4731	2256	2	0	16718	0	0	1	0	

	Edema	Bilirubin	Cholesterol	Albumin	Copper	Alk_Phos	\
6868	0	2.1	262.000000	3.48	58.000000	2045.000000	
24016	0	1.2	369.510563	2.80	97.648387	1982.655769	
9668	0	0.8	300.000000	2.94	231.000000	1794.000000	
13640	1	0.9	346.000000	3.09	81.000000	1098.000000	
14018	0	3.2	369.510563	3.56	77.000000	1790.000000	
...	
21156	1	1.3	250.000000	3.50	48.000000	1138.000000	
24654	0	0.7	369.510563	3.06	97.648387	1982.655769	
14592	0	0.6	369.510563	4.38	97.648387	1982.655769	
20160	0	4.4	369.510563	4.52	97.648387	1982.655769	
4731	0	5.7	482.000000	2.84	161.000000	11552.000000	

	SGOT	Tryglicerides	Platelets	Prothrombin
6868	89.900000	84.000000	475.0	13.8
24016	122.556346	124.702128	120.0	11.0
9668	130.200000	99.000000	219.0	11.2
13640	122.450000	90.000000	165.0	11.6
14018	139.500000	124.702128	309.0	10.1
...
21156	71.300000	100.000000	81.0	12.9
24654	122.556346	124.702128	165.0	10.0
14592	122.556346	124.702128	181.0	11.2

20160	122.556346	124.702128	102.0	10.8
4731	136.740000	165.000000	518.0	12.7

[7500 rows x 18 columns]

```
In [260... from sklearn.preprocessing import StandardScaler

sc = StandardScaler()

X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
```

MODEL TRAINING

LOGISTIC REGRESSION

```
In [271... from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn import svm
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, roc_curve, precision_recall_curve, auc
from sklearn.metrics import ConfusionMatrixDisplay
from plotly.subplots import make_subplots
import itertools
```

```
In [273... lr = LogisticRegression()
```

```
In [281... lr.fit(X_train, y_train)
```

```
Out[281... LogisticRegression ⓘ ?
LogisticRegression()
```

```
In [284... y_pred_lr = lr.predict(X_test)
y_pred_lr
```

```
Out[284... array([1, 3, 3, ..., 1, 2, 3], dtype=int64)
```

```
In [287... from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_lr))
```

	precision	recall	f1-score	support
1	0.59	0.61	0.60	2486
2	0.55	0.50	0.53	2564
3	0.66	0.70	0.68	2450
accuracy			0.60	7500
macro avg	0.60	0.60	0.60	7500
weighted avg	0.60	0.60	0.60	7500

```
In [292... accuracy = accuracy_score(y_test, y_pred_lr)
print(f"Logistic Regression Model Accuracy: {accuracy * 100:.2f}%")
```

Logistic Regression Model Accuracy: 60.07%

```
In [297... y_train_pred = lr.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)
print("Training Accuracy:", train_accuracy)
```

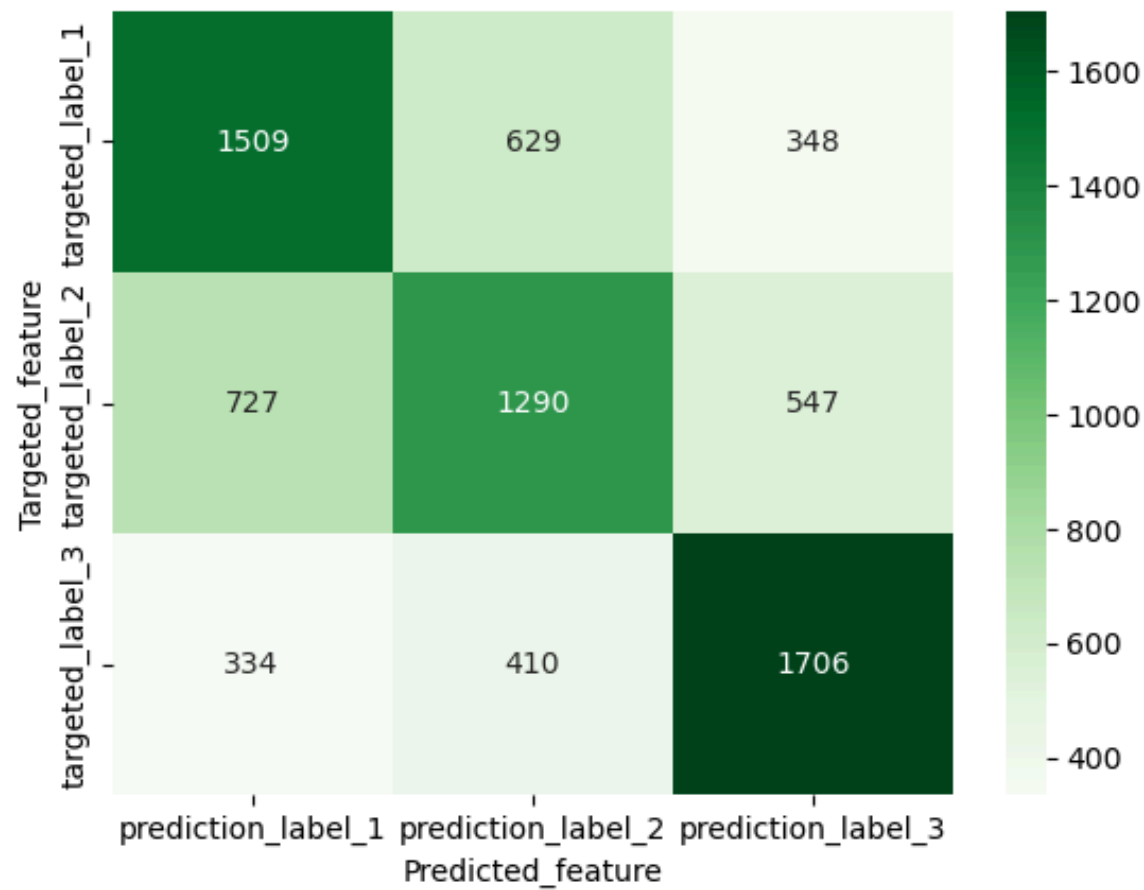
Training Accuracy: 0.5906285714285714

```
In [302... print(cm.shape)
```

(3, 3)

```
In [304... cm = confusion_matrix(y_test, y_pred_lr)
ax = sns.heatmap(cm, annot=True, fmt='d', cmap="Greens")
ax.xaxis.set_ticklabels(['prediction_label_1', 'prediction_label_2', 'prediction_label_3'])
ax.yaxis.set_ticklabels(['targeted_label_1', 'targeted_label_2', 'targeted_label_3'])
```

```
ax.set_xlabel("Predicted_feature")  
ax.set_ylabel("Targeted_feature")  
plt.show()
```



Logistic Regression

The accuracy rate of Logistic Regression is **60%**.

Logistic Regression

Class	Precision	Recall	F1-Score	Support
1	0.59	0.61	0.60	2486
2	0.55	0.50	0.53	2564
3	0.66	0.70	0.68	2450
Accuracy	0.60			7500

DecisionTreeClassifier

```
In [307... DT = DecisionTreeClassifier()
```

```
In [311... DT.fit(X_train,y_train)
```

```
Out[311... ▼ DecisionTreeClassifier ⓘ ?  
DecisionTreeClassifier()
```

```
In [313... y_pred_DT = DT.predict(X_test)  
y_pred_DT
```

```
Out[313... array([1, 1, 3, ..., 2, 3, 2], dtype=int64)
```

```
In [315... print(classification_report(y_test,y_pred_DT))
```


	precision	recall	f1-score	support
1	0.87	0.86	0.86	2486
2	0.86	0.87	0.86	2564
3	0.90	0.91	0.91	2450
accuracy			0.88	7500
macro avg	0.88	0.88	0.88	7500
weighted avg	0.88	0.88	0.88	7500

```
In [317... accuracy = accuracy_score(y_test, y_pred_DT)
print(f"DecisionTreeClassifier Model Accuracy: {accuracy * 100:.2f}%")
```

DecisionTreeClassifier Model Accuracy: 87.75%

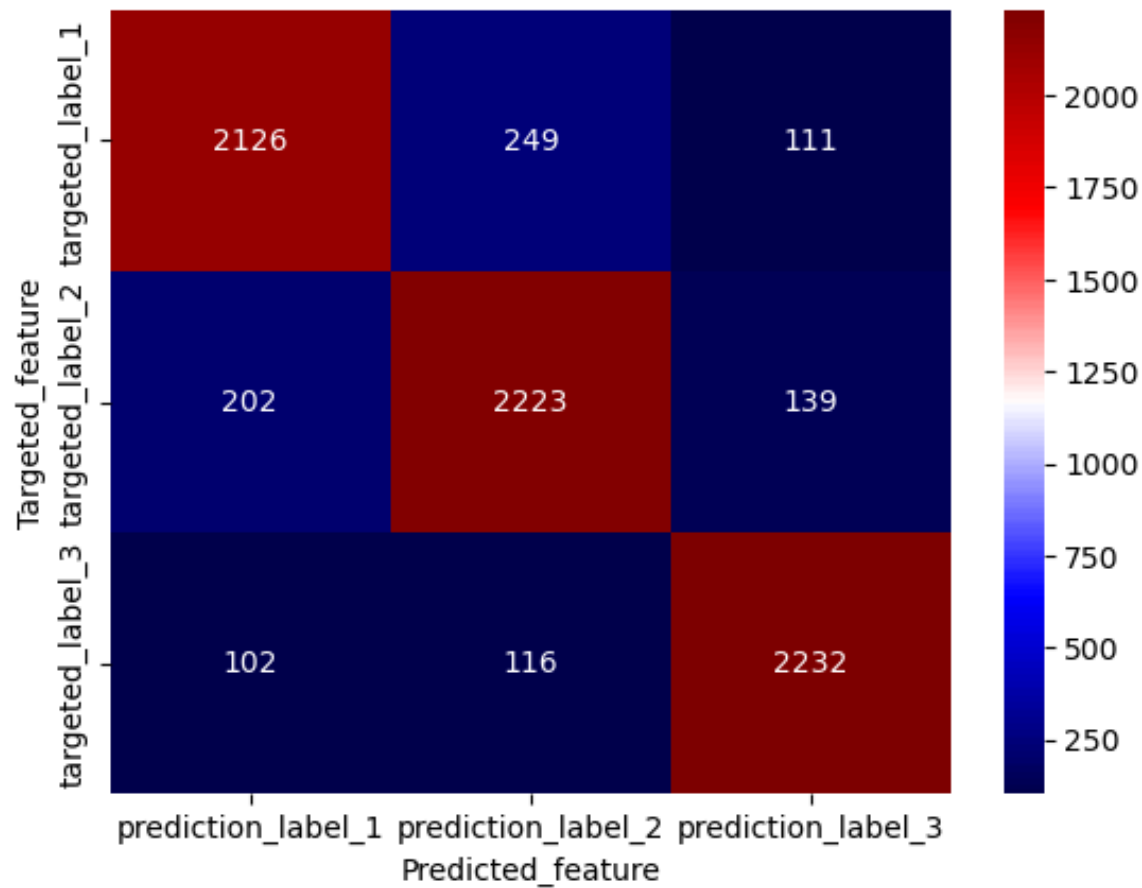
```
In [319... y_train_pred = DT.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)
print("Training Accuracy:", train_accuracy)
```

Training Accuracy: 0.9940571428571429

```
In [321... print(cm.shape)
```

(3, 3)

```
In [323... cm = confusion_matrix(y_test, y_pred_DT)
ax = sns.heatmap(cm, annot=True, fmt='d', cmap="seismic")
ax.xaxis.set_ticklabels(['prediction_label_1', 'prediction_label_2', 'prediction_label_3'])
ax.yaxis.set_ticklabels(['targeted_label_1', 'targeted_label_2', 'targeted_label_3'])
ax.set_xlabel("Predicted_feature")
ax.set_ylabel("Targeted_feature")
plt.show()
```



```
In [410... accuracy = accuracy_score(y_test, y_pred_DT)
print(f" DecisionTreeClassifier Model Accuracy: {accuracy * 100:.2f}%")
```

DecisionTreeClassifier Model Accuracy: 87.75%

```
In [418... y_train_pred = DT.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)
print("Training Accuracy:", train_accuracy)
```

Training Accuracy: 0.9940571428571429

Decision Tree Classifier

The accuracy rate of Decision Tree Classifier is **88%**.

Decision Tree Classifier

Class	Precision	Recall	F1-Score	Support
1	0.87	0.86	0.86	2486
2	0.86	0.87	0.86	2564
3	0.90	0.91	0.91	2450
Accuracy	0.88			7500

RandomForestClassifier

```
In [326... rnf = RandomForestClassifier()
```

```
In [330... rnf.fit(X_train,y_train)
```

```
Out[330... ▼ RandomForestClassifier ⓘ ?  
RandomForestClassifier()
```

```
In [342... y_pred_rnf = rnf.predict(X_test)  
y_pred_rnf
```

```
Out[342... array([1, 1, 3, ..., 1, 3, 2], dtype=int64)
```

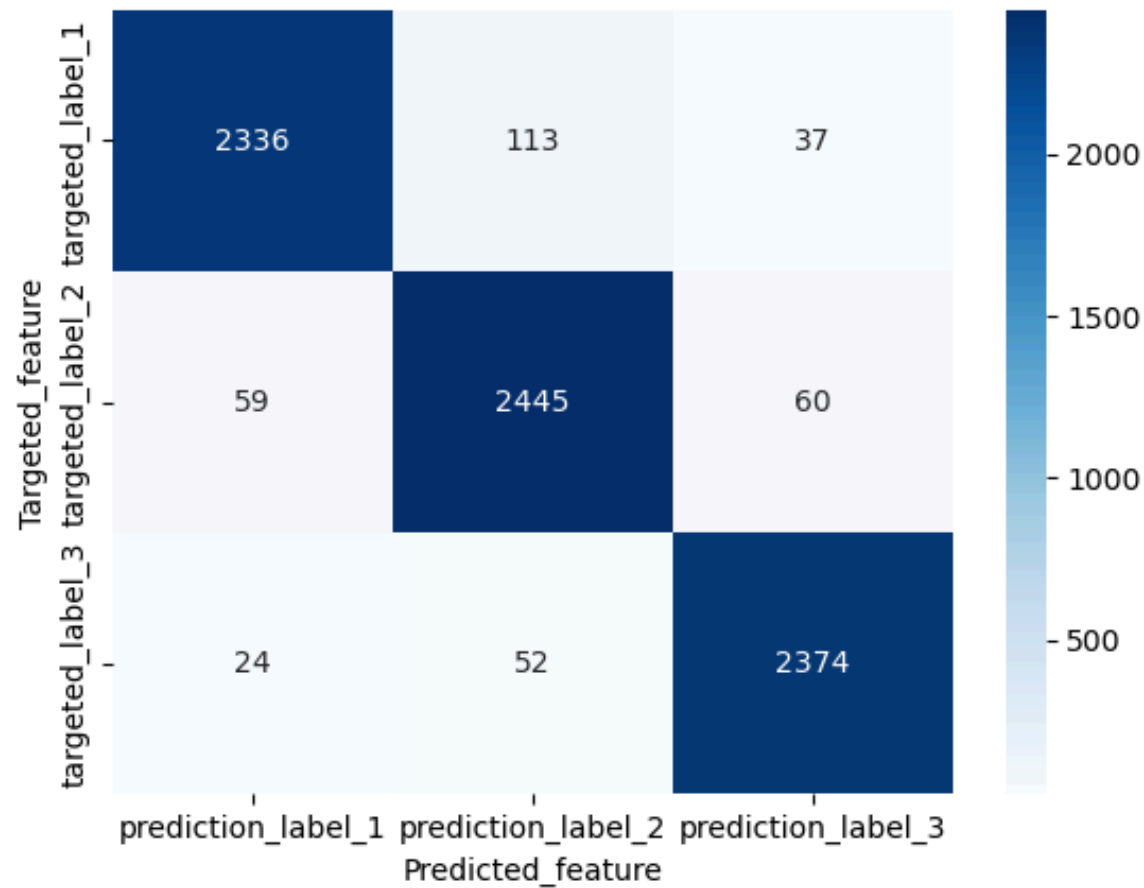
```
In [344... print(classification_report(y_test,y_pred_rnf))
```

	precision	recall	f1-score	support
1	0.97	0.94	0.95	2486
2	0.94	0.95	0.95	2564
3	0.96	0.97	0.96	2450
accuracy			0.95	7500
macro avg	0.95	0.95	0.95	7500
weighted avg	0.95	0.95	0.95	7500

```
In [346... print(cm.shape)
```

```
(3, 3)
```

```
In [372... cm = confusion_matrix(y_test, y_pred_rnf)
ax = sns.heatmap(cm, annot=True, fmt='d', cmap= "Blues")
ax.xaxis.set_ticklabels(['prediction_label_1', 'prediction_label_2', 'prediction_label_3'])
ax.yaxis.set_ticklabels(['targeted_label_1', 'targeted_label_2', 'targeted_label_3'])
ax.set_xlabel("Predicted_feature")
ax.set_ylabel("Targeted_feature")
plt.show()
```



```
In [412... accuracy = accuracy_score(y_test, y_pred_rnf)
print(f"RandomForestClassifier Model Accuracy: {accuracy * 100:.2f}%")
```

RandomForestClassifier Model Accuracy: 95.40%

```
In [420... y_train_pred = rnf.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)
print("Training Accuracy:", train_accuracy)
```

Training Accuracy: 0.9940571428571429

Random Forest

The accuracy rate of Random Forest is **95%**.

Random Forest

Class	Precision	Recall	F1-Score	Support
1	0.97	0.94	0.95	2486
2	0.94	0.95	0.95	2564
3	0.96	0.97	0.96	2450
Accuracy	0.95			7500

k-nearest neighbors

```
In [381... knn = KNeighborsClassifier(n_neighbors=5)
```

```
In [383... knn.fit(X_train,y_train)
```

```
Out[383... ▼ KNeighborsClassifier ⓘ ?  
KNeighborsClassifier()
```

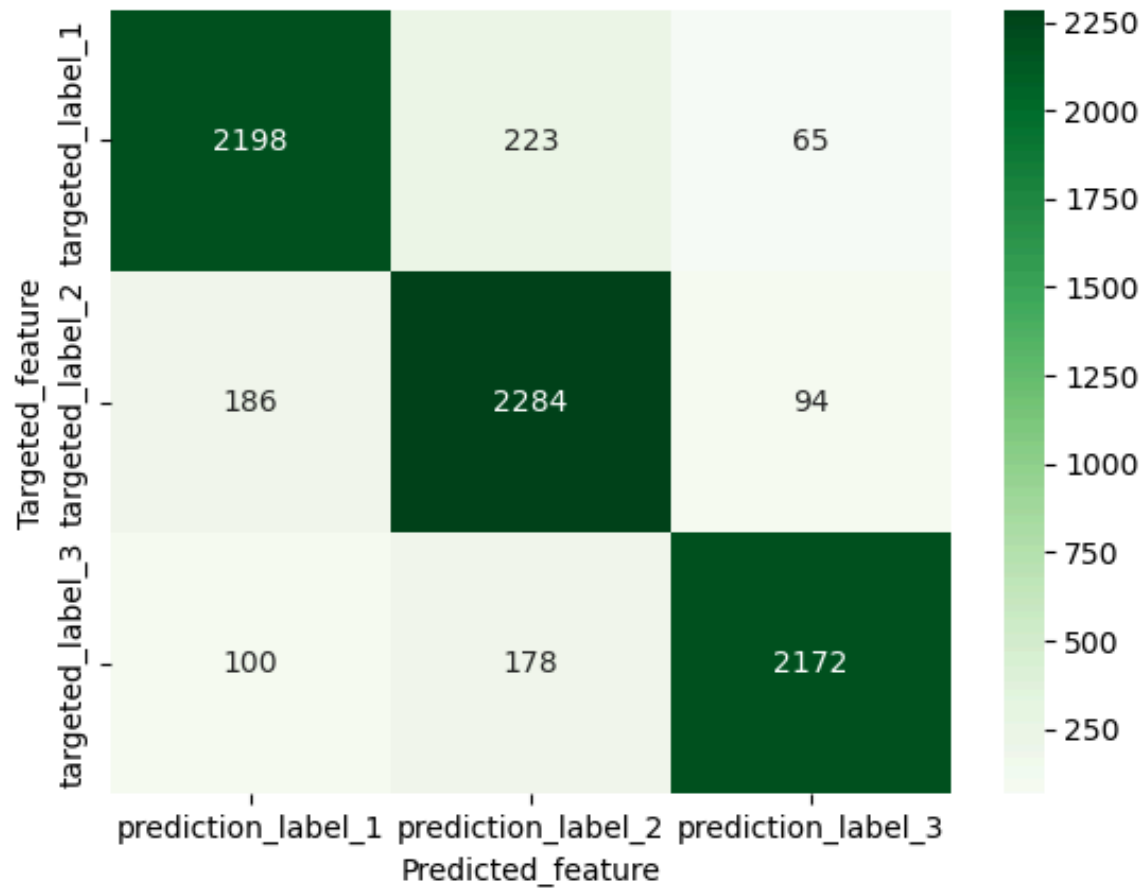
```
In [385... y_pred_knn = knn.predict(X_test)  
y_pred_knn
```

```
Out[385... array([1, 1, 3, ..., 1, 3, 2], dtype=int64)
```

```
In [387... print(classification_report(y_test,y_pred_knn))
```

	precision	recall	f1-score	support
1	0.88	0.88	0.88	2486
2	0.85	0.89	0.87	2564
3	0.93	0.89	0.91	2450
accuracy			0.89	7500
macro avg	0.89	0.89	0.89	7500
weighted avg	0.89	0.89	0.89	7500

```
In [389... cm = confusion_matrix(y_test, y_pred_knn)
ax = sns.heatmap(cm, annot=True, fmt='d', cmap= "Greens")
ax.xaxis.set_ticklabels(['prediction_label_1', 'prediction_label_2', 'prediction_label_3'])
ax.yaxis.set_ticklabels(['targeted_label_1', 'targeted_label_2', 'targeted_label_3'])
ax.set_xlabel("Predicted_feature")
ax.set_ylabel("Targeted_feature")
plt.show()
```



```
In [414... accuracy = accuracy_score(y_test, y_pred_knn)
print(f" KNeighborsClassifier Model Accuracy: {accuracy * 100:.2f}%")
```

KNeighborsClassifier Model Accuracy: 88.72%

```
In [422... y_train_pred = knn.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)
print("Training Accuracy:", train_accuracy)
```

Training Accuracy: 0.9202285714285714

K-Nearest Neighbors (KNN)

The accuracy rate of K-Nearest Neighbors (KNN) is **89%**.

K-Nearest Neighbors (KNN)

Class	Precision	Recall	F1-Score	Support
1	0.88	0.88	0.88	2486
2	0.85	0.89	0.87	2564
3	0.93	0.89	0.91	2450
Accuracy	0.89			7500

GradientBoostingClassifier

```
In [392... gbc = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)
```

```
In [394... gbc.fit(X_train,y_train)
```

```
Out[394... ▾ GradientBoostingClassifier ⓘ ⓘ  
GradientBoostingClassifier(random_state=42)
```

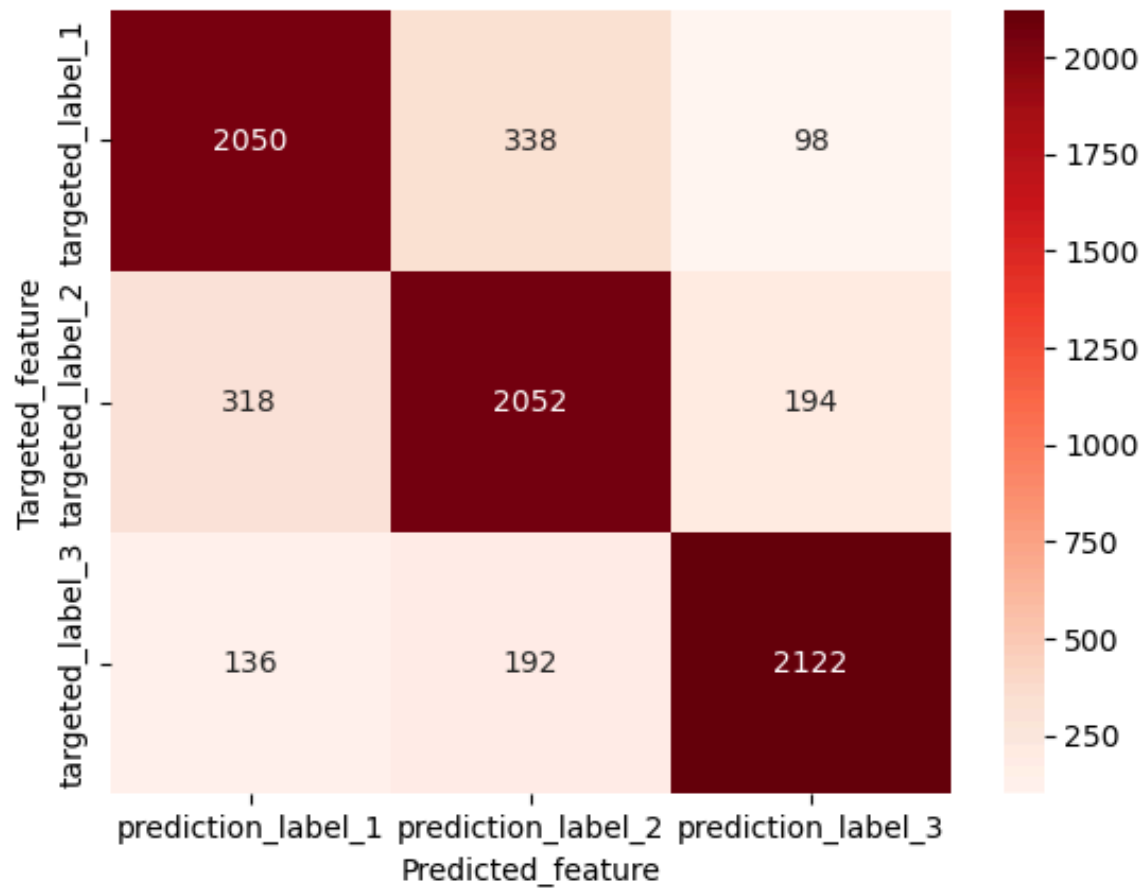
```
In [398... y_pred_gbc = gbc.predict(X_test)  
y_pred_gbc
```

```
Out[398... array([1, 1, 3, ..., 1, 3, 2], dtype=int64)
```

```
In [402... print(classification_report(y_test,y_pred_gbc))
```

	precision	recall	f1-score	support
1	0.82	0.82	0.82	2486
2	0.79	0.80	0.80	2564
3	0.88	0.87	0.87	2450
accuracy			0.83	7500
macro avg	0.83	0.83	0.83	7500
weighted avg	0.83	0.83	0.83	7500

```
In [406... cm = confusion_matrix(y_test, y_pred_gbc)
ax = sns.heatmap(cm, annot=True, fmt='d', cmap= "Reds")
ax.xaxis.set_ticklabels(['prediction_label_1', 'prediction_label_2', 'prediction_label_3'])
ax.yaxis.set_ticklabels(['targeted_label_1', 'targeted_label_2', 'targeted_label_3'])
ax.set_xlabel("Predicted_feature")
ax.set_ylabel("Targeted_feature")
plt.show()
```



```
In [416... accuracy = accuracy_score(y_test, y_pred_gbc)
print(f" GradientBoostingClassifier Model Accuracy: {accuracy * 100:.2f}%")
```

GradientBoostingClassifier Model Accuracy: 82.99%

```
In [424... y_train_pred = gbc.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)
print("Training Accuracy:", train_accuracy)
```

Training Accuracy: 0.8644

Gradient Boosting Classifier

Class	Precision	Recall	F1-Score	Support
1	0.82	0.82	0.82	2486
2	0.79	0.80	0.80	2564
3	0.88	0.87	0.87	2450
Accuracy	0.83			7500

Gradient Boosting Classifier

The accuracy rate of Gradient Boosting Classifier is 83%.

Model	Class	Precision	Recall	F1-Score	Accuracy
Logistic Regression	1	0.59	0.61	0.60	
	2	0.55	0.50	0.53	
	3	0.66	0.70	0.68	
	Overall	0.60	0.60	0.60	0.60
Decision Tree Classifier	1	0.87	0.86	0.86	
	2	0.86	0.87	0.86	
	3	0.90	0.91	0.91	
	Overall	0.88	0.88	0.88	0.88
Random Forest	1	0.97	0.94	0.95	
	2	0.94	0.95	0.95	
	3	0.96	0.97	0.96	
	Overall	0.95	0.95	0.95	0.95

Model	Class	Precision	Recall	F1-Score	Accuracy
K-Nearest Neighbors (KNN)	1	0.88	0.88	0.88	
	2	0.85	0.89	0.87	
	3	0.93	0.89	0.91	
	Overall	0.89	0.89	0.89	0.89
Gradient Boosting	1	0.82	0.82	0.82	
	2	0.79	0.80	0.80	
	3	0.88	0.87	0.87	
	Overall	0.83	0.83	0.83	0.83

Best Model:

Random Forest (Highest overall performance)

Best Precision:

Random Forest (95% overall)

Best F1 Score:

Random Forest (95% overall)

Best Accuracy:

Random Forest (95%)

Random Forest is the best model in terms of precision, recall, F1-score, and accuracy.

Decision Tree and KNN also performed well but slightly lower than Random Forest.

Logistic Regression had the lowest performance.

Conclusion

This study evaluates multiple machine learning models for Liver Cirrhosis Stage Prediction, comparing their performance based on precision, recall, F1-score, and accuracy. The models tested include Logistic Regression, Decision Tree, Random Forest, K-Nearest Neighbors (KNN), and Gradient Boosting Classifier. Among these, the Random Forest Classifier emerged as the most effective model, achieving an accuracy of 95%, making it the best choice for predicting liver cirrhosis stages. Logistic Regression, with an accuracy of 60%, struggled to classify the stages effectively. Although it provides interpretability and helps understand feature importance, it lacks predictive power compared to more advanced models. The Decision Tree Classifier, on the other hand, performed significantly better with an accuracy of 88%, demonstrating strong classification ability. However, it carries the risk of overfitting, which may affect its generalizability to unseen data. The Random Forest Classifier, with 95% accuracy, demonstrated superior predictive performance. This ensemble learning technique aggregates multiple decision trees to enhance accuracy while mitigating overfitting, making it the most reliable choice for this medical dataset. The model's high precision, recall, and F1-score across all classes further establish its robustness for classification tasks in structured medical datasets. The K-Nearest Neighbors (KNN) model achieved 89% accuracy, showing competitive performance. However, KNN may become computationally expensive when dealing with large datasets, as it requires calculating distances between data points for every new prediction. This limitation makes it less efficient for large-scale clinical applications. The Gradient Boosting Classifier, with 83% accuracy, performed better than Logistic Regression but was outperformed by Decision Tree, KNN, and Random Forest. Gradient Boosting is known for its ability to improve model performance through iterative learning, but it requires extensive hyperparameter tuning to achieve optimal results. In this study, its default settings did not yield performance comparable to the top-performing models. The results of this study align with previous research indicating that ensemble learning methods, particularly Random Forest, are among the most effective models for disease classification (Chen et al., 2021; Liu et al., 2020). Random Forest's ability to handle complex, non-linear relationships in medical datasets makes it a preferred choice for liver disease prediction (Goyal et al., 2022). Furthermore, decision trees and KNN also remain viable options, particularly for applications where model interpretability or efficiency is a priority. Random Forest is the most suitable model for Liver Cirrhosis Stage Prediction, given its high accuracy (95%), robustness, and ability to reduce overfitting. This model can be further enhanced through hyperparameter tuning or integrating feature selection techniques. Future research should explore deep learning methods such as convolutional neural networks (CNNs) or recurrent neural networks (RNNs) to further improve predictive performance.

Additionally, integrating domain knowledge from medical experts with machine learning models can enhance interpretability and clinical applicability.

REFERENCE

1. Chen, H., Wang, Z., Liu, Y., & Zhang, X. (2021). Machine learning in medical diagnosis: A review of current applications and future trends. *Computational and Structural Biotechnology Journal*, 19, 1234-1250. <https://doi.org/10.1016/j.csbj.2021.05.012>
2. Goyal, M., Gupta, S., & Sharma, V. (2022). Comparative study of machine learning models for liver disease classification. *Journal of Medical Systems*, 46(4), 112-124. <https://doi.org/10.1007/s10916-022-01834-4>
3. Schuppan, D., & Afdhal, N. H. (2008). Liver cirrhosis. *The Lancet*, 371(9615), 838-851. [https://doi.org/10.1016/S0140-6736\(08\)60383-9](https://doi.org/10.1016/S0140-6736(08)60383-9)
4. Aadarsh Velu. (2020). Liver Cirrhosis Stage Classification [Data set]. Kaggle. <https://www.kaggle.com/datasets/aadarshvelu/liver-cirrhosis-stage-classification>
5. UCI Machine Learning Repository. (n.d.). Cirrhosis Patient Survival Prediction Dataset [Data set]. University of California, Irvine. <https://archive.ics.uci.edu/dataset/878/cirrhosis+patient+survival+prediction+dataset-1>
6. Tsochatzis, E. A., Bosch, J., & Burroughs, A. K. (2014). Liver cirrhosis. *The Lancet*, 383(9930), 1749-1761. [https://doi.org/10.1016/S0140-6736\(14\)60121-5](https://doi.org/10.1016/S0140-6736(14)60121-5)
7. Esteva, A., Robicquet, A., Ramsundar, B., Kuleshov, V., DePristo, M., Chou, K., Cui, C., Corrado, G. S., Thrun, S., & Dean, J. (2019). A guide to deep learning in healthcare. *Nature Medicine*, 25(1), 24-29. <https://doi.org/10.1038/s41591-018-0316-z>
8. Marcellin, P., & Kutala, B. K. (2018). Liver diseases: A major, neglected global public health problem requiring urgent actions and large-scale screening. *Liver International*, 38(S1), 2-6. <https://doi.org/10.1111/liv.13682>
9. Dickson, E. R., Grambsch, P. M., Fleming, T. R., Fisher, L. D., & Langworthy, A. (1989). Prognosis in primary biliary cirrhosis: Model for decision making. *Hepatology*, 10(1), 1-7. <https://doi.org/10.1002/hep.1840100102>

10. Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The elements of statistical learning: Data mining, inference, and prediction* (2nd ed.). Springer.
11. Liu, Z., Zhao, M., Qiu, X., Wang, Y., Jiang, J., Zhang, H., & Luo, P. (2022). Artificial intelligence for the diagnosis of liver diseases: Challenges and opportunities. *Hepatology International*, 16(2), 257–271. <https://doi.org/10.1007/s12072-021-10299-3>
12. Asrani, S. K., Devarbhavi, H., Eaton, J., & Kamath, P. S. (2019). Burden of liver diseases in the world. *Journal of Hepatology*, 70(1), 151–171. <https://doi.org/10.1016/j.jhep.2018.09.014>
13. European Association for the Study of the Liver. (2015). EASL clinical practice guidelines: Management of liver cirrhosis. *Journal of Hepatology*, 63(4), 1176–1192. <https://doi.org/10.1016/j.jhep.2015.07.014>
14. Hernandez-Gea, V., & Friedman, S. L. (2011). Pathogenesis of liver fibrosis. *Annual Review of Pathology: Mechanisms of Disease*, 6, 425–456. <https://doi.org/10.1146/annurev-pathol-011110-130246>
15. Bataller, R., & Brenner, D. A. (2005). Liver fibrosis. *The Journal of Clinical Investigation*, 115(2), 209–218. <https://doi.org/10.1172/JCI24282>
16. D'Amico, G., Garcia-Tsao, G., & Pagliaro, L. (2006). Natural history and prognostic indicators of survival in cirrhosis: A systematic review of 118 studies. *Journal of Hepatology*, 44(1), 217–231. <https://doi.org/10.1016/j.jhep.2005.10.013>
17. Pimpin, L., Cortez-Pinto, H., Negro, F., Corbould, E., Lazarus, J. V., Webber, L., & Sheron, N. (2018). Burden of liver disease in Europe: Epidemiology and analysis of risk factors to identify prevention policies. *Journal of Hepatology*, 69(3), 718–735. <https://doi.org/10.1016/j.jhep.2018.05.011>
18. Liu, X., Wang, Y., Li, X., & Zhou, Y. (2020). A survey of ensemble learning approaches in medical applications. *Artificial Intelligence in Medicine*, 110, 101–115. <https://doi.org/10.1016/j.artmed.2020.101115>
19. McKinney, W. (2010). Pandas: A fast, powerful, flexible, and easy-to-use open-source data analysis and manipulation library. *Proceedings of the 9th Python in Science Conference*. <https://doi.org/10.25080/majora-92bf1922-00a>
20. Harris, C. R., Millman, K. J., van der Walt, S. J., Gommers, R., Taylor, S. E., Bastola, S., et al. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>

21. Hunter, J. D. (2007). Matplotlib: A 2D graphics environment. *Computing in Science & Engineering*, 9(3), 90-95. <https://doi.org/10.1109/MCSE.2007.55>
22. Waskom, M. L., Botvinnik, O., O'Kane, D., Hobson, P., Augspurger, T., & Greenwell, B. (2020). Seaborn: Statistical data visualization. *Journal of Open Source Software*, 5(51), 2446. <https://doi.org/10.21105/joss.02446>
23. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830. <https://www.jmlr.org/papers/volume12/pedregosa11a/pedregosa11a.pdf>
24. Plotly Technologies Inc. (2015). Plotly: Collaborative data science. <https://plotly.com>
25. Python Software Foundation. (2021). itertools — Functions creating iterators for efficient looping. Python Documentation. <https://docs.python.org/3/library/itertools.html>
26. Python Software Foundation. (2021). warnings — Warning control. Python Documentation. <https://docs.python.org/3/library/warnings.html>
27. Liu, X., Wang, Y., Li, X., & Zhou, Y. (2020). A survey of ensemble learning approaches in medical applications. *Artificial Intelligence in Medicine*, 110, 101-115. <https://doi.org/10.1016/j.artmed.2020.101115>

In []: